

# Direct Pose Estimation and Refinement

Hatem Alismail  
halismai@cs.cmu.edu

Thesis Oral  
July 28, 2016

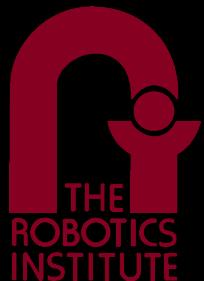
Brett Browning (Co-chair)

Simon Lucey (Co-Chair)

Michael Kaess

Martial Hebert

Ian D. Reid, The University of Adelaide



# Promised Completed Work

## Chapters 5 and 6

Analysis of the proposed idea of  
direct alignment of binary  
descriptors (Bit-Planes)

# Promised Completed Work

## Chapters 5 and 6

Analysis of the proposed idea of direct alignment of binary descriptors (Bit-Planes)

## Chapter 6

Applications to visual odometry in poorly lit underground mines

# Promised Completed Work

## Chapters 5 and 6

Analysis of the proposed idea of direct alignment of binary descriptors (Bit-Planes)

## Chapter 6

Applications to visual odometry in poorly lit underground mines

## Chapter 7

Direct (photometric) refinement of pose and structure jointly  
*[Bundle adjustment without correspondences]*

# Promised Completed Work

## Chapters 5 and 6

Analysis of the proposed idea of direct alignment of binary descriptors (Bit-Planes)

## Chapter 6

Applications to visual odometry in poorly lit underground mines

## Chapter 7

Direct (photometric) refinement of pose and structure jointly  
*[Bundle adjustment without correspondences]*

## Chapter 4 (extra)

Evaluation of important details in direct visual odometry

# Contributions

- Robust and real-time pose estimation in challenging environments
  - Near darkness, blur, specular reflections, ...
- First formulation of a direct (photometric) bundle adjustment for VSLAM
  - No correspondences required

# Geometric Pose Estimation

One of the most successful applications of  
Computer Vision

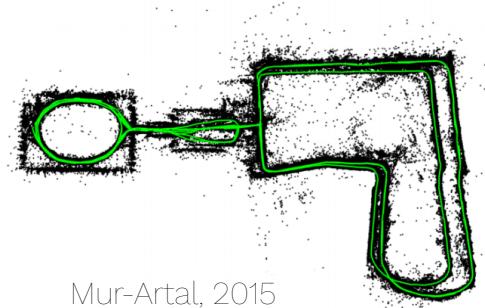
AR/VR



Structure-from-Motion (SfM)



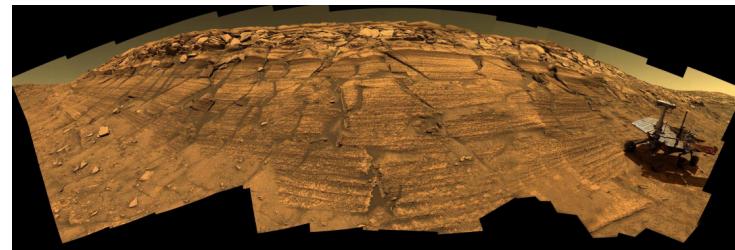
Visual Odometry & VSLAM



Robotic Inspection



Space Exploration



Maimone, 2004

<http://www.superdroidrobots.com/>

Autonomous Driving



Google

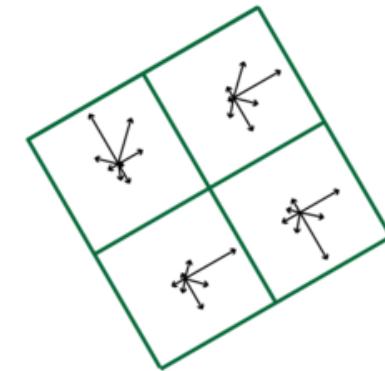
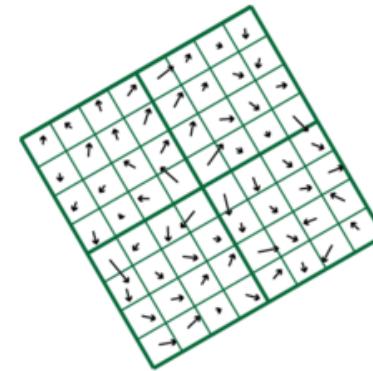
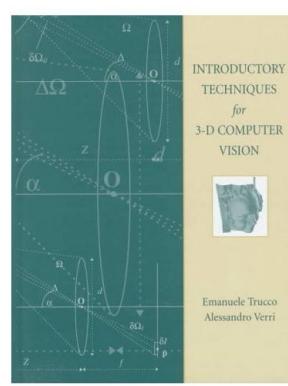
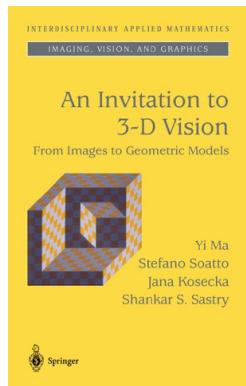
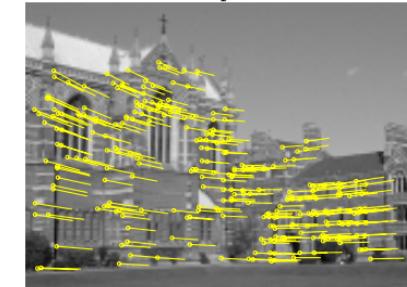
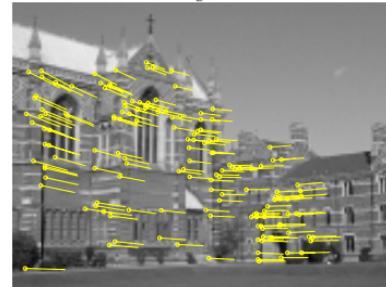
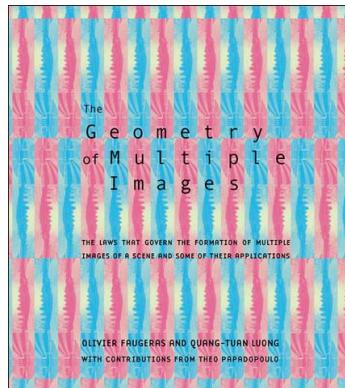
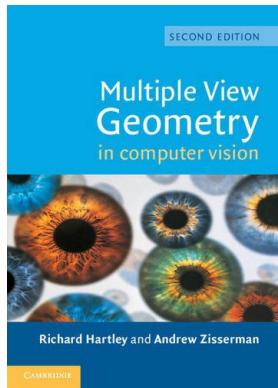
# Pose Estimation too



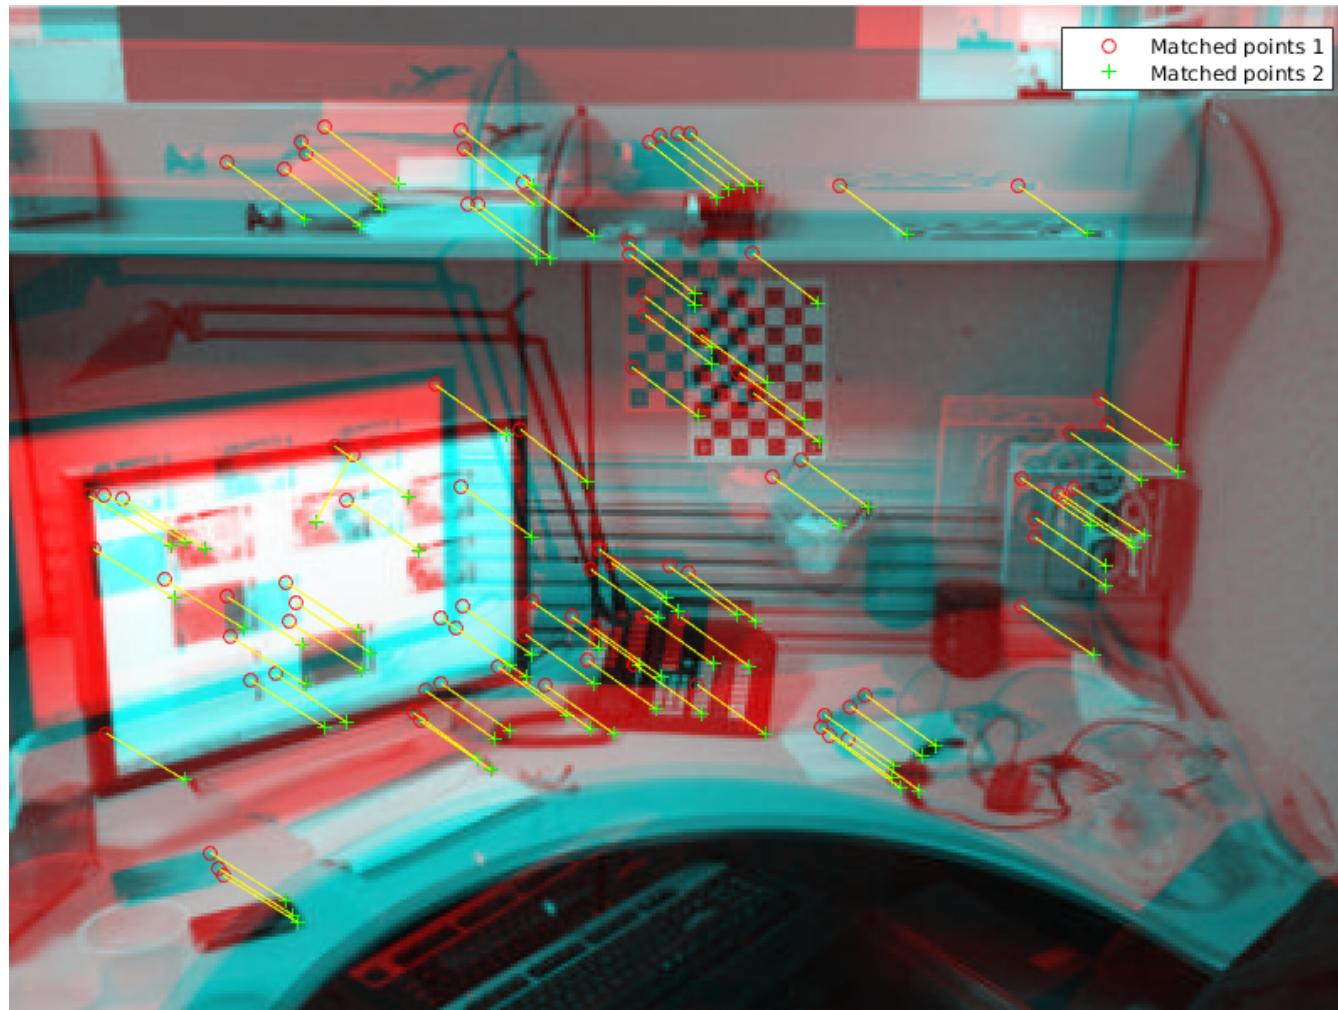
Image credit [thewardrobedoor.com](http://thewardrobedoor.com) / Google

# Feature-based Pipeline

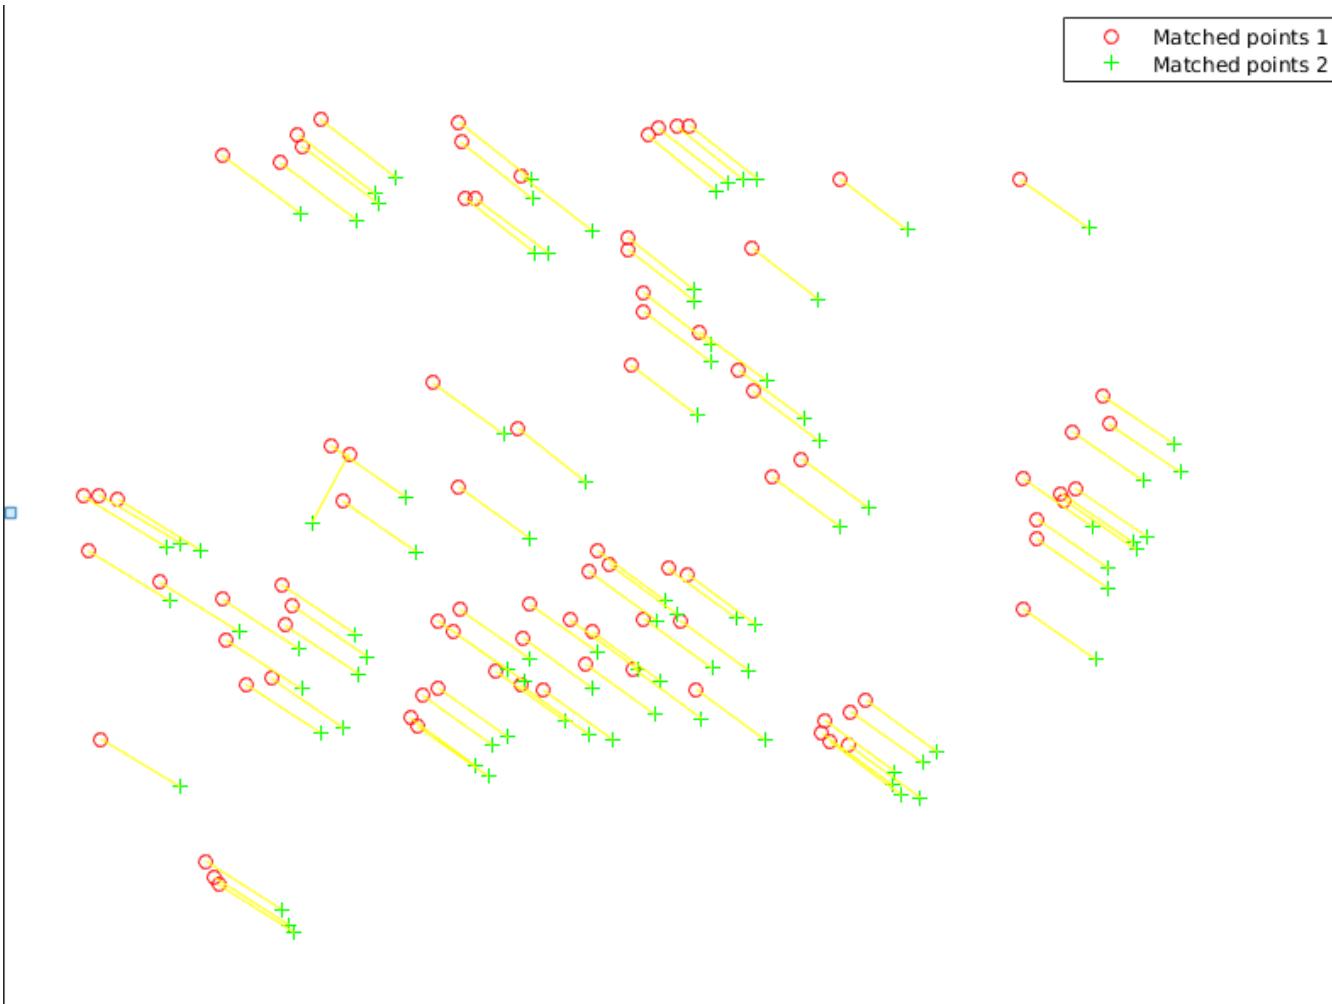
Major Contributor to the success of geometric pose estimation



# Feature-based Pipeline

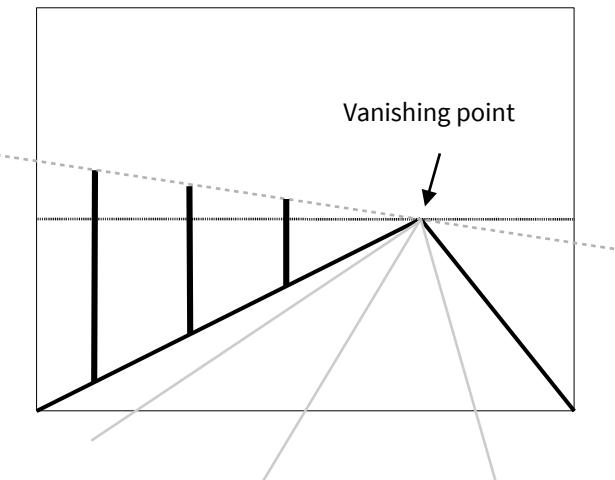


# Feature-based Pipeline



# Feature-based Pipeline

Natural transition from images to geometry



Viewpoint  
invariance



Mikolajczyk, 2007

Illumination  
invariance



Mikolajczyk, 2007

# Limitations of Keypoints

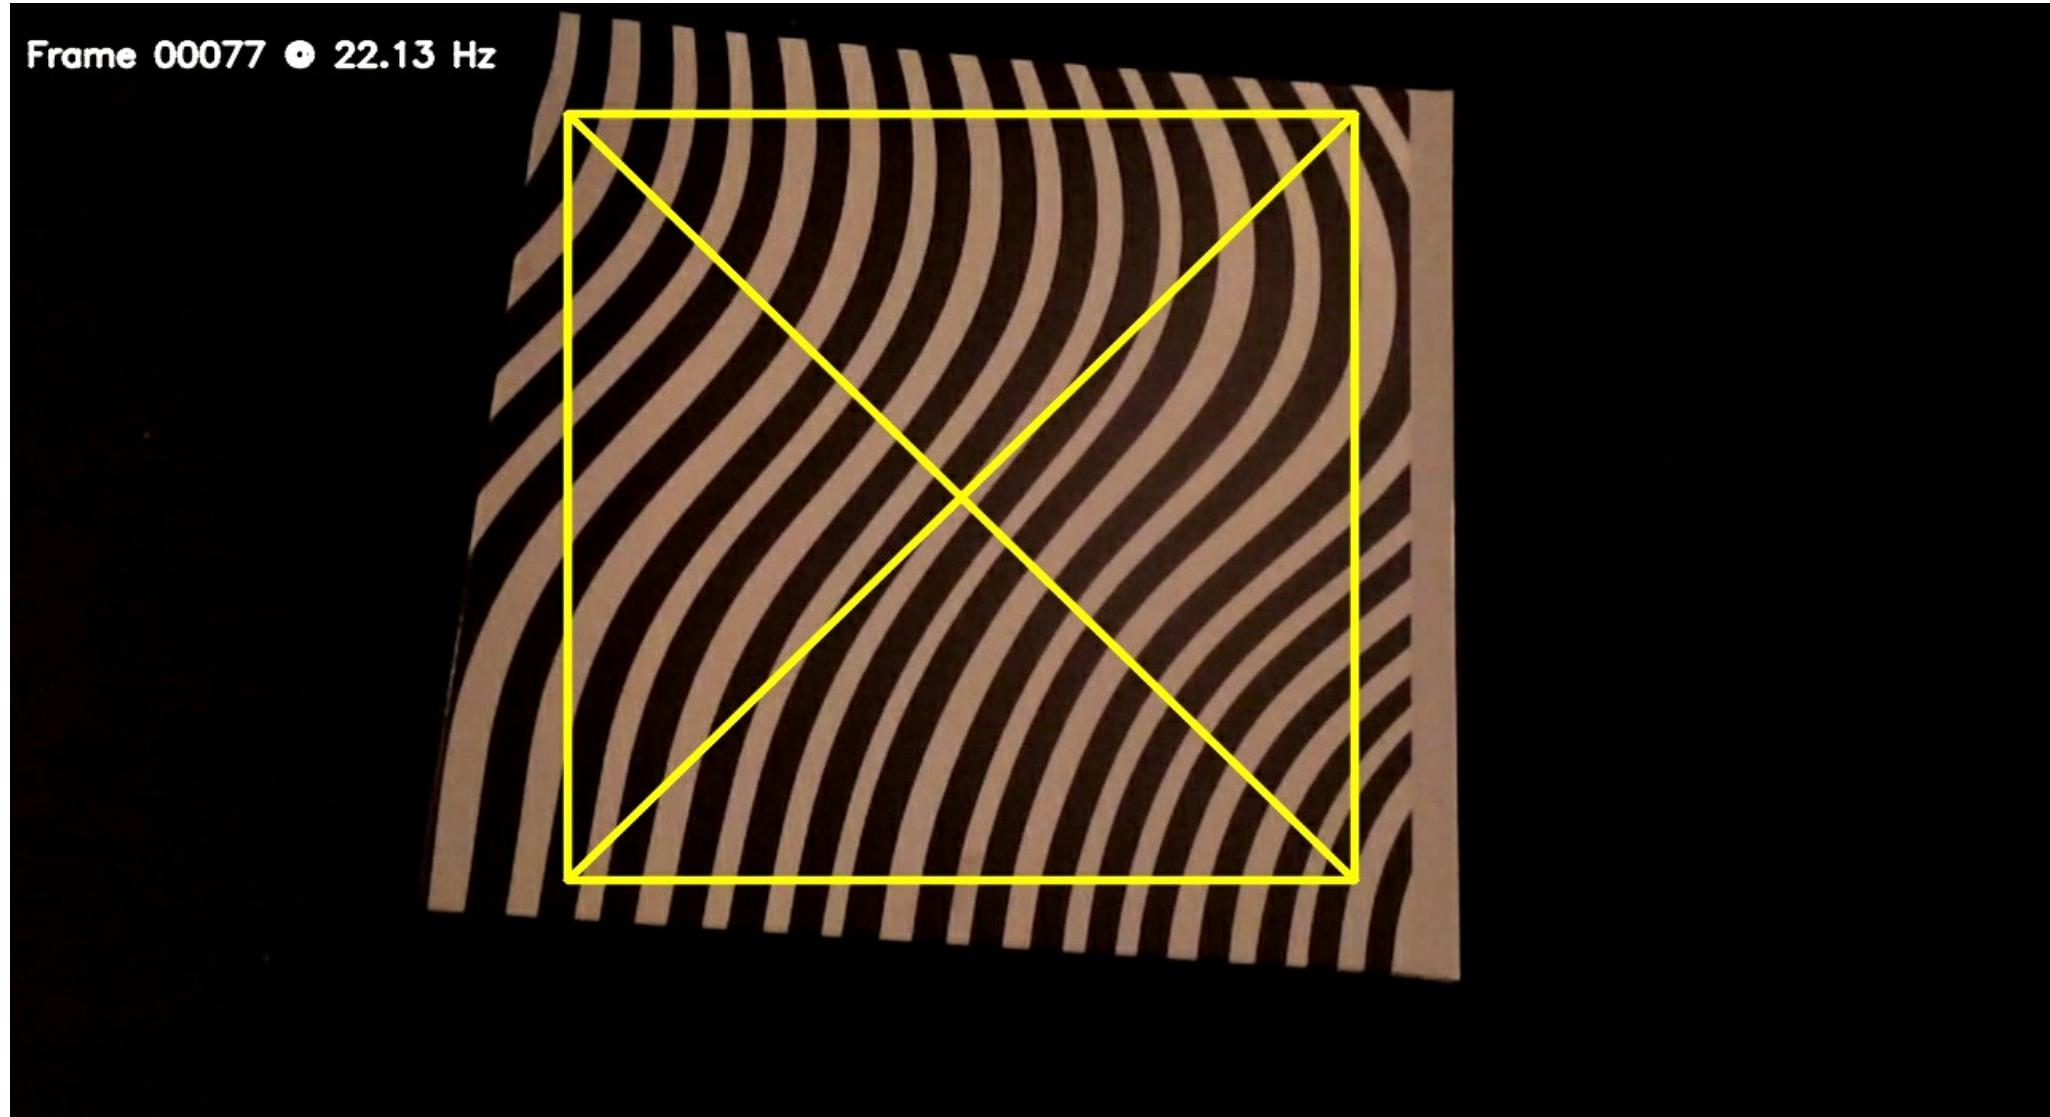


# Limitations of Keypoints

Frame 00077 • 28.07 Hz



# Limitations of Keypoints



# A step towards vision in challenging environments



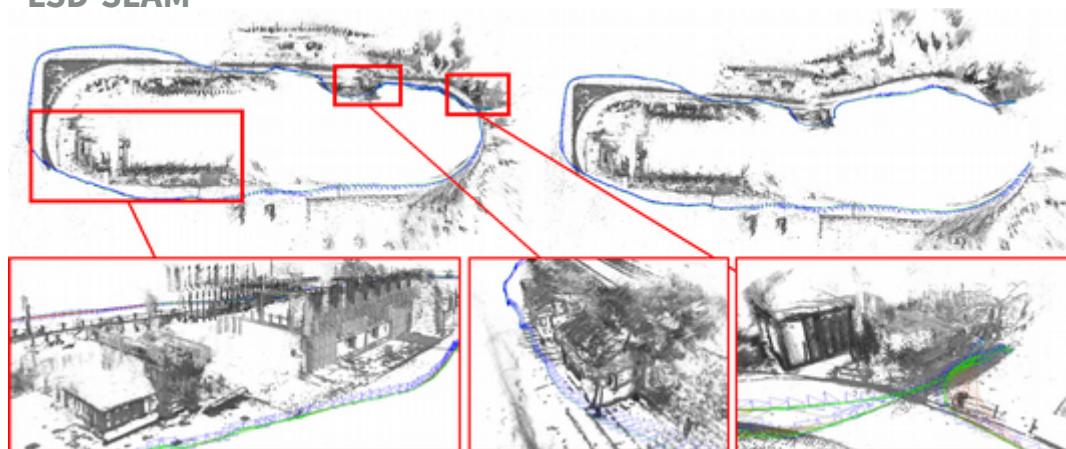
[http://hci.iwr.uni-heidelberg.de/benchmarks/Challenging\\_Data\\_for\\_Stereo\\_and\\_Optical\\_Flow](http://hci.iwr.uni-heidelberg.de/benchmarks/Challenging_Data_for_Stereo_and_Optical_Flow)

# Direct Methods

Increasing availability of high frame-rate data



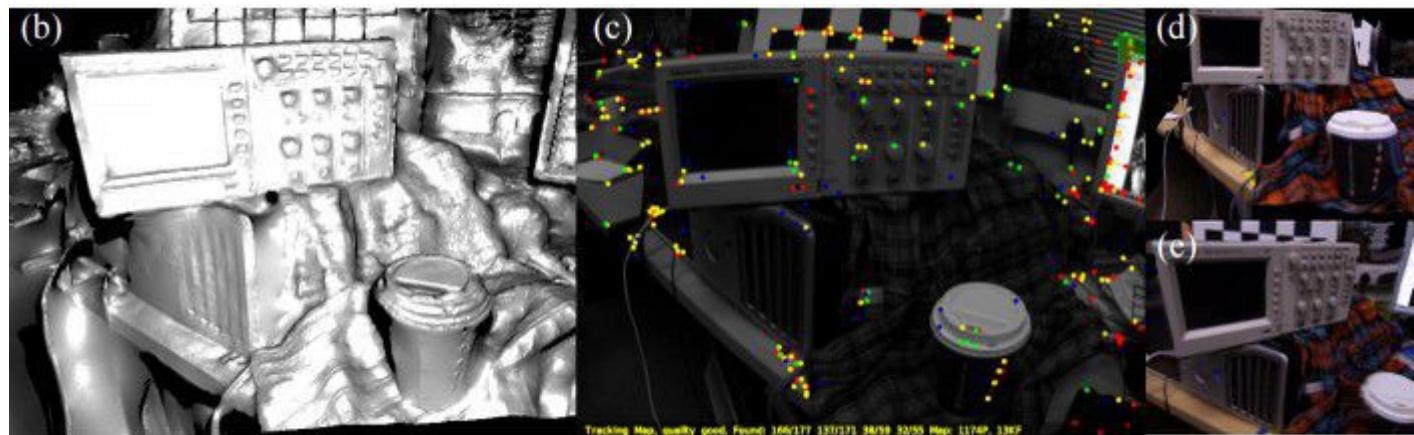
LSD-SLAM



RGBD-SLAM

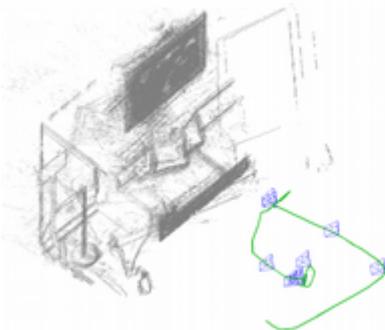


DTAM

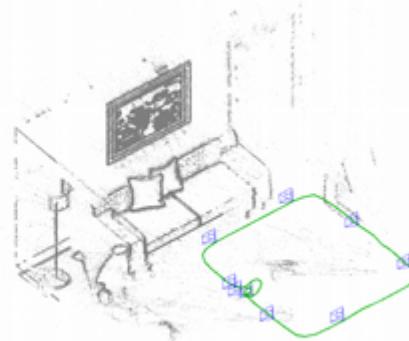


# Increased Interest in Direct Methods

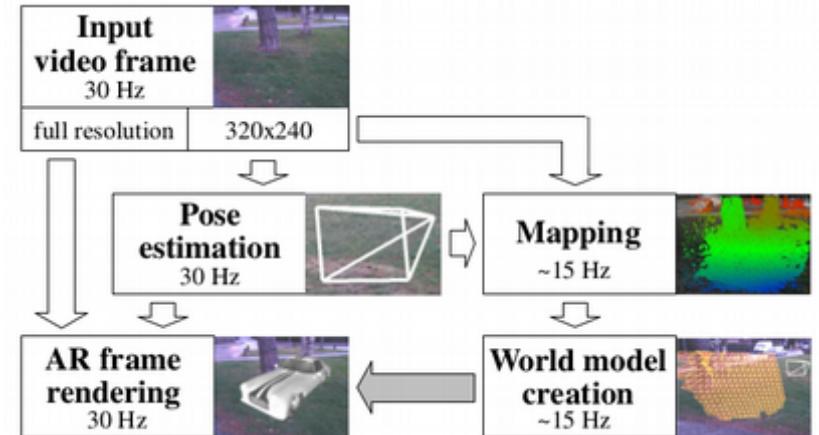
Rolling shutter



Kim, Cadena, Reid 2016

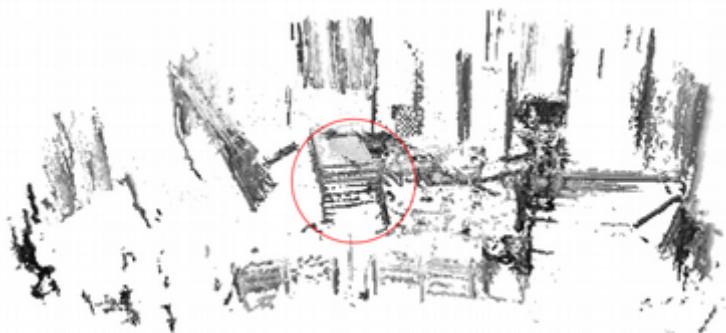


Mobile phones



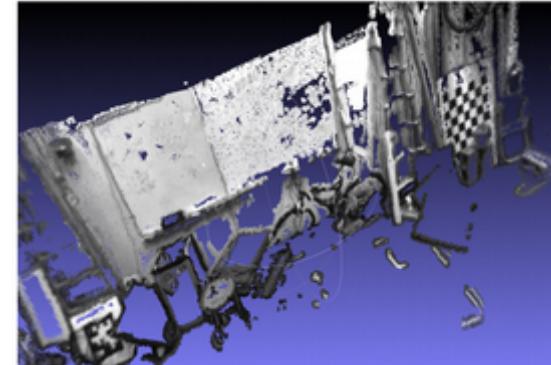
Schoppe, Engel, Cremers, 2014

Visual-Inertial stereo SLAM



Usenko, Engel, Stuckler, Cremers, 2016

Visual-Inertial Mono SLAM



Concha, Loianno, Kumar, Civera, 2016

# Direct Pose Estimation

The Lucas & Kanade Algorithm



Gradient-based search using  
image data directly to  
minimize a photometric error

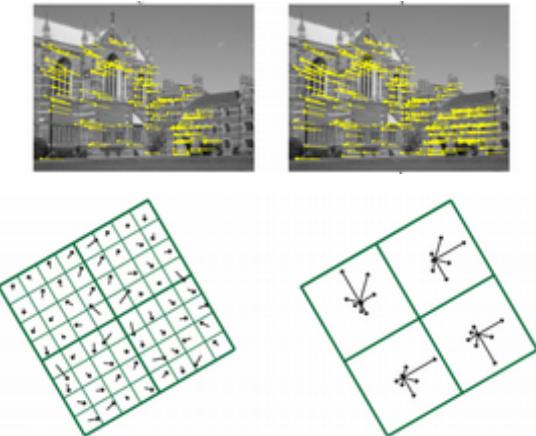
No keypoints necessary

More robust because we can  
use most of the image

# The Correspondence Problem

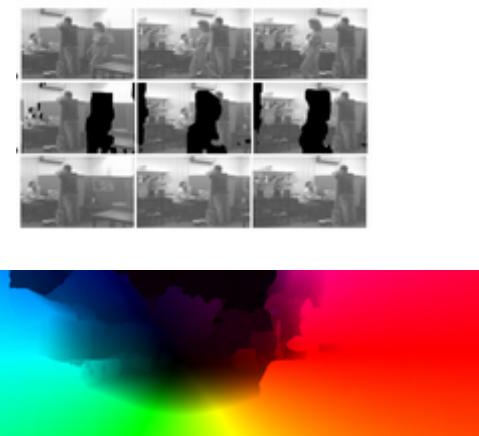
## Feature-based

Use pre-computed corrs.



## Direct

Estimate corrs. with pose



## Feature Based Methods for Structure and Motion Estimation

P. H. S. Torr<sup>1</sup> and A. Zisserman<sup>2</sup>

<sup>1</sup> Microsoft Research Ltd, 1 Guildhall St  
Cambridge CB2 3NH, UK  
[philtorr@microsoft.com](mailto:philtorr@microsoft.com)

<sup>2</sup> Department of Engineering Science, University of Oxford  
Oxford, OX1 3PJ, UK  
[az@robots.ox.ac.uk](mailto:az@robots.ox.ac.uk)

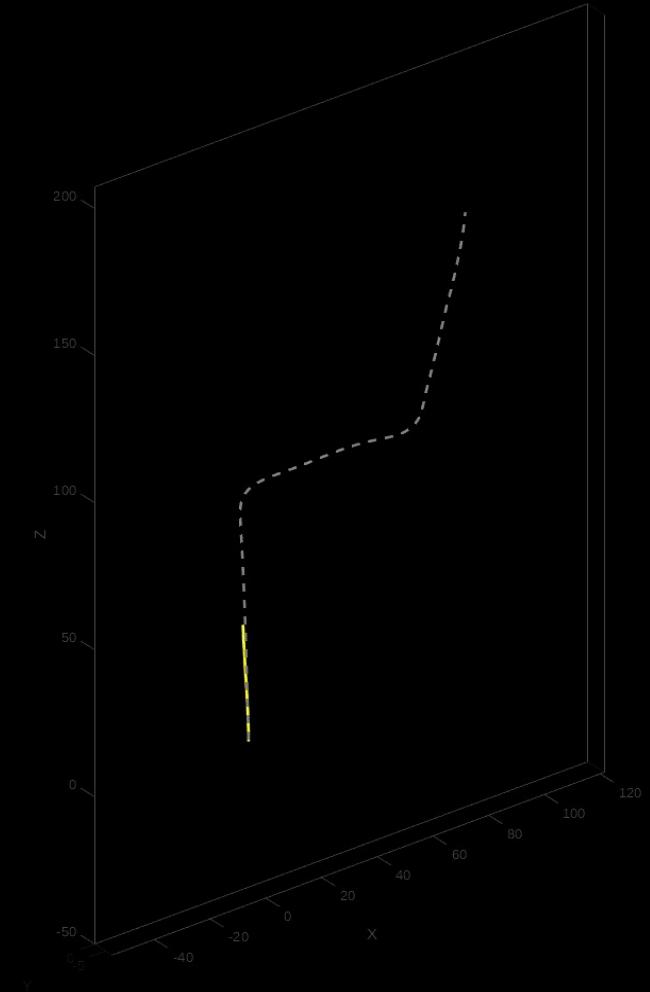
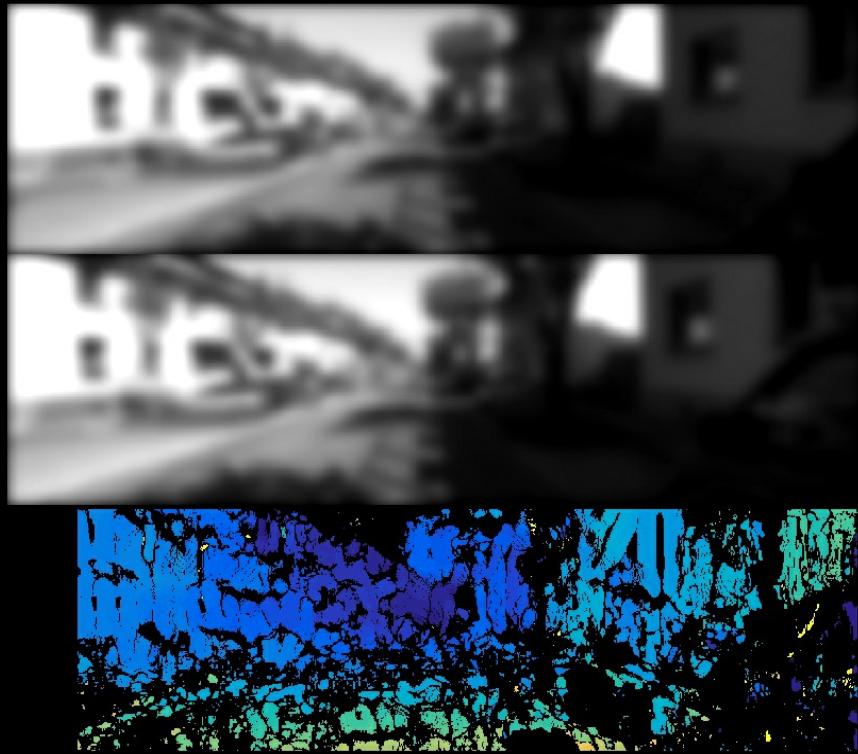
## All About Direct Methods

M. Irani<sup>1</sup> and P. Anandan<sup>2</sup>

<sup>1</sup> Dept. of Computer Science and Applied Mathematics,  
The Weizmann Inst. of Science, Rehovot, Israel.  
[irani@wisdom.weizmann.ac.il](mailto:irani@wisdom.weizmann.ac.il)

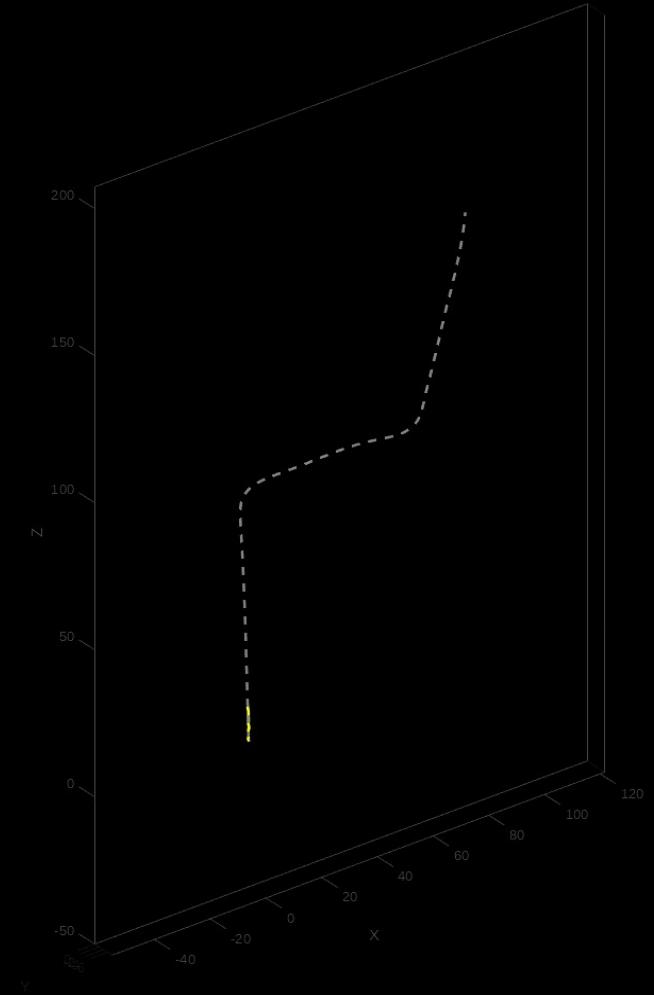
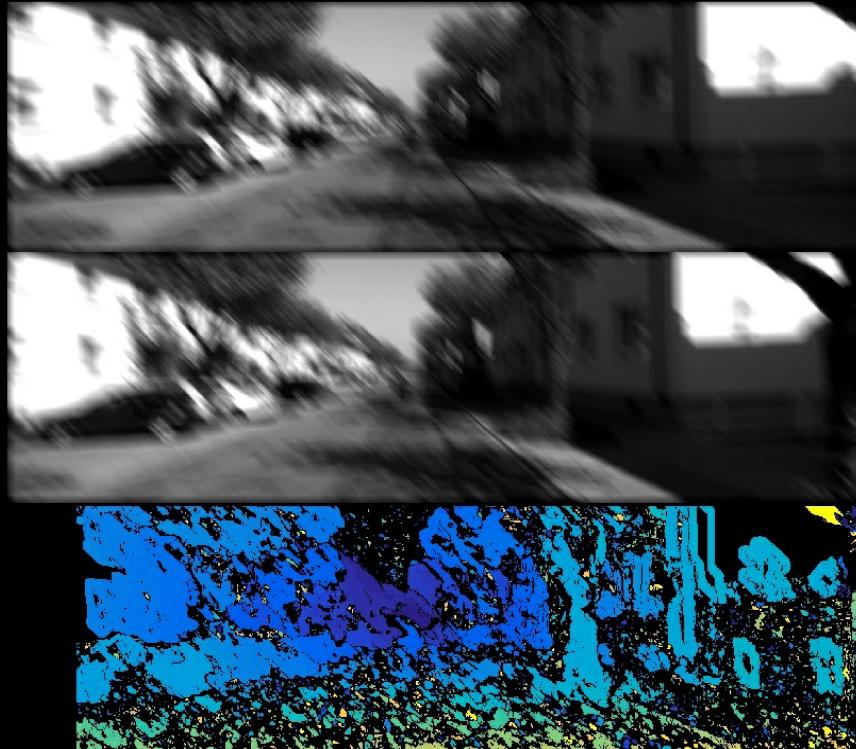
<sup>2</sup> Microsoft Research, One Microsoft Way,  
Redmond, WA 98052, USA.  
[anandan@microsoft.com](mailto:anandan@microsoft.com)

# Robustness of Direct Pose Estimation



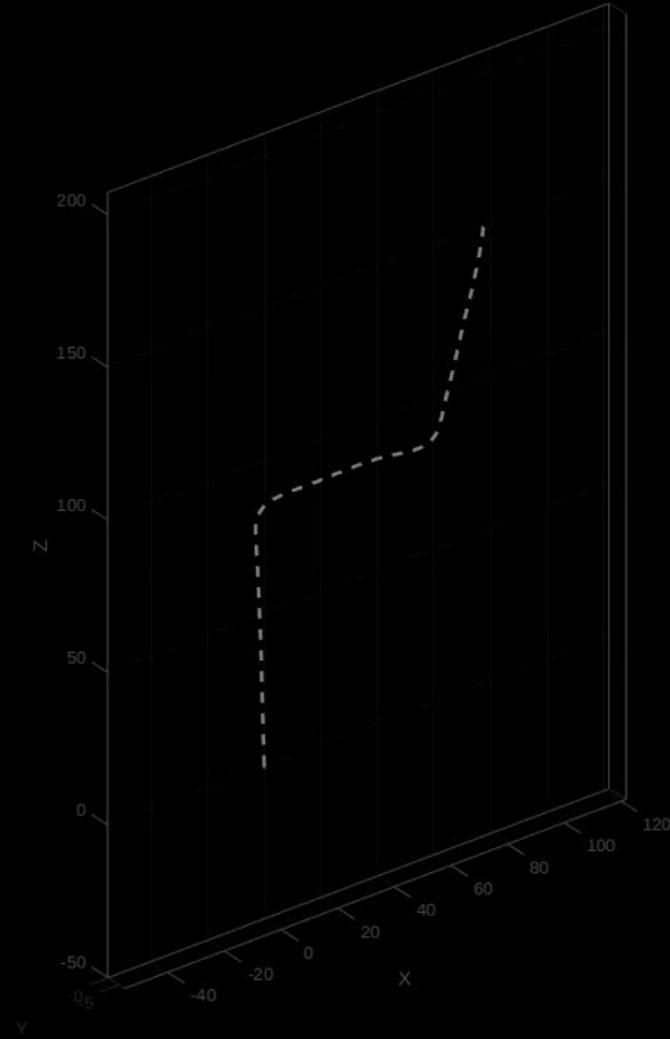
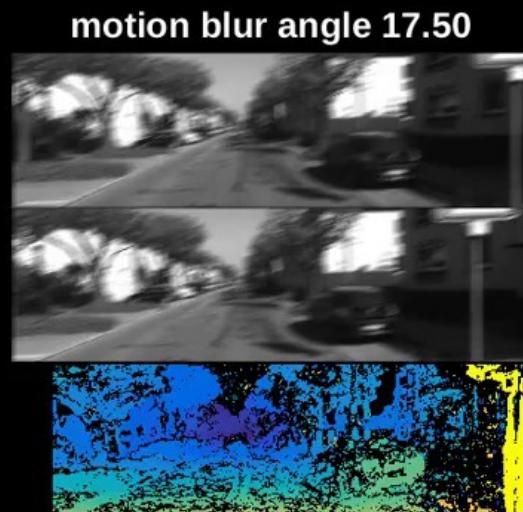
Using the Direct Disparity Space (DDS) algorithm in Chapter 3

# Robustness of Direct Pose Estimation



Using the Direct Disparity Space (DDS) algorithm in Chapter 3

# Robustness of Direct Pose Estimation

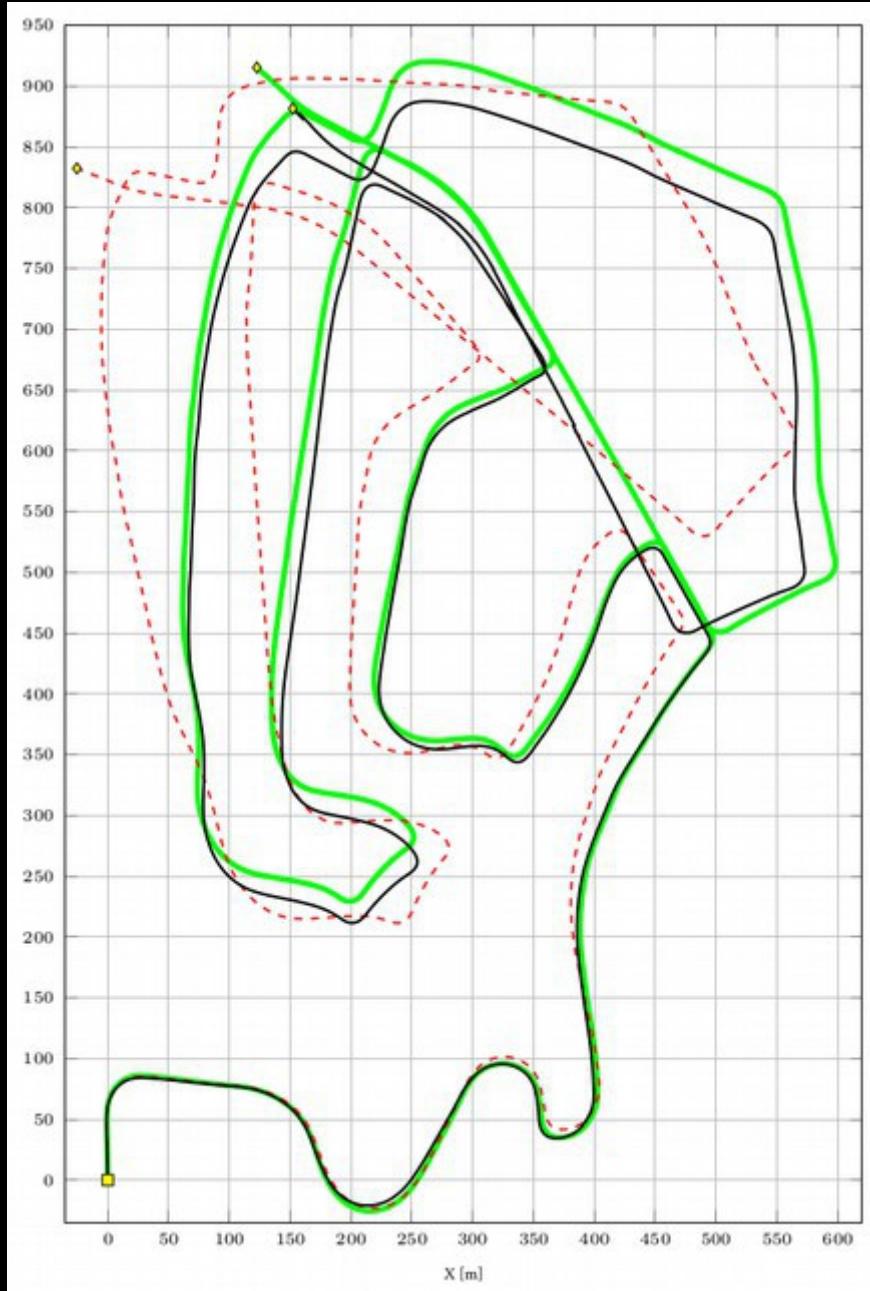


Using the Direct Disparity Space (DDS) algorithm in Chapter 3

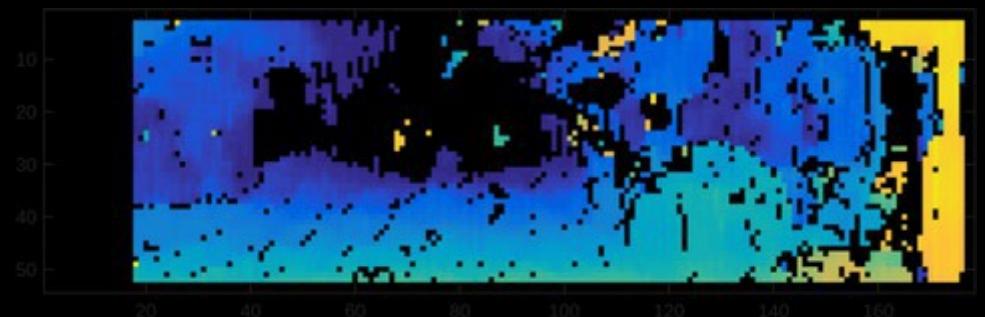
# Input to Our Algorithm



# KITTI with thumbnail stereo

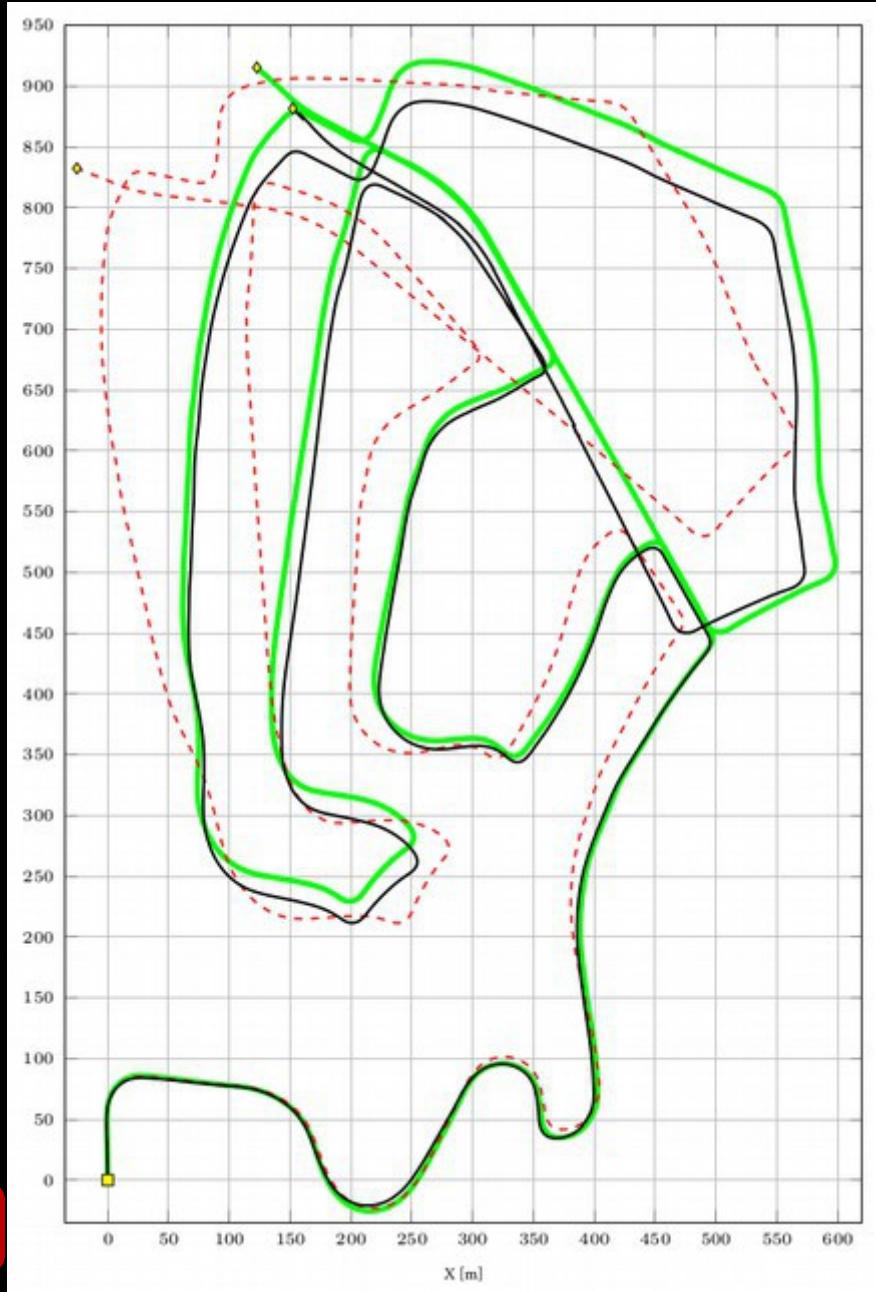


Green, DDS with 178x54 stereo images  
Red, VISO with 1241x376 resolution  
Black, GPS ground truth

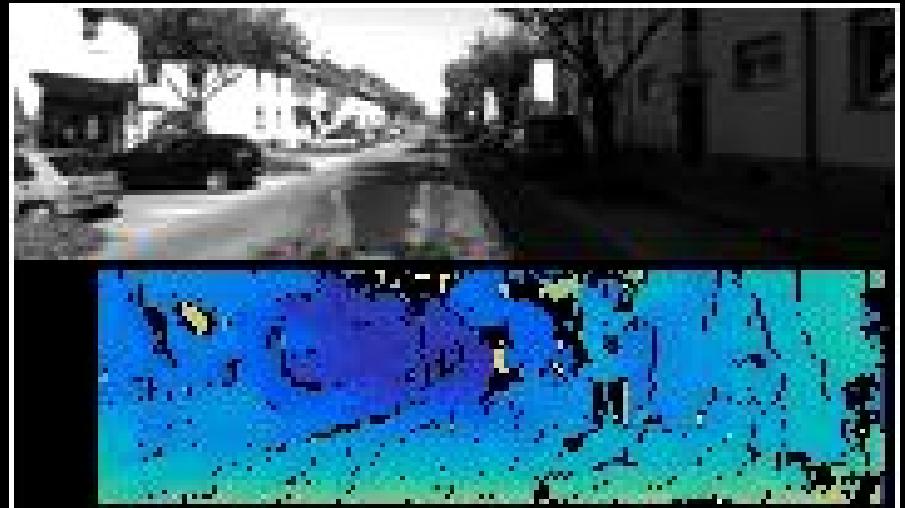


Stereo with opencv block matching  
5x5 SAD window and 16 disparities search range

# KITTI with thumbnail stereo

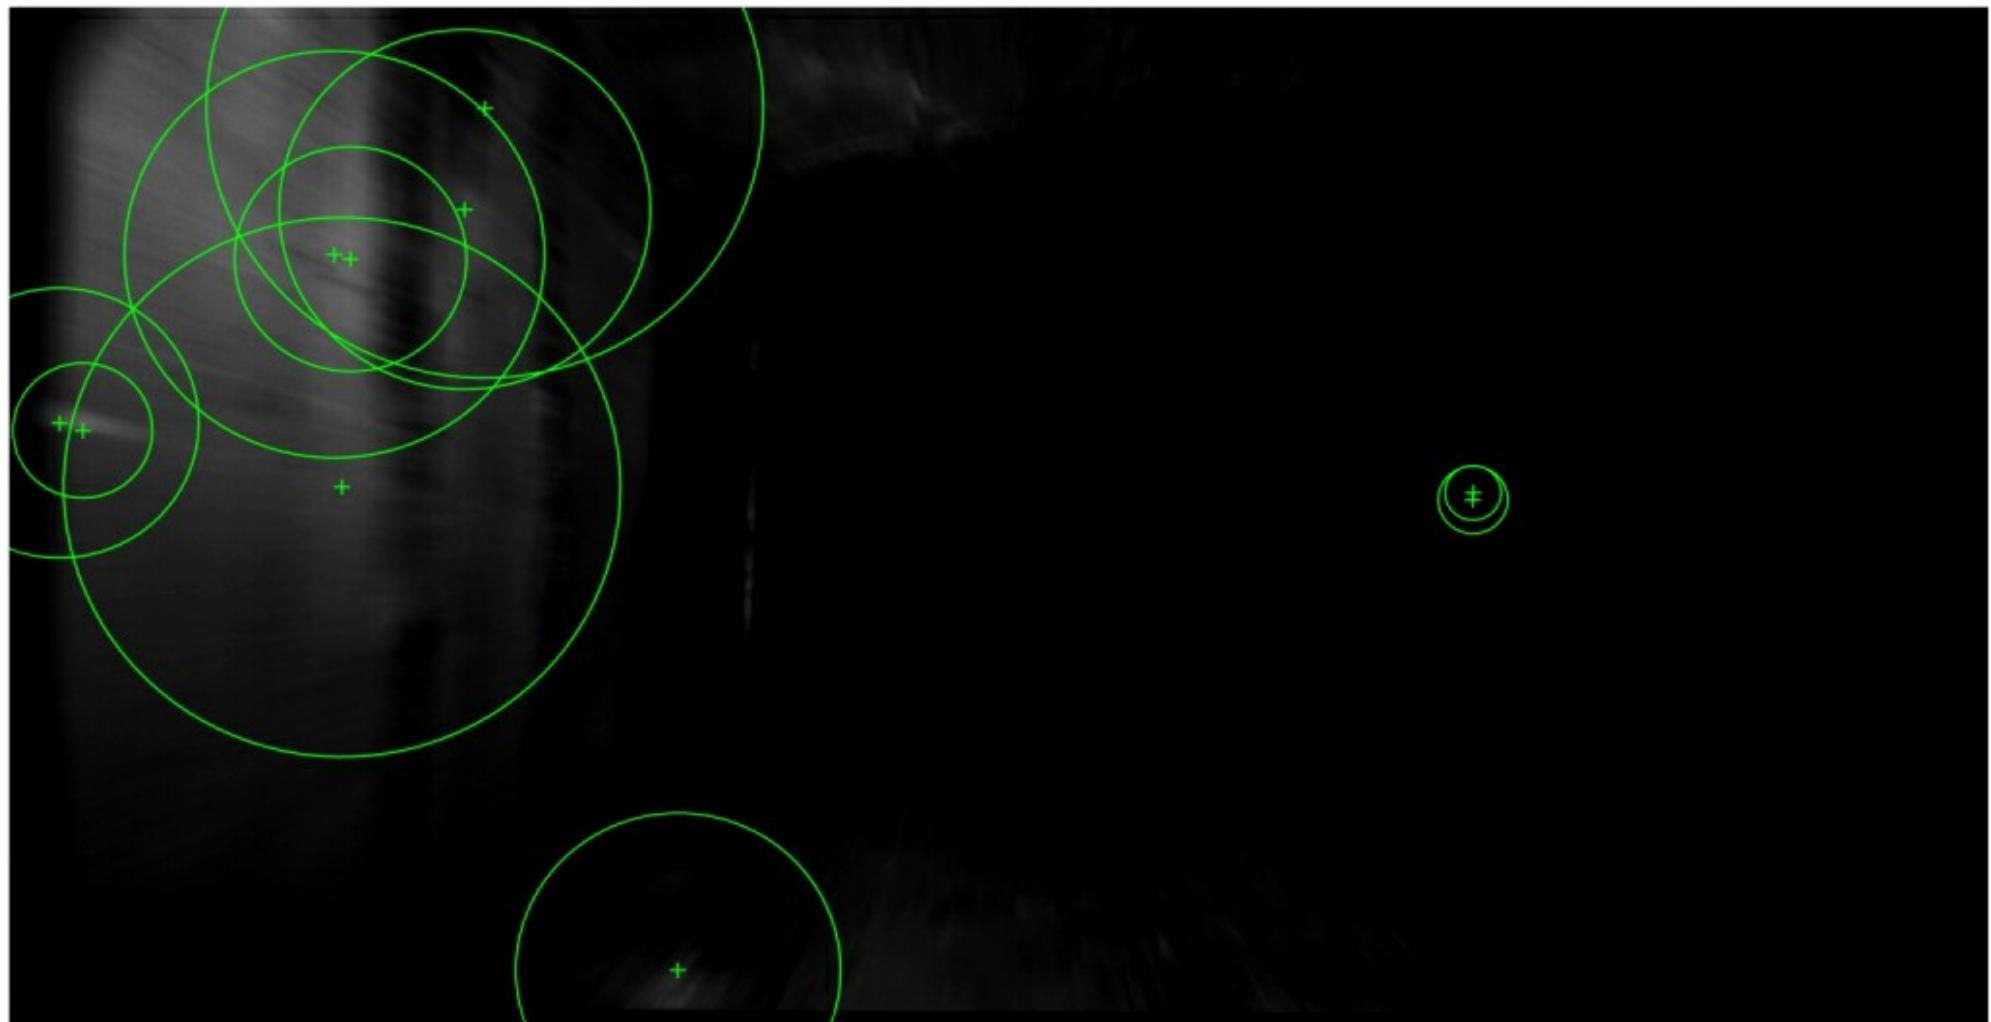


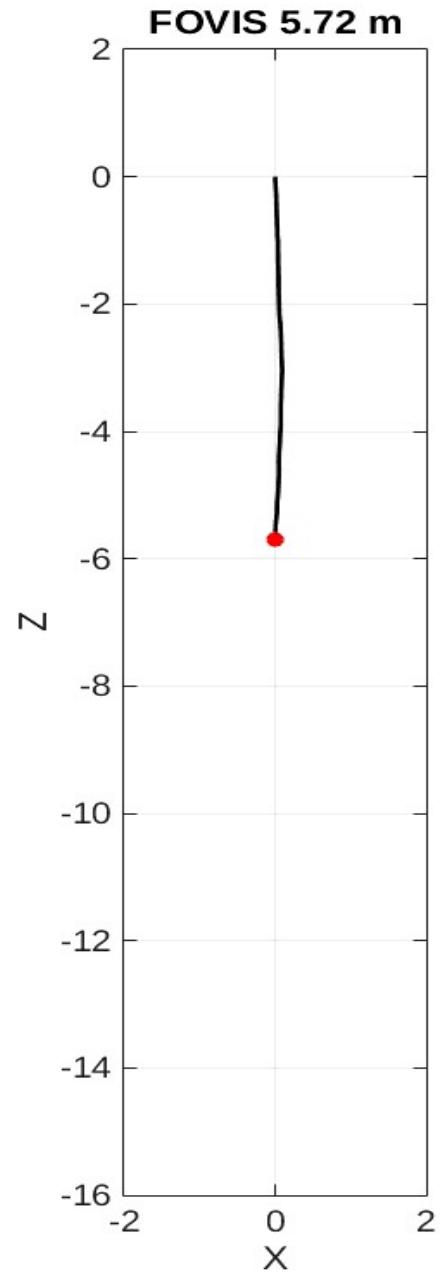
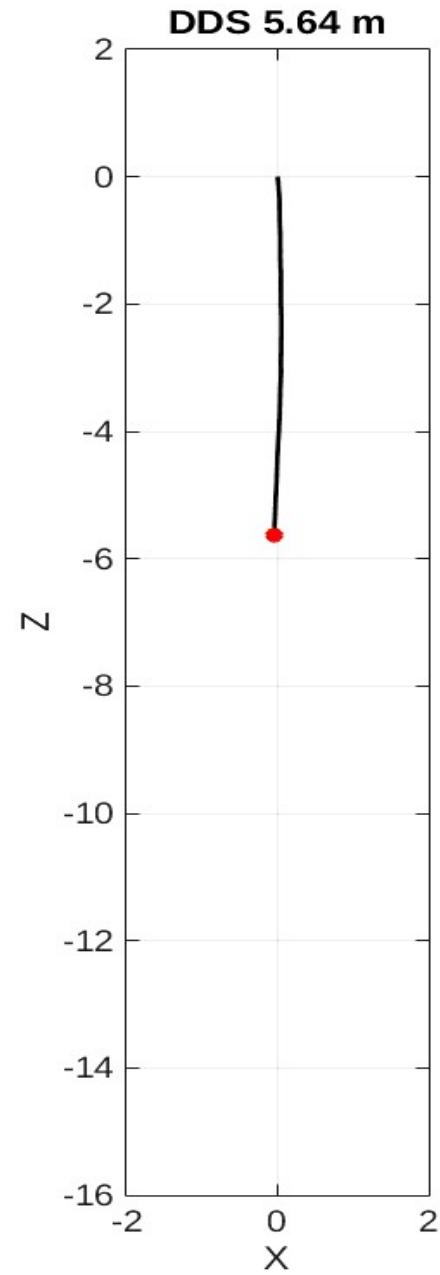
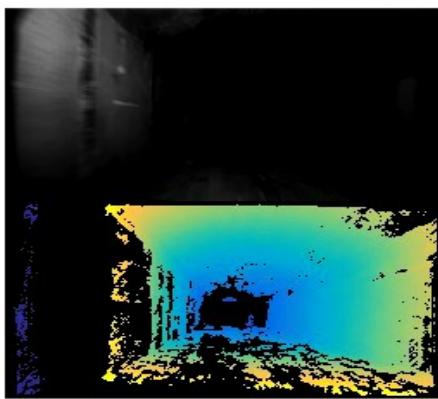
Green, DDS with 178x54 stereo images  
Red, VISO with 1241x376 resolution  
Black, GPS ground truth



Stereo with opencv block matching  
5x5 SAD window and 16 disparities search range

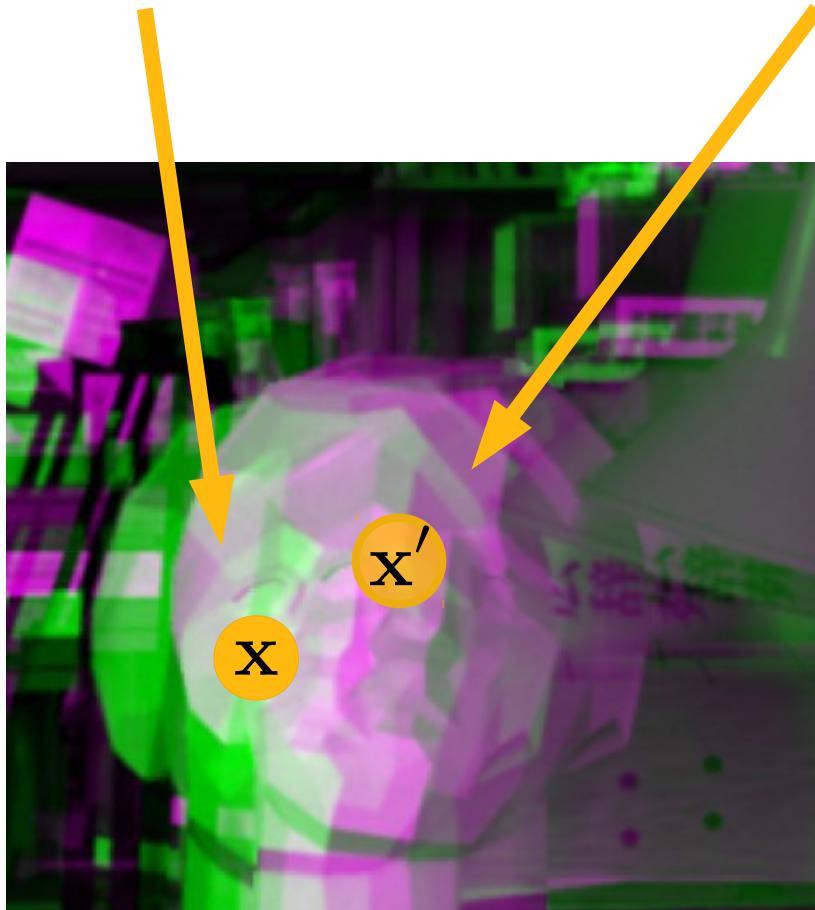
# Real Data from Underground Mines





# Limitations of Direct Methods

Pixel location in the reference image with intensity  $\mathbf{I}_0(\mathbf{x})$



There exists an (unknown) corresponding location in input image with the same intensity

$$\mathbf{I}_0(\mathbf{x}) = \mathbf{I}_1(\mathbf{x}')$$

## Brightness Constancy

The two views are related via a warp

$$\mathbf{x}' = \mathbf{w}(\mathbf{x}; \theta)$$

# Brightness Constancy

$$\Delta\theta^* = \operatorname{argmin}_{\Delta\theta} \sum_{\mathbf{x}} \|\mathbf{I}_0(\mathbf{x}) - \mathbf{I}_1(\mathbf{w}(\mathbf{x}; \theta + \Delta\theta))\|^2$$

Does not hold most of the time



Effective solution to  
the brightness  
constancy assumption

+

Direct Alignment  
*Lucas-Kanade*

---

Robust Pose estimation

**Effective** solution to  
the brightness  
constancy assumption

+

Direct Alignment  
*Lucas-Kanade*

Parameter-free

Invariant to arbitrary  
changes in illumination

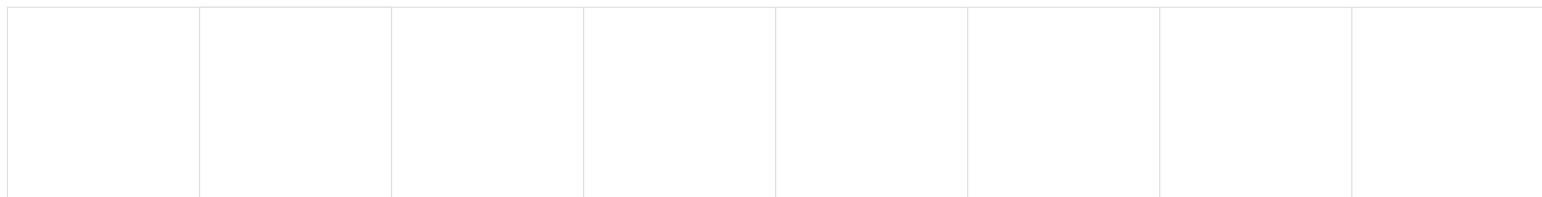
Efficient to compute

Robust Pose estimation

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

	>	



# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1		
	>	

1								
---	--	--	--	--	--	--	--	--

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1	1	
	>	

1	1								
---	---	--	--	--	--	--	--	--	--

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1	1	0
	>	

1	1	0						
---	---	---	--	--	--	--	--	--

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1	1	0
	>	0

1	1	0	0					
---	---	---	---	--	--	--	--	--

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1	1	0
	>	0
		1

1	1	0	0	1				
---	---	---	---	---	--	--	--	--

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1	1	0
	>	0
	1	1

1	1	0	0	1	1		
---	---	---	---	---	---	--	--

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1	1	0
	>	0
0	1	1

1	1	0	0	1	1	0	
---	---	---	---	---	---	---	--

# Binary Feature Descriptors

8	12	200
56	42	55
128	16	11

1	1	0
0	>	0
0	1	1

1	1	0	0	1	1	0	0
---	---	---	---	---	---	---	---

# Photometric Invariance

8	12	200
56	42	55
128	16	11

40	12	200
56	42	55
128	16	11

{1,1,0,0,1,1,0,0}

{1,1,0,0,1,1,0,0}

If using SSD, residual is  
 $(40-8)^2 = 1024$

# Must be Matching with a Binary Norm

8	12	200
56	42	55
128	16	11

255	12	200
56	42	55
128	16	11

{**1**,1,0,0,1,1,0,0}

{**0**,1,0,0,1,1,0,0}

Hamming distance = 1  
Descriptors remain close

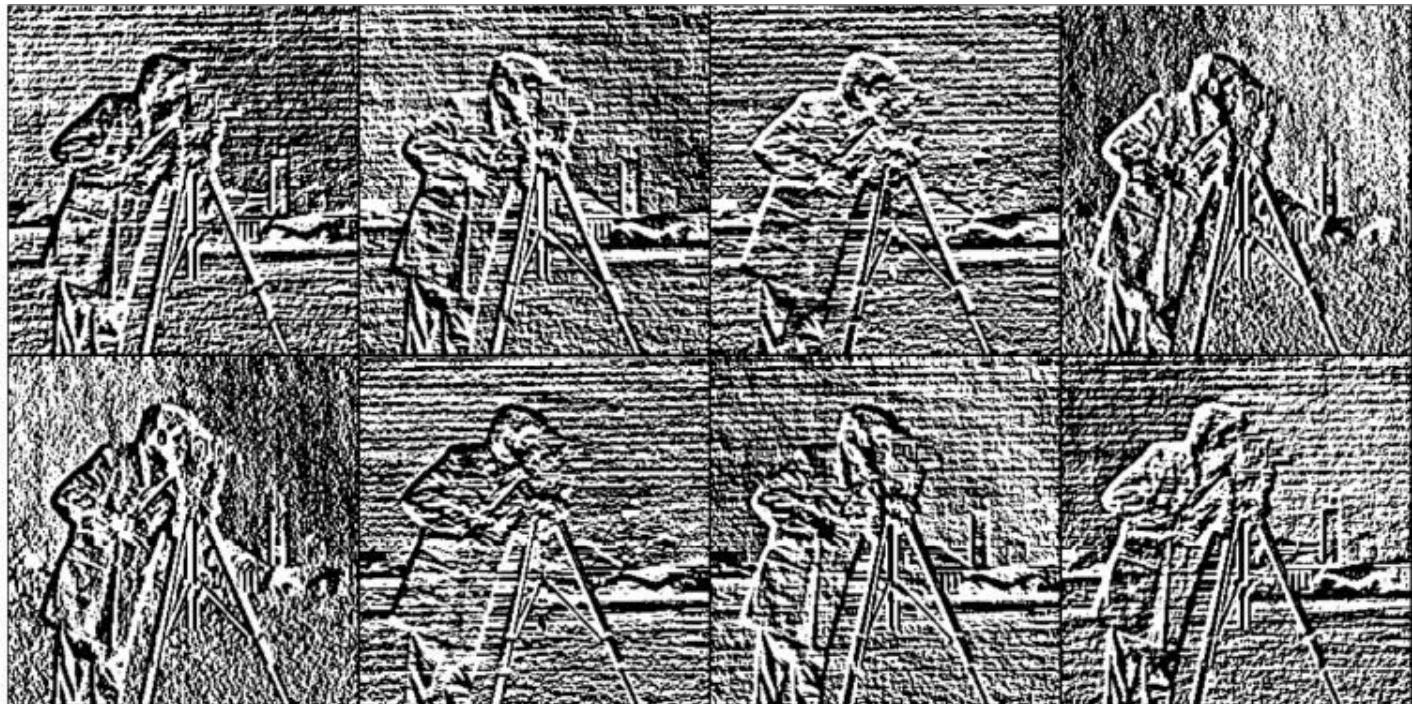
# Using Binary Descriptors in LK

$$\Delta\theta^* = \operatorname{argmin}_{\Delta\theta} \sum_{\mathbf{x}} \|\mathbf{I}_0(\mathbf{x}) - \mathbf{I}_1(\mathbf{w}(\mathbf{x}; \theta + \Delta\theta))\|^2$$

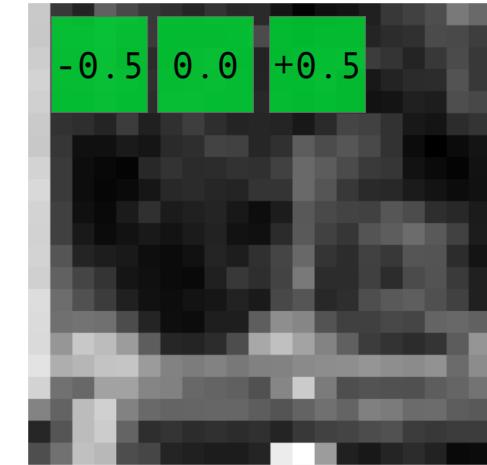
We can approximate the Hamming distance,  
but lose invariance.

# Bit-Planes Representation

Elegant Solution



Achieves two goals



Gradients with simple convolution

Squared norm becomes equivalent to Hamming

$$\begin{array}{r} [1, 1, 0, 0, 1, 0, 0] \\ \oplus \\ [1, 1, 0, 0, 0, 1, 0] \\ \hline [0, 0, 0, 0, 1, 1, 0] \end{array}$$

# Does it work?



**Intensity only**



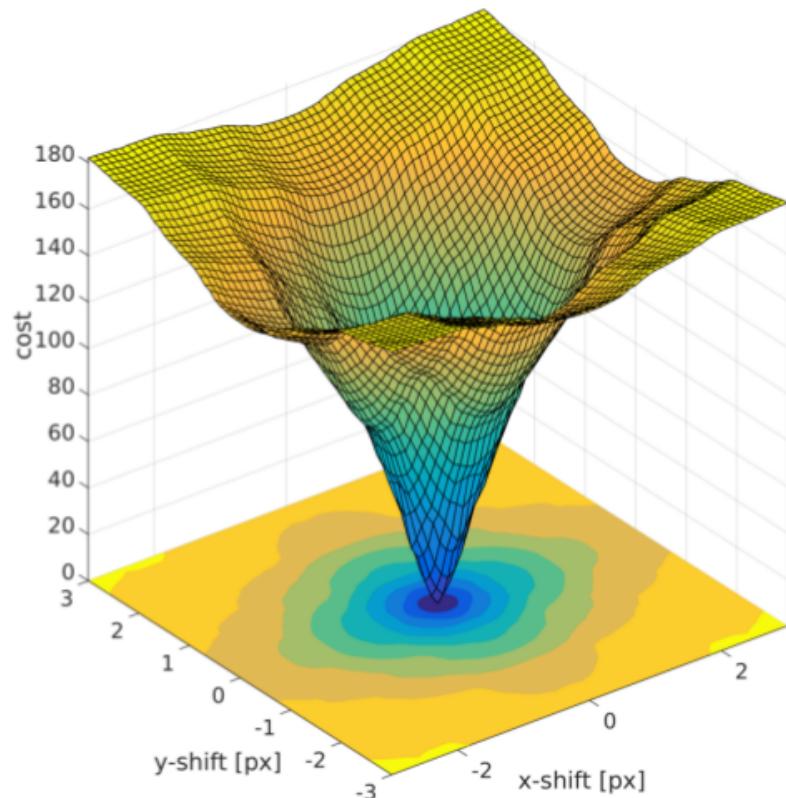
**BitPlanes**

8DOF homography tracking

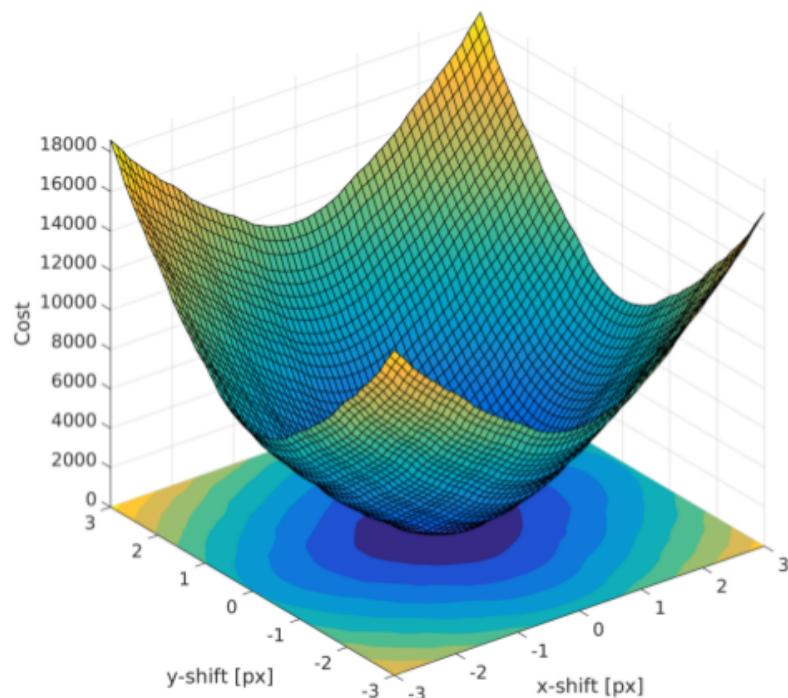


# Why does it work?

Cost surface is quasi-quadratic



(a) Bit-Planes.



(b) Raw intensity.

# Real time Tracking

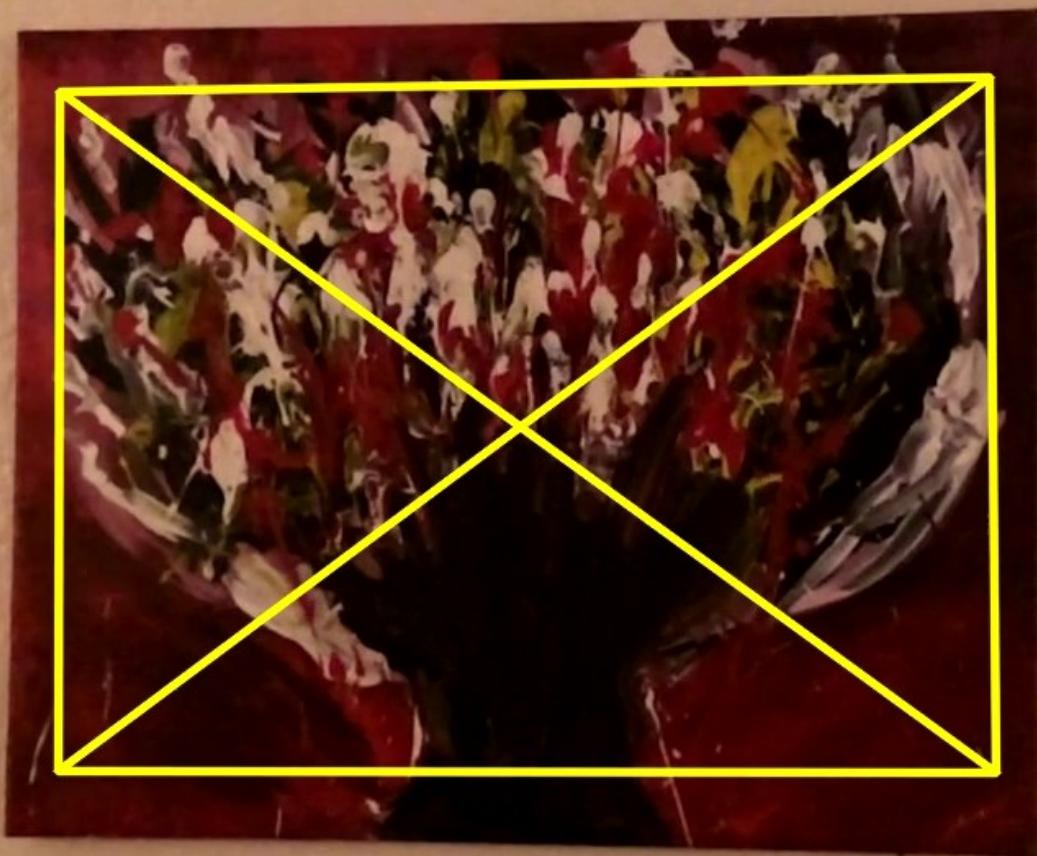
## BitPlanes

*Direct Tracking with Binary Descriptors*



# Real time Tracking

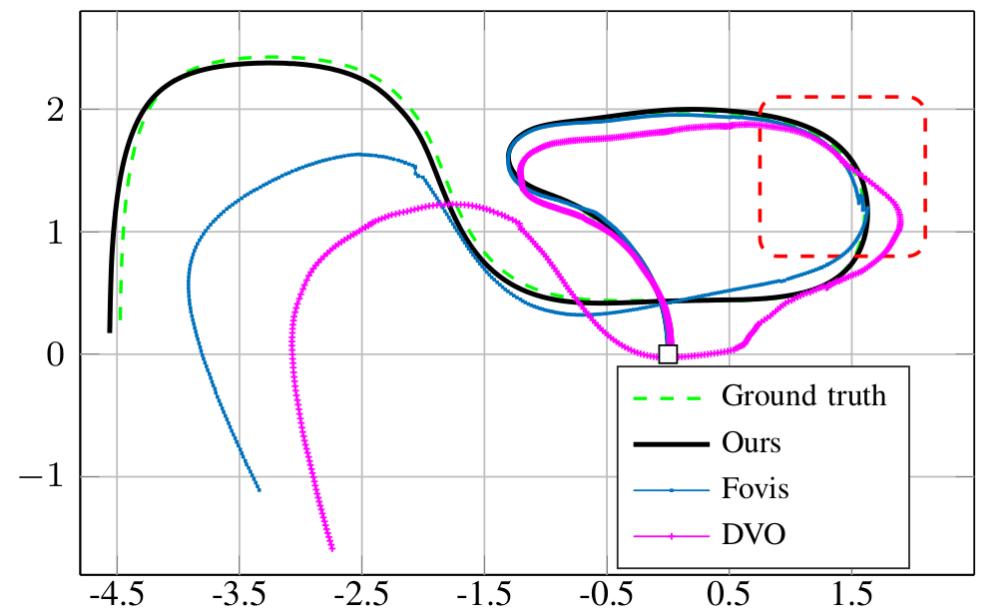
Frame 00226 • 122.23 Hz



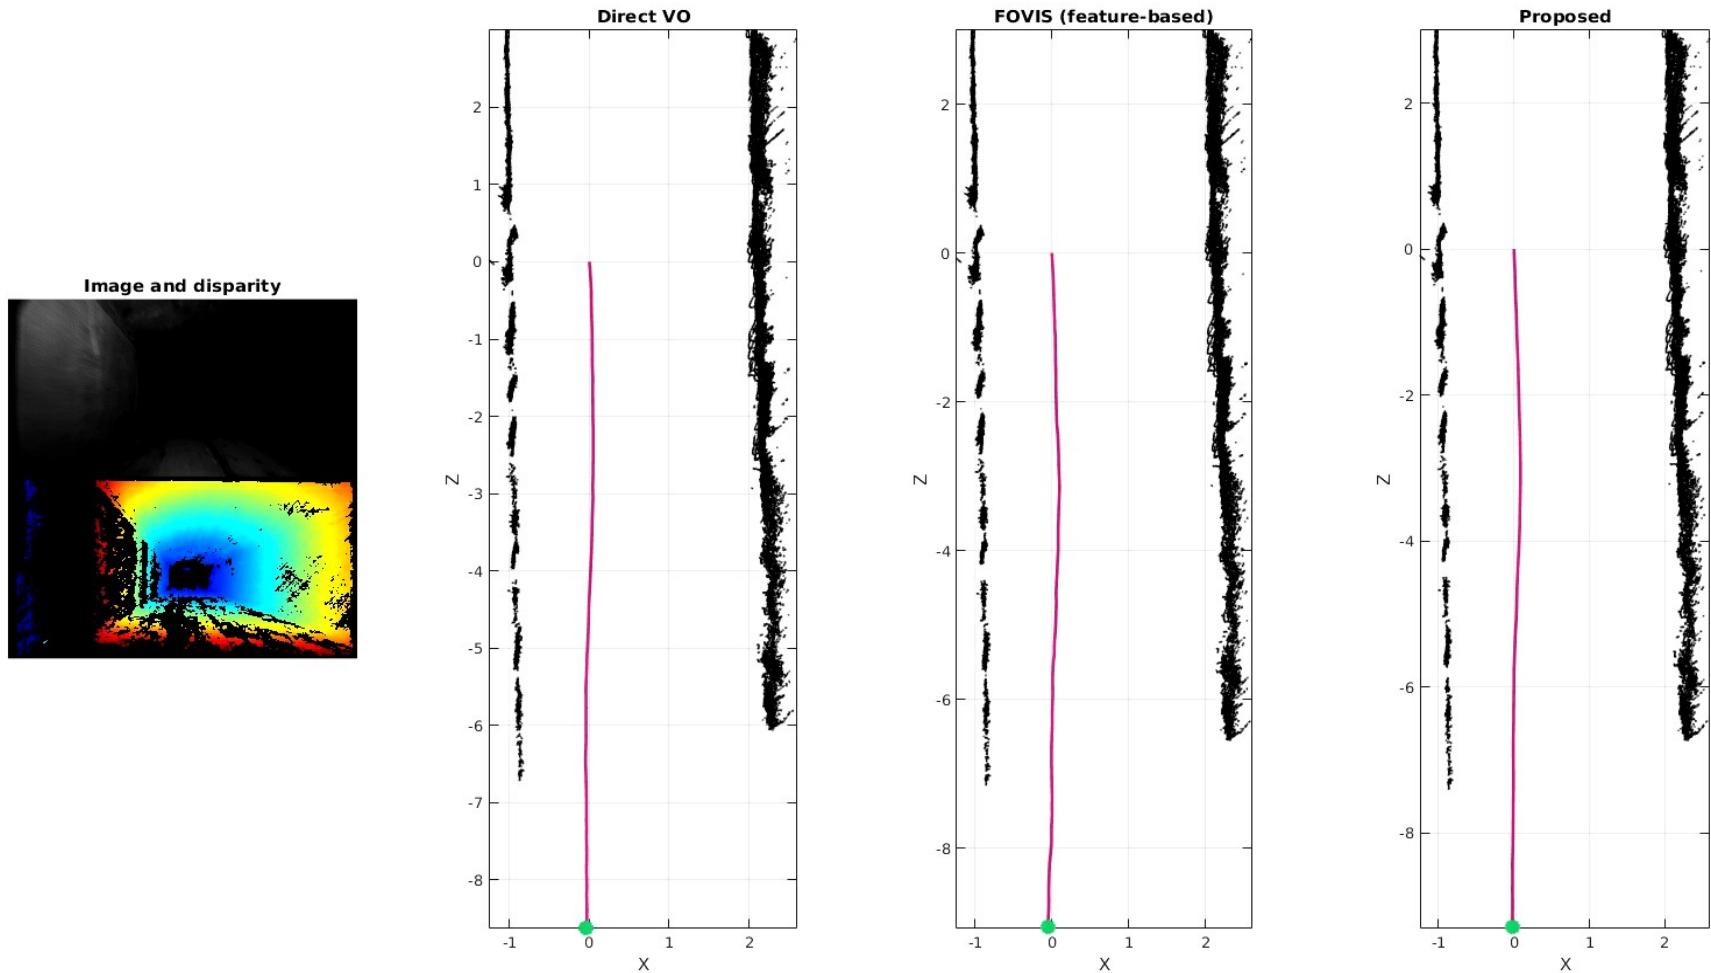
# Application to Robust Visual Odometry



<https://github.com/halismai/bpvo>

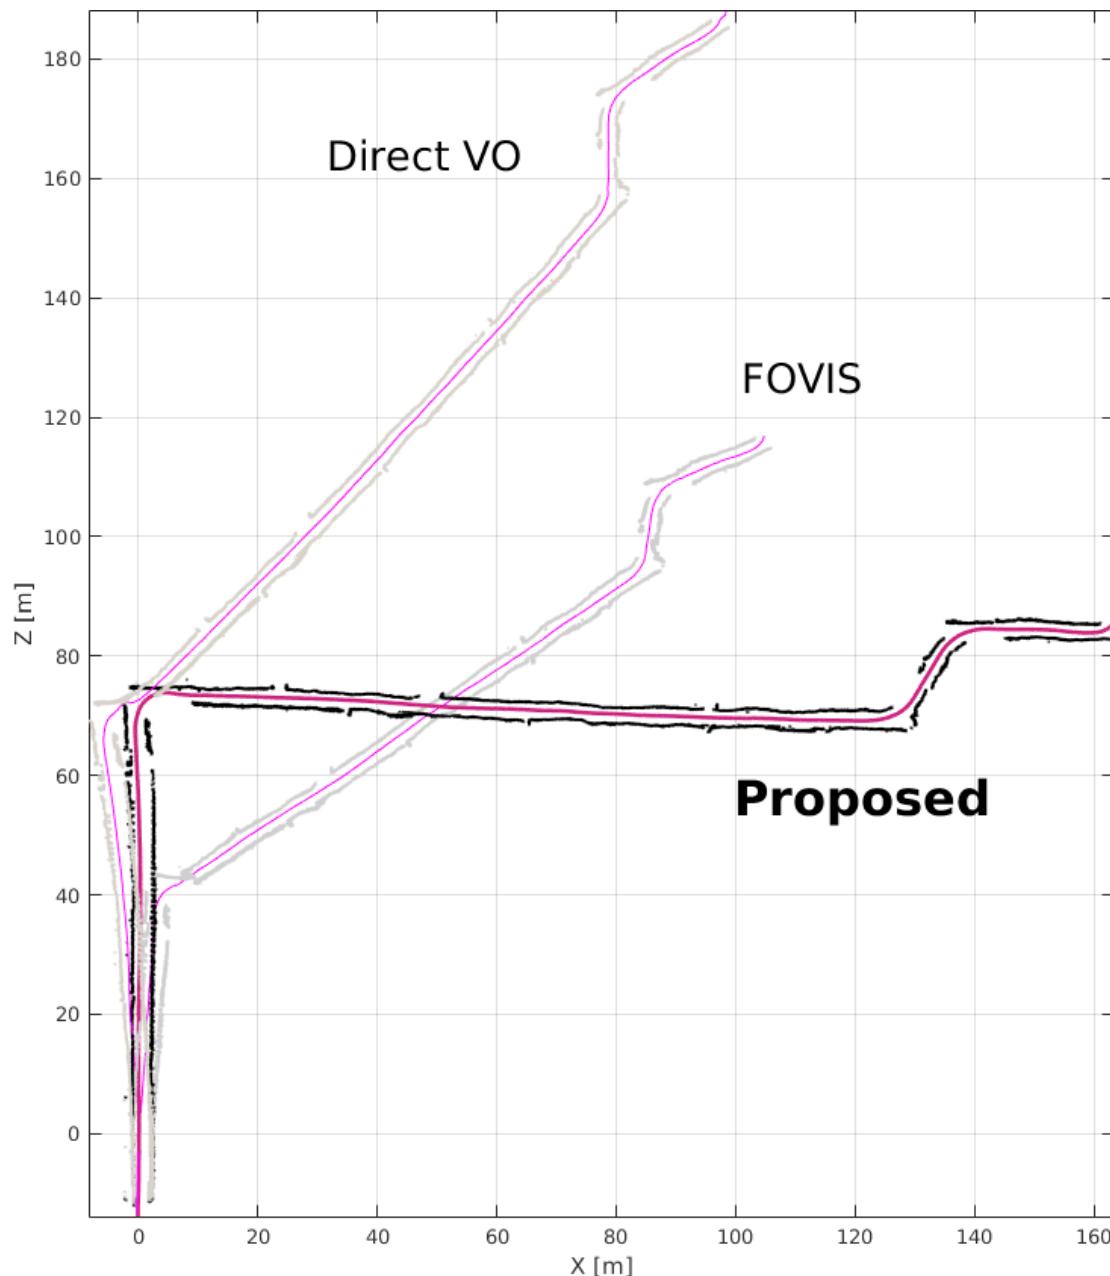


# Application to Underground Mines

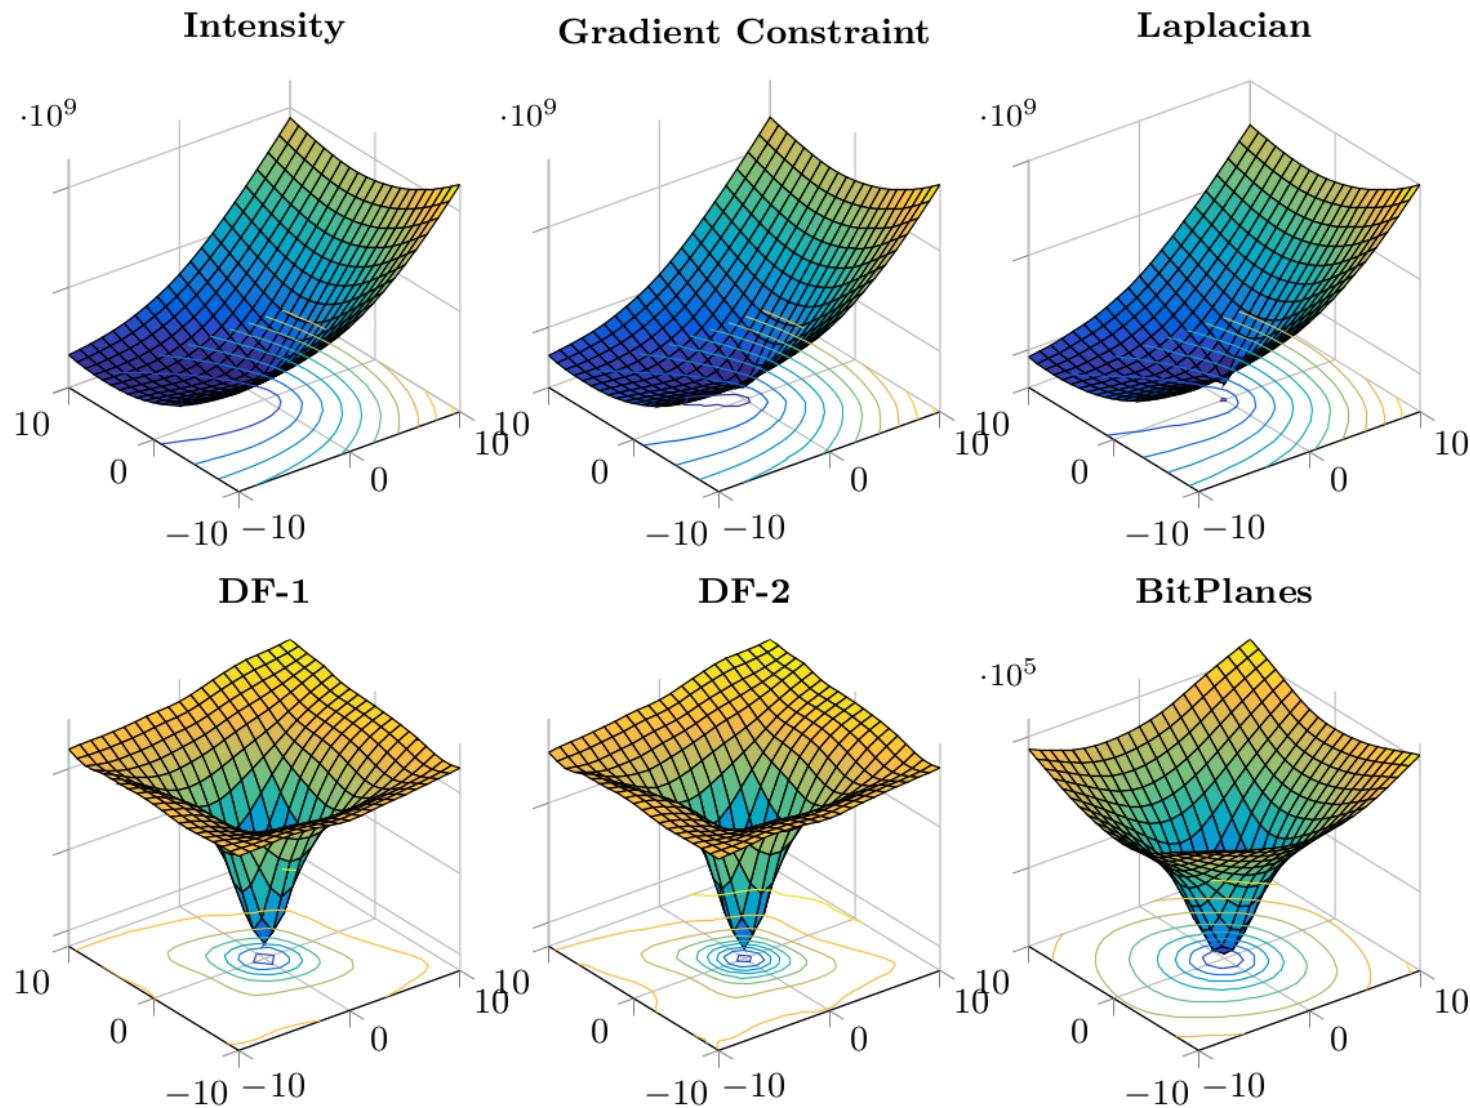


Sped up 2X

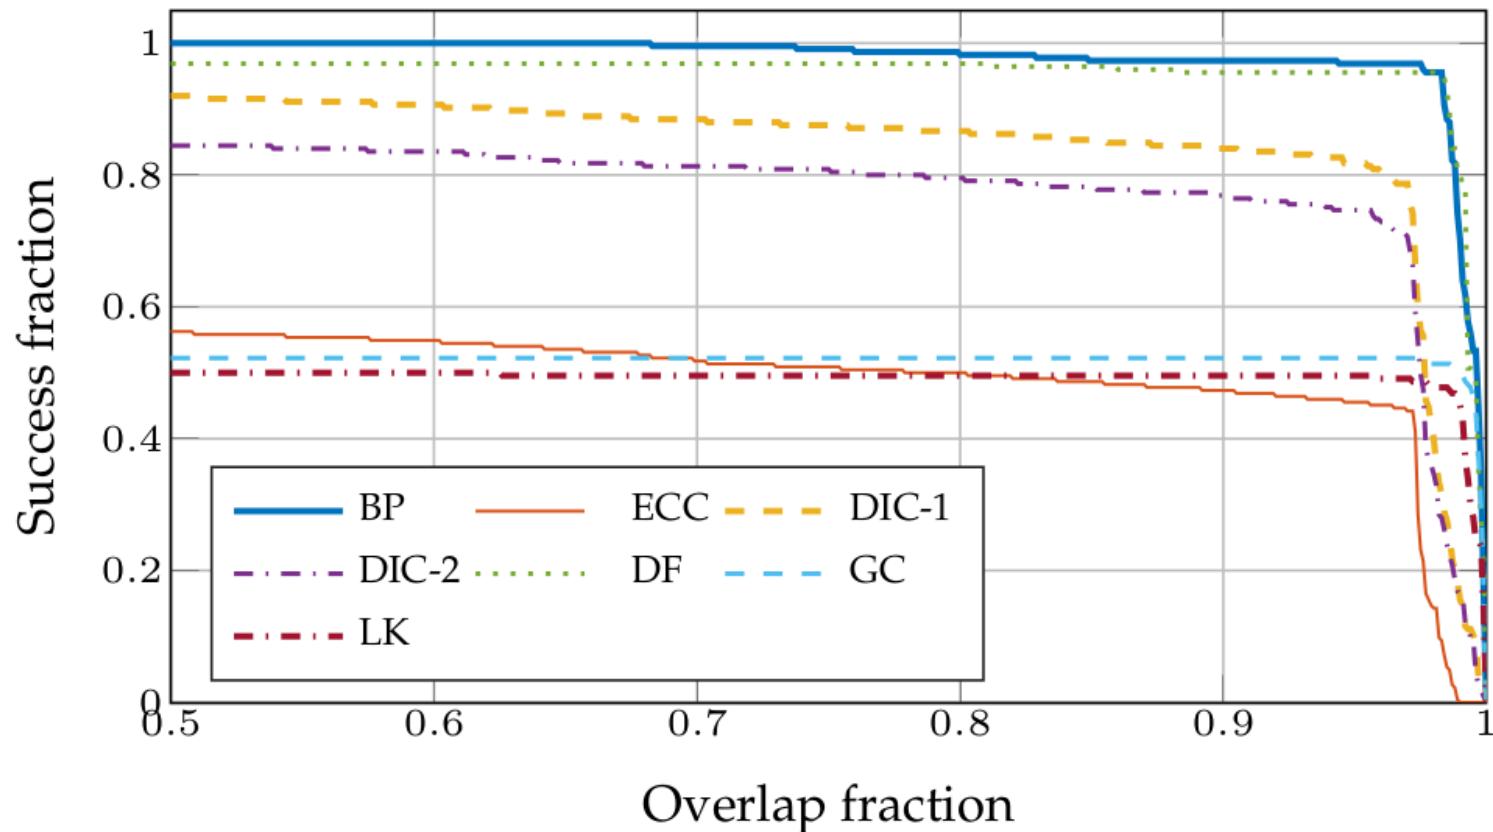
# Application to Underground Mines



# Comparison with other descriptors



# Quantitative Benchmark



Benchmark by:  
Gauglitz, et al 2011

Chapters 5 and 6 in dissertation include additional evaluation and quantitative results with applications to VO and parametric image alignment

# Comparison with other descriptors

**Bit-Planes**



**DF**

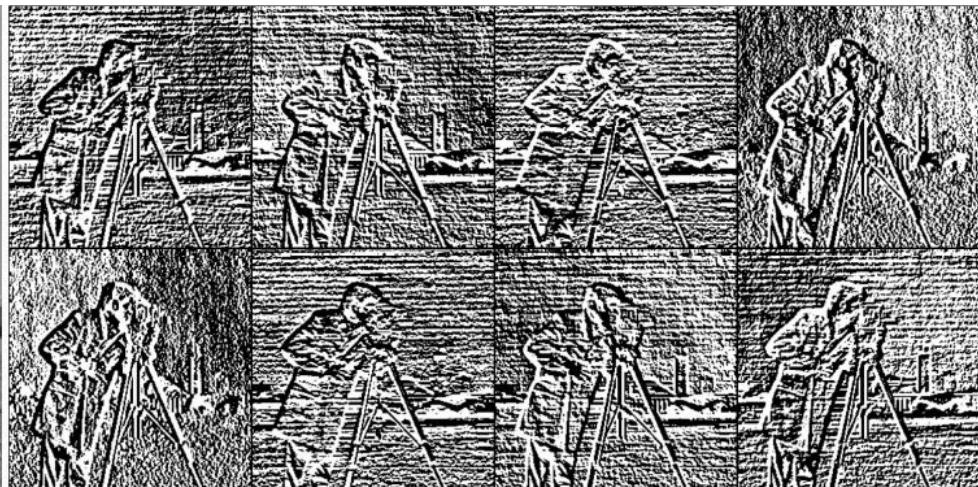
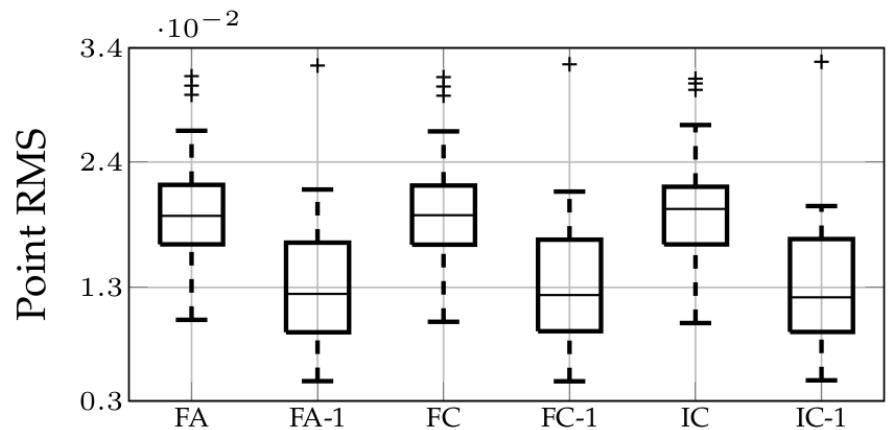


**GC**



# Additional Details

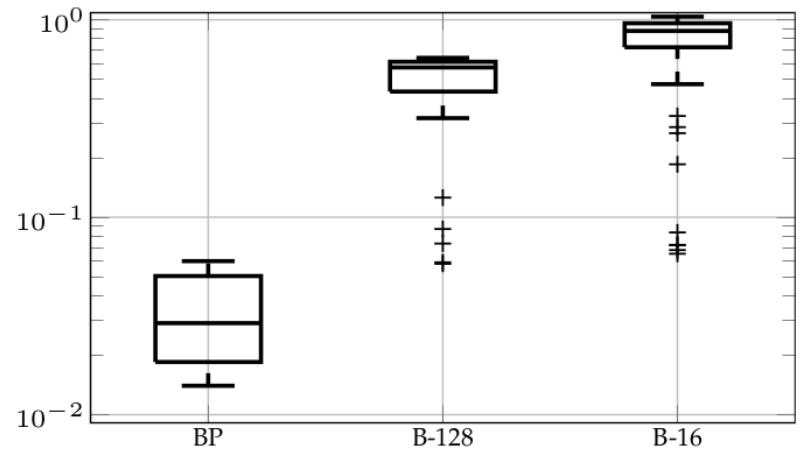
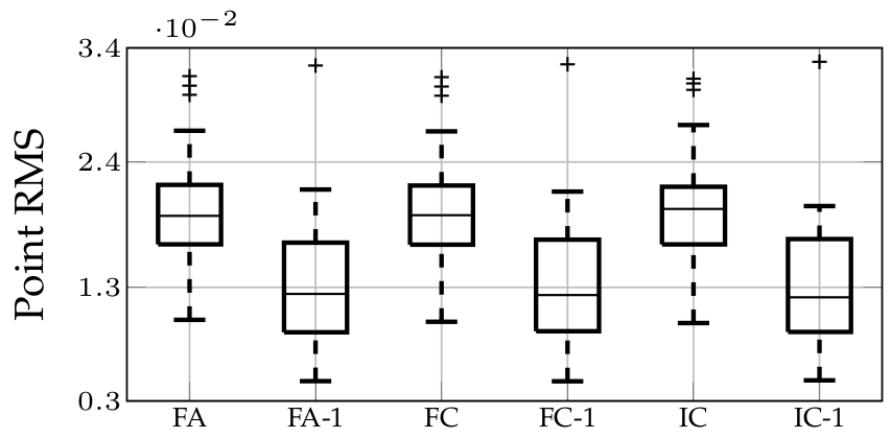
- Warp descriptor images, or warp image then recompute descriptors?



Chapters 5 and 6

# Additional Details

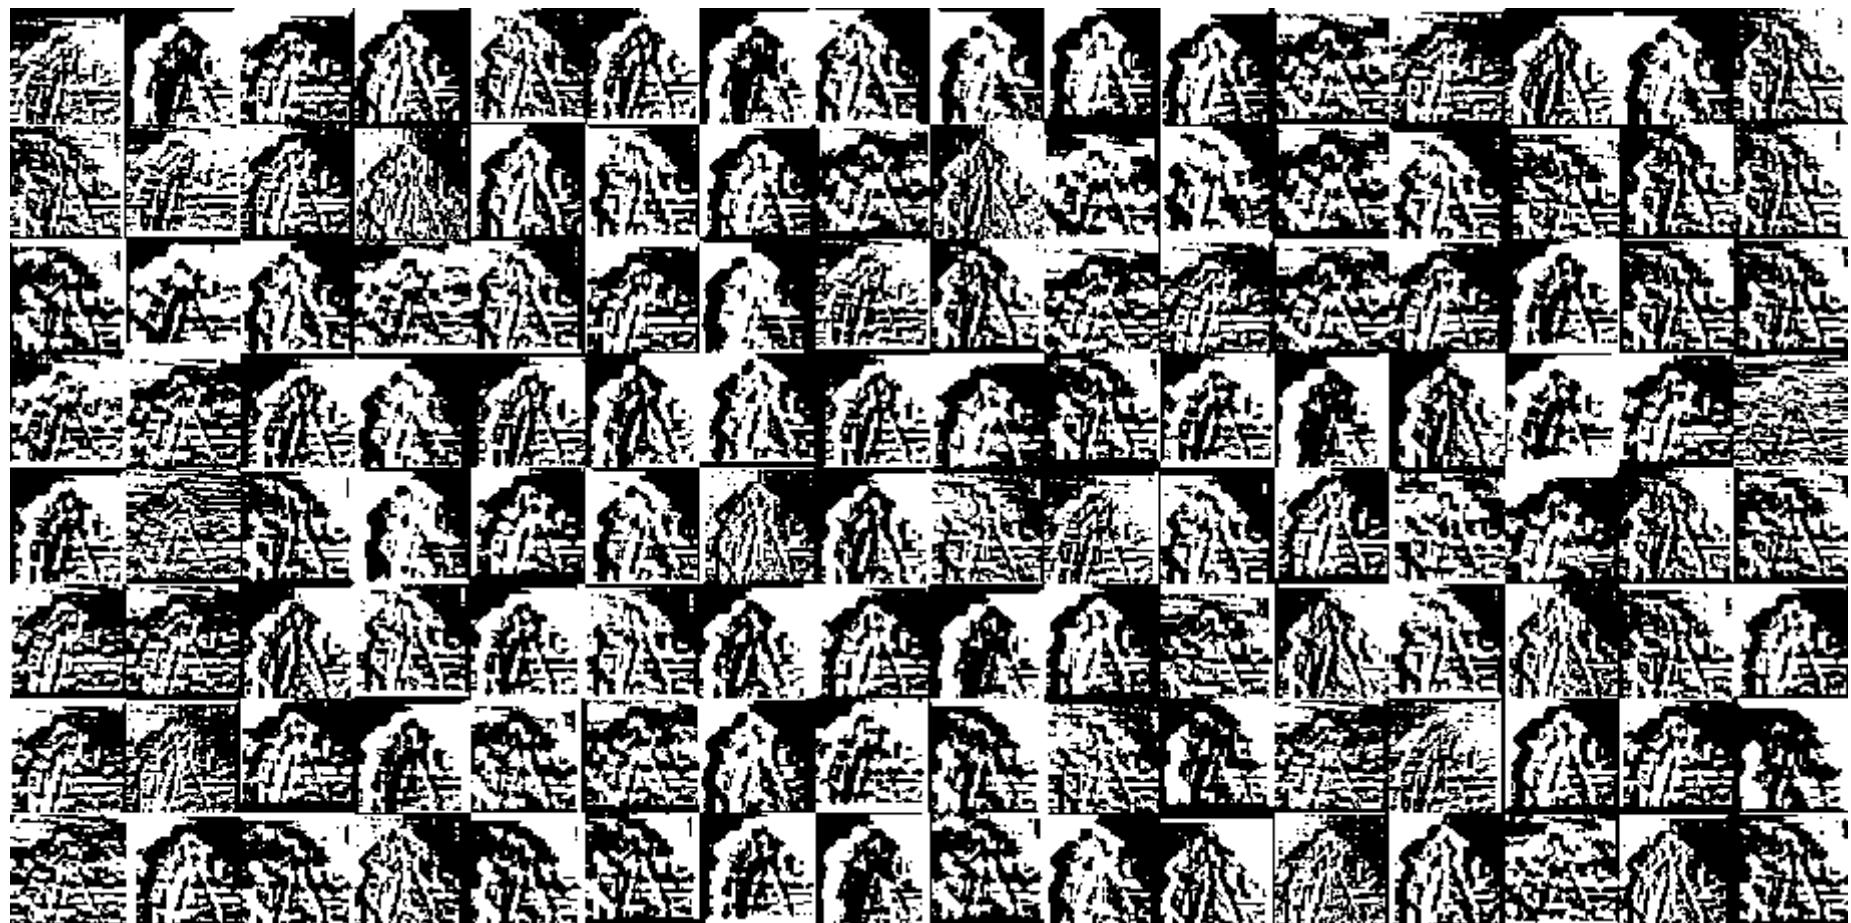
- Warp descriptor images, or warp image then recompute descriptors?
- Does it work with more sophisticated binary descriptors?



Chapters 5 and 6

# Additional Details

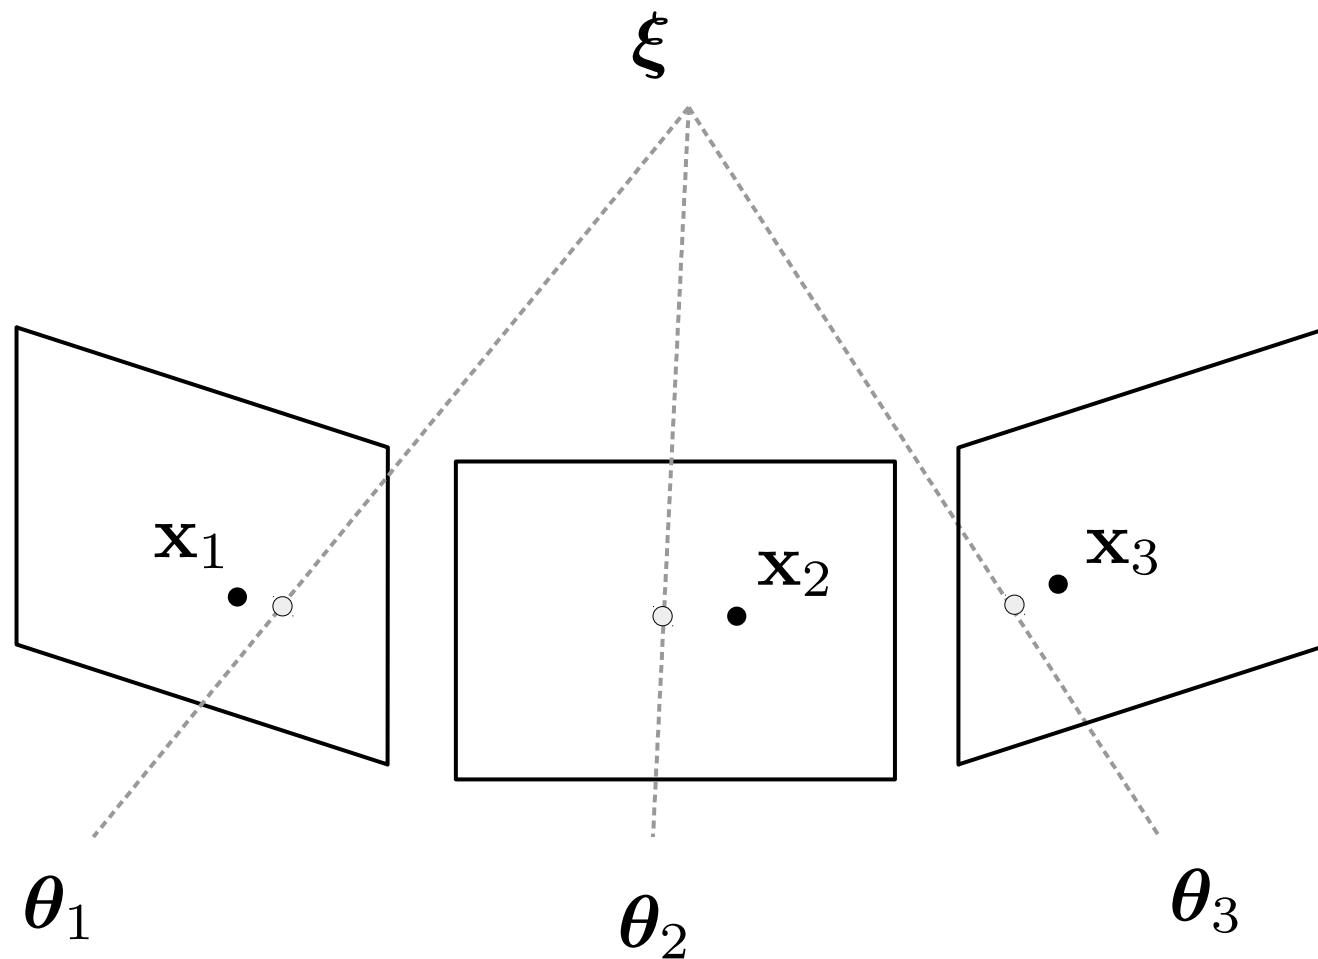
e.g.: BRIEF (128 Bit-Planes)



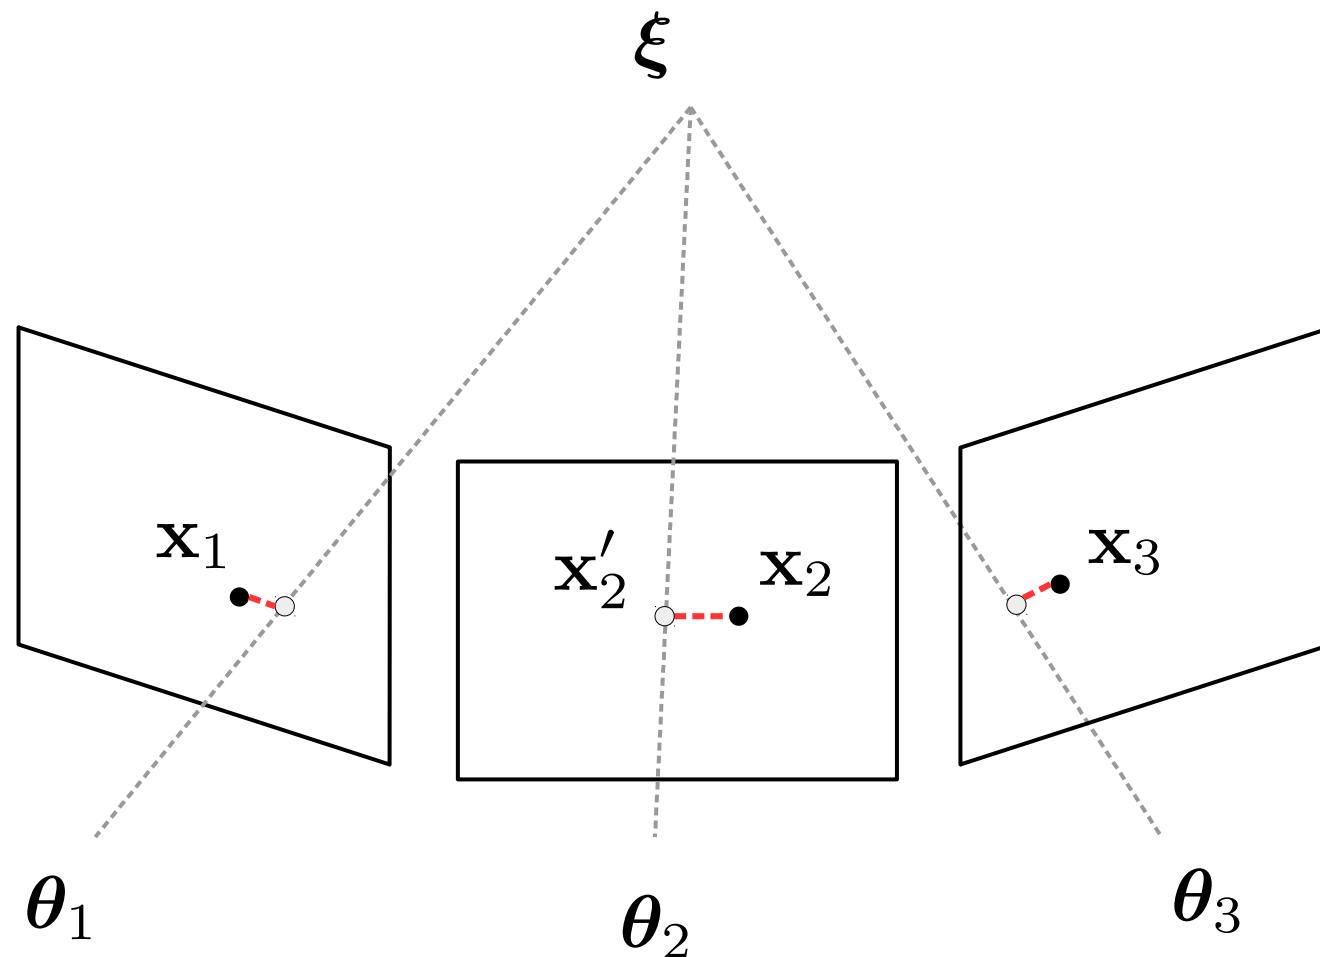
# Multi-view Refinement

- ▶ Features remain the method of choice in VSLAM
  - Due to joint refinement of pose and structure over multiple views using **bundle adjustment**
- ▶ Can we formulate a *Direct Bundle Adjustment*?
  - Bundle adjustment without correspondences
  - No longer need corner/edge structures

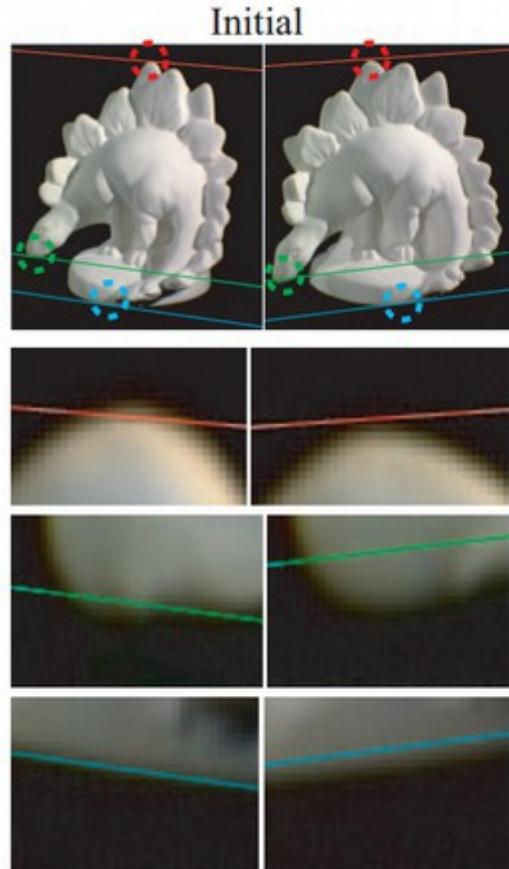
# Geometric Bundle Adjustment



# Geometric Bundle Adjustment



# Problems with the Reprojection Error

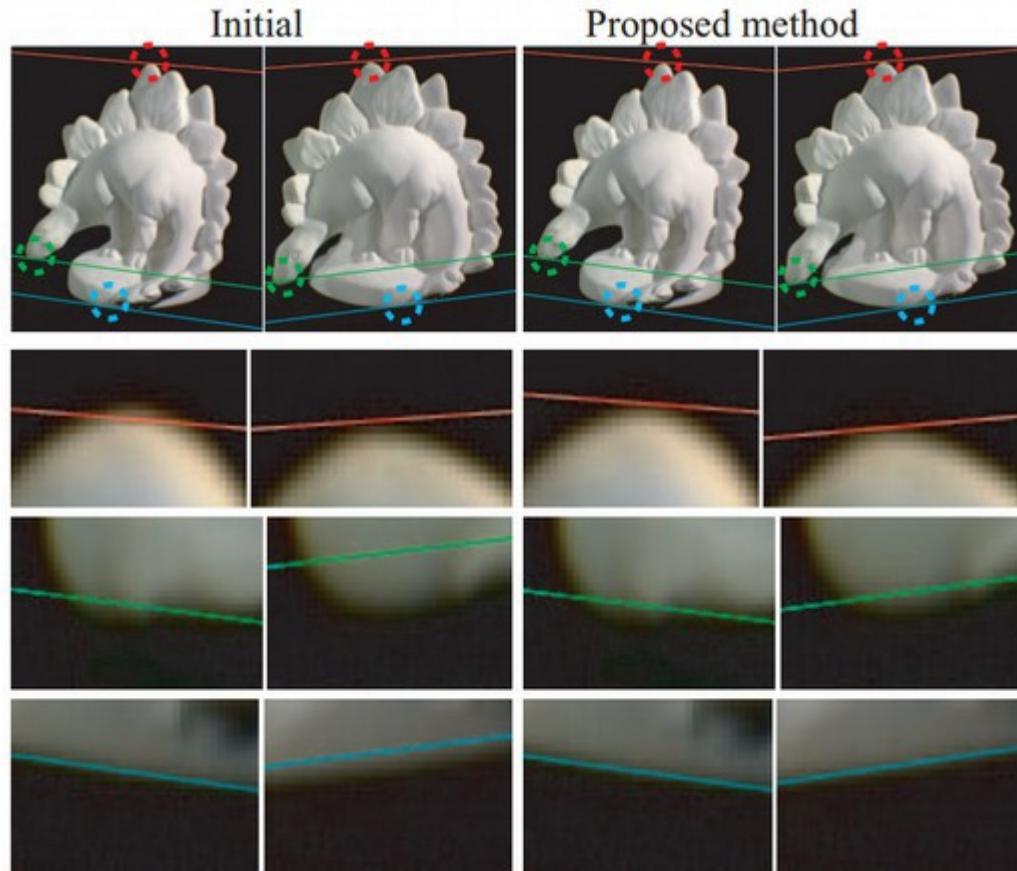


Attribute errors in feature localization to slight miscalibration

Interleave correspondence refinement with geometric BA

Furukawa & Ponce, CVPR 2008

# Problems with the Reprojection Error

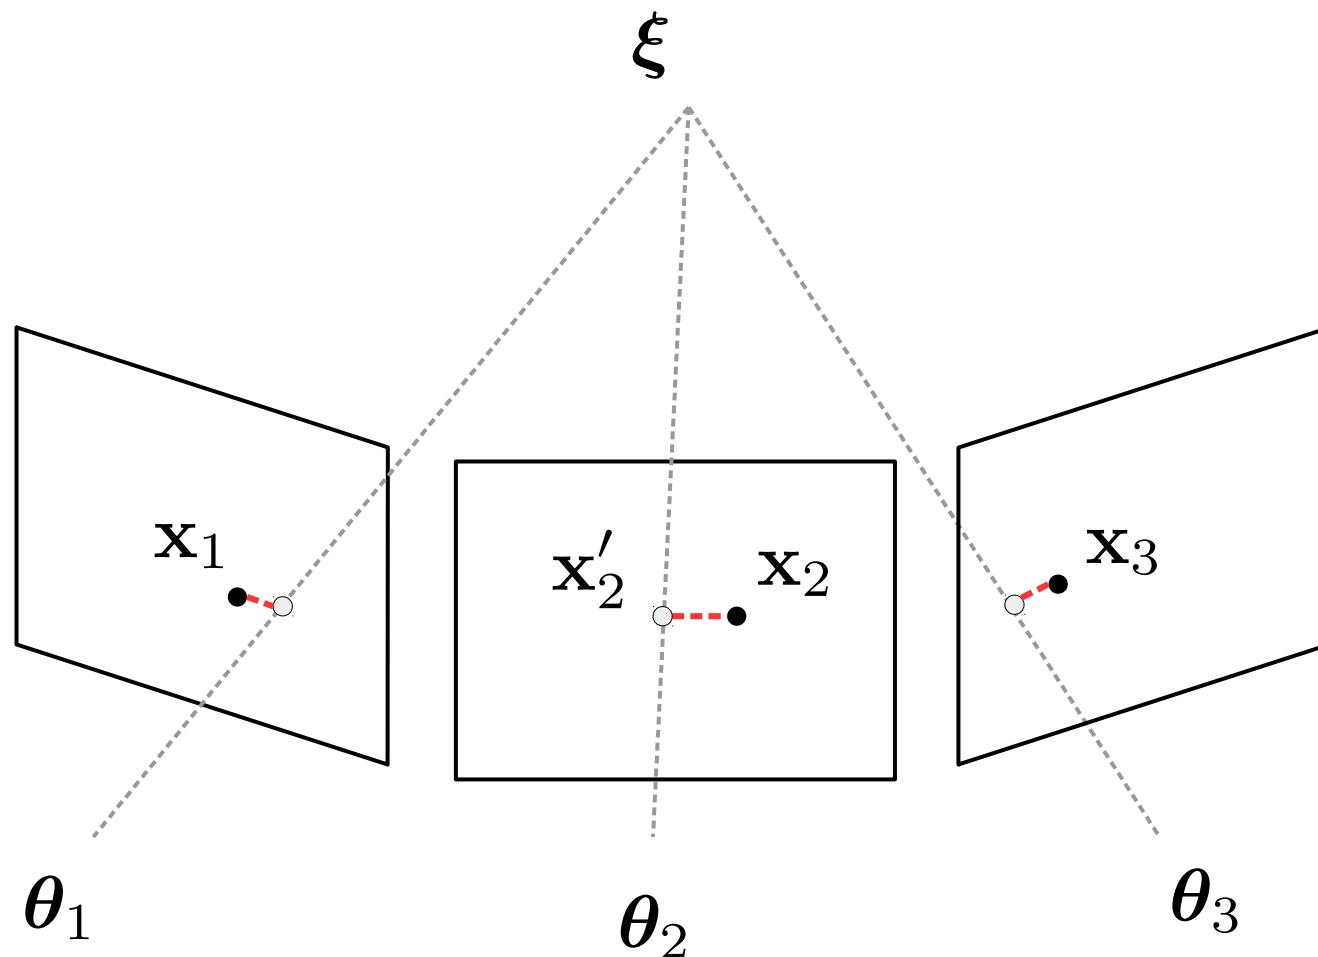


Attribute errors in feature localization to slight miscalibration

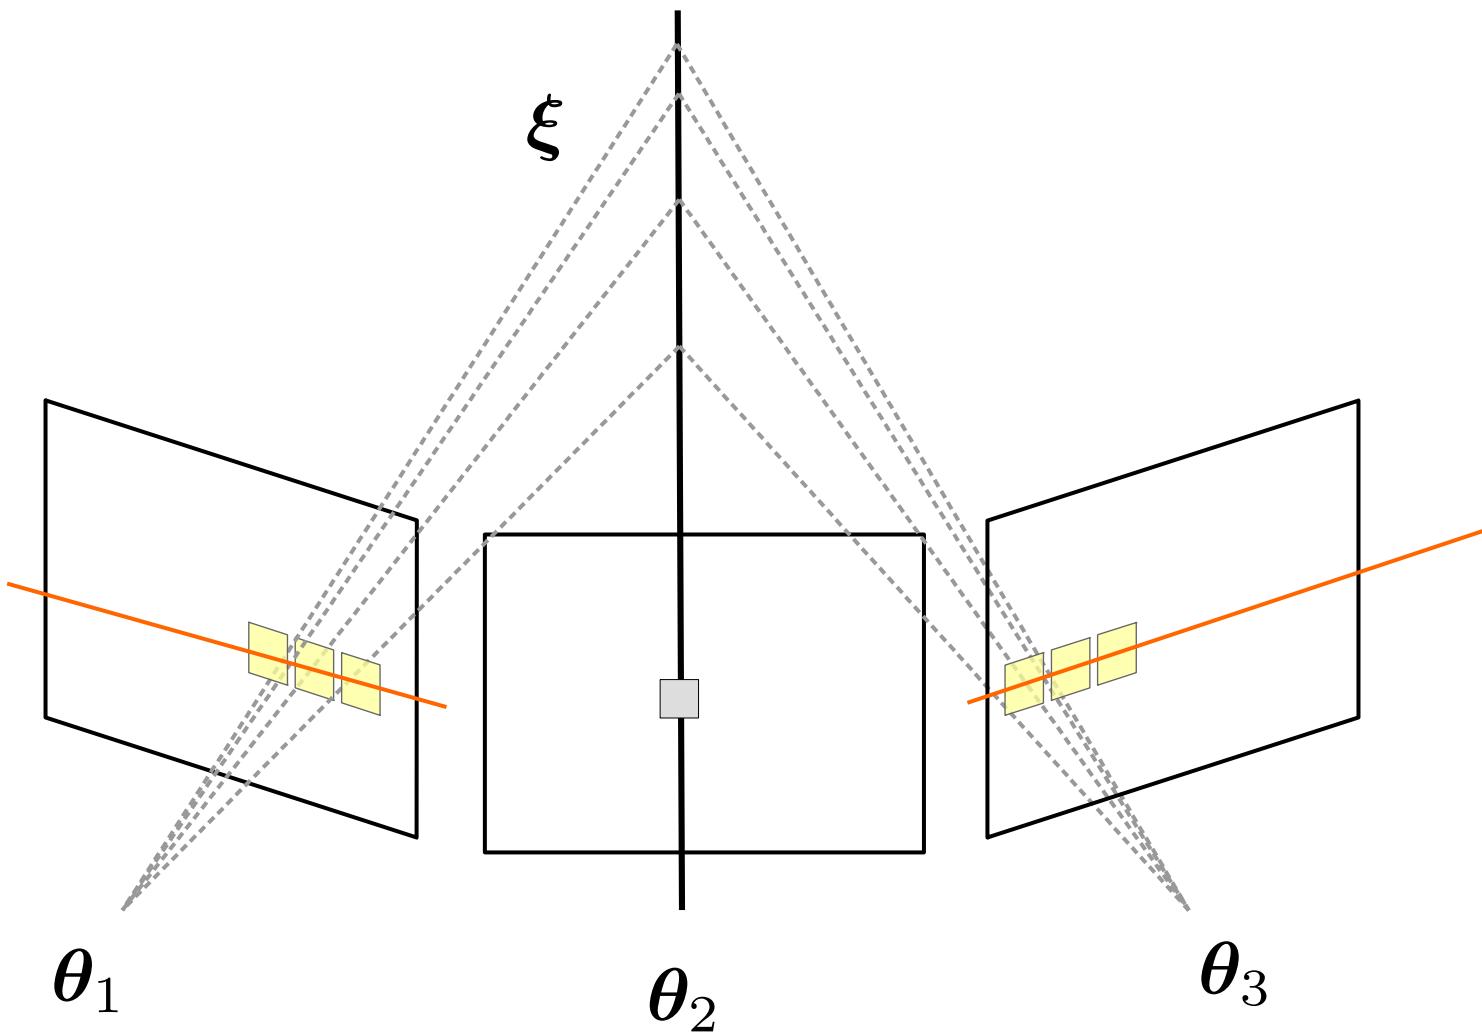
Interleave correspondence refinement with geometric BA

Furukawa & Ponce, CVPR 2008

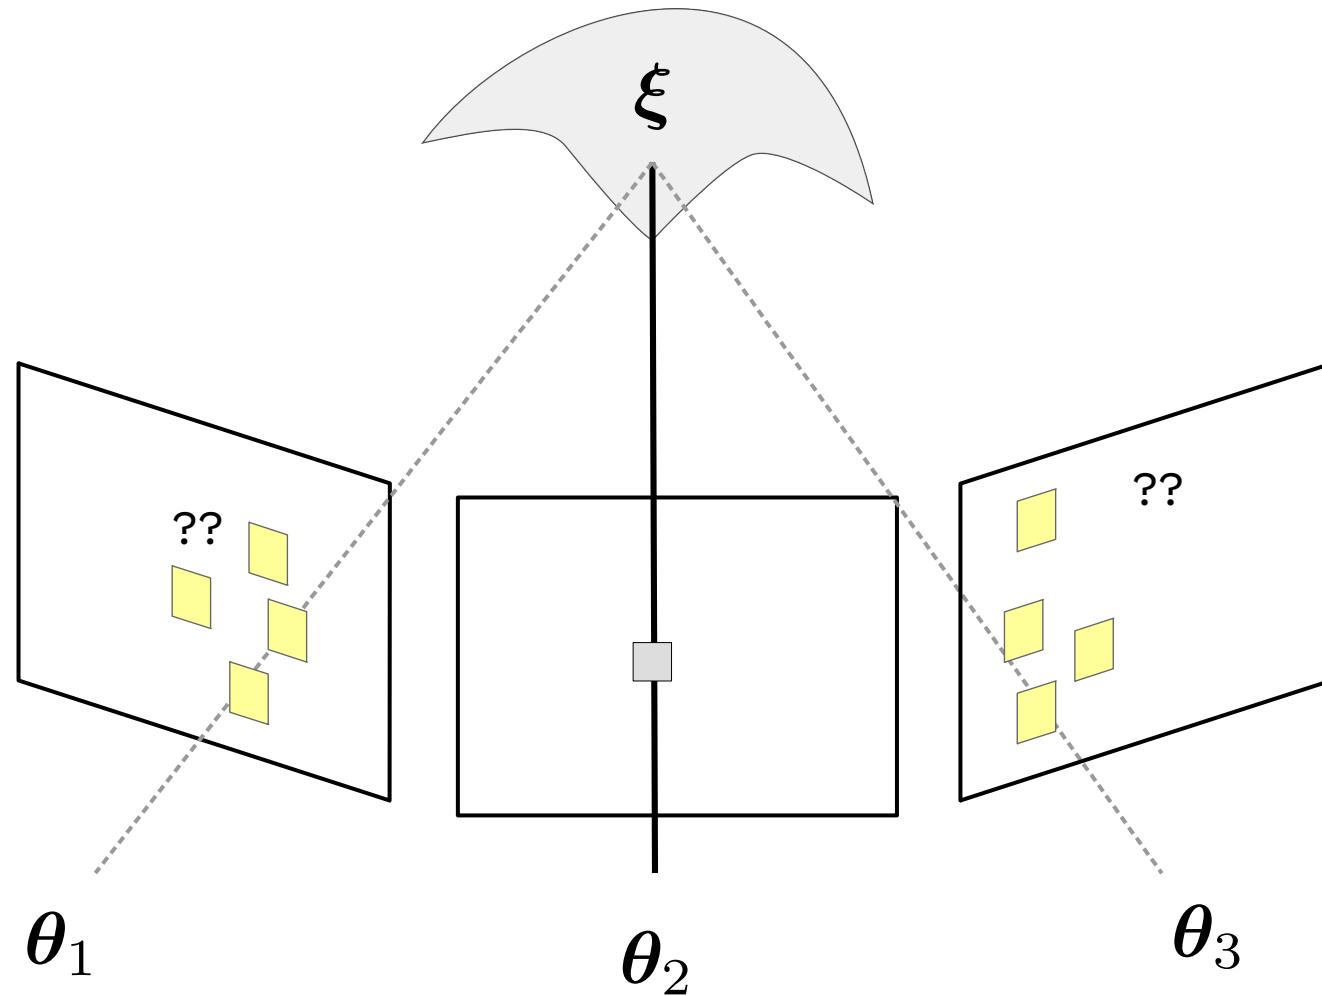
# Geometric Bundle Adjustment



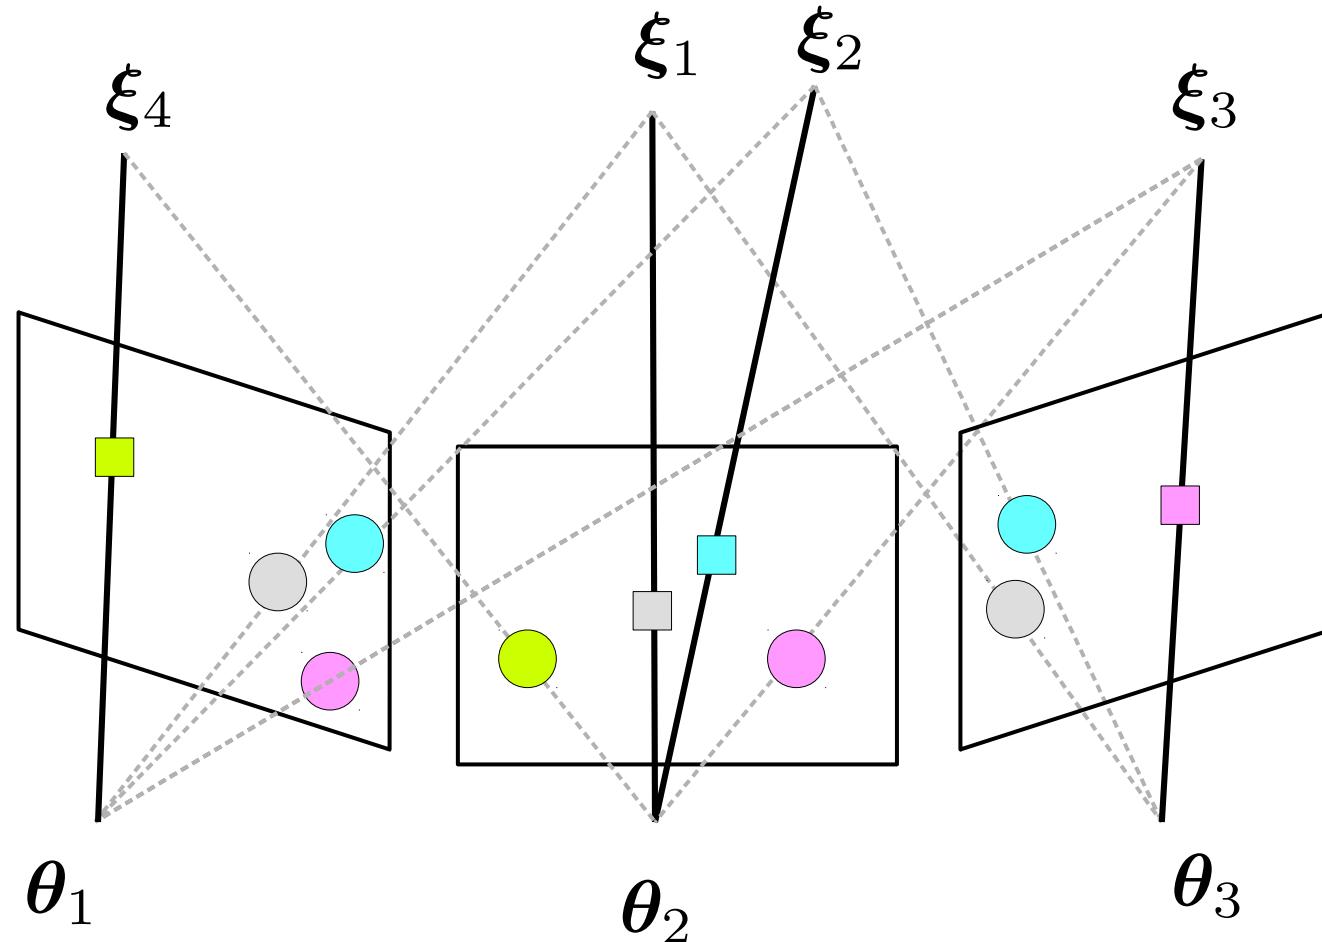
# Multi-view Stereo



# Photometric Bundle Adjustment



# Photometric Bundle Adjustment



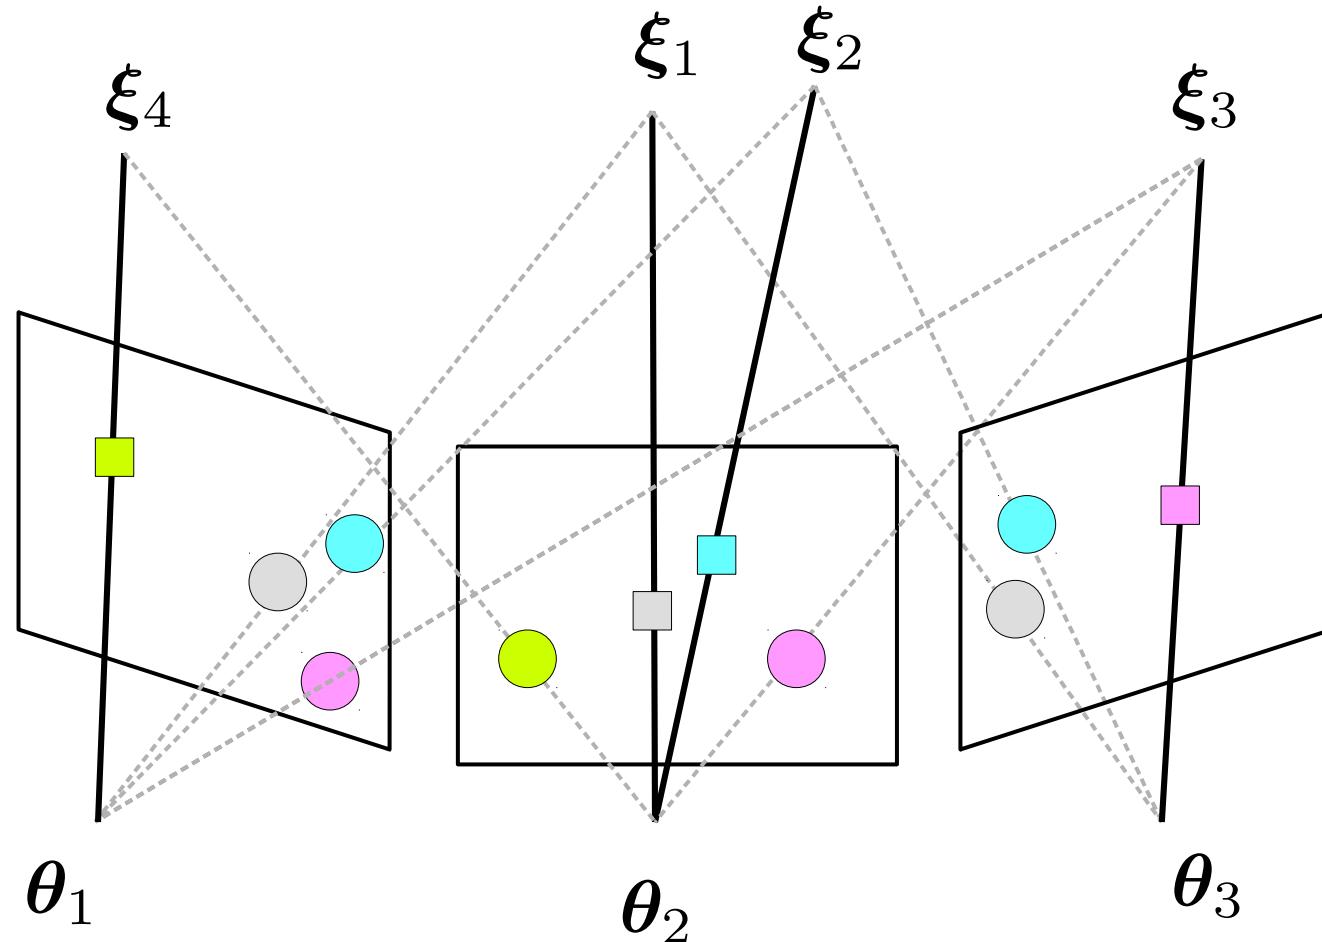
Patch at reference frame



Potential search area

Need to determine  
visibility info

# Photometric Bundle Adjustment



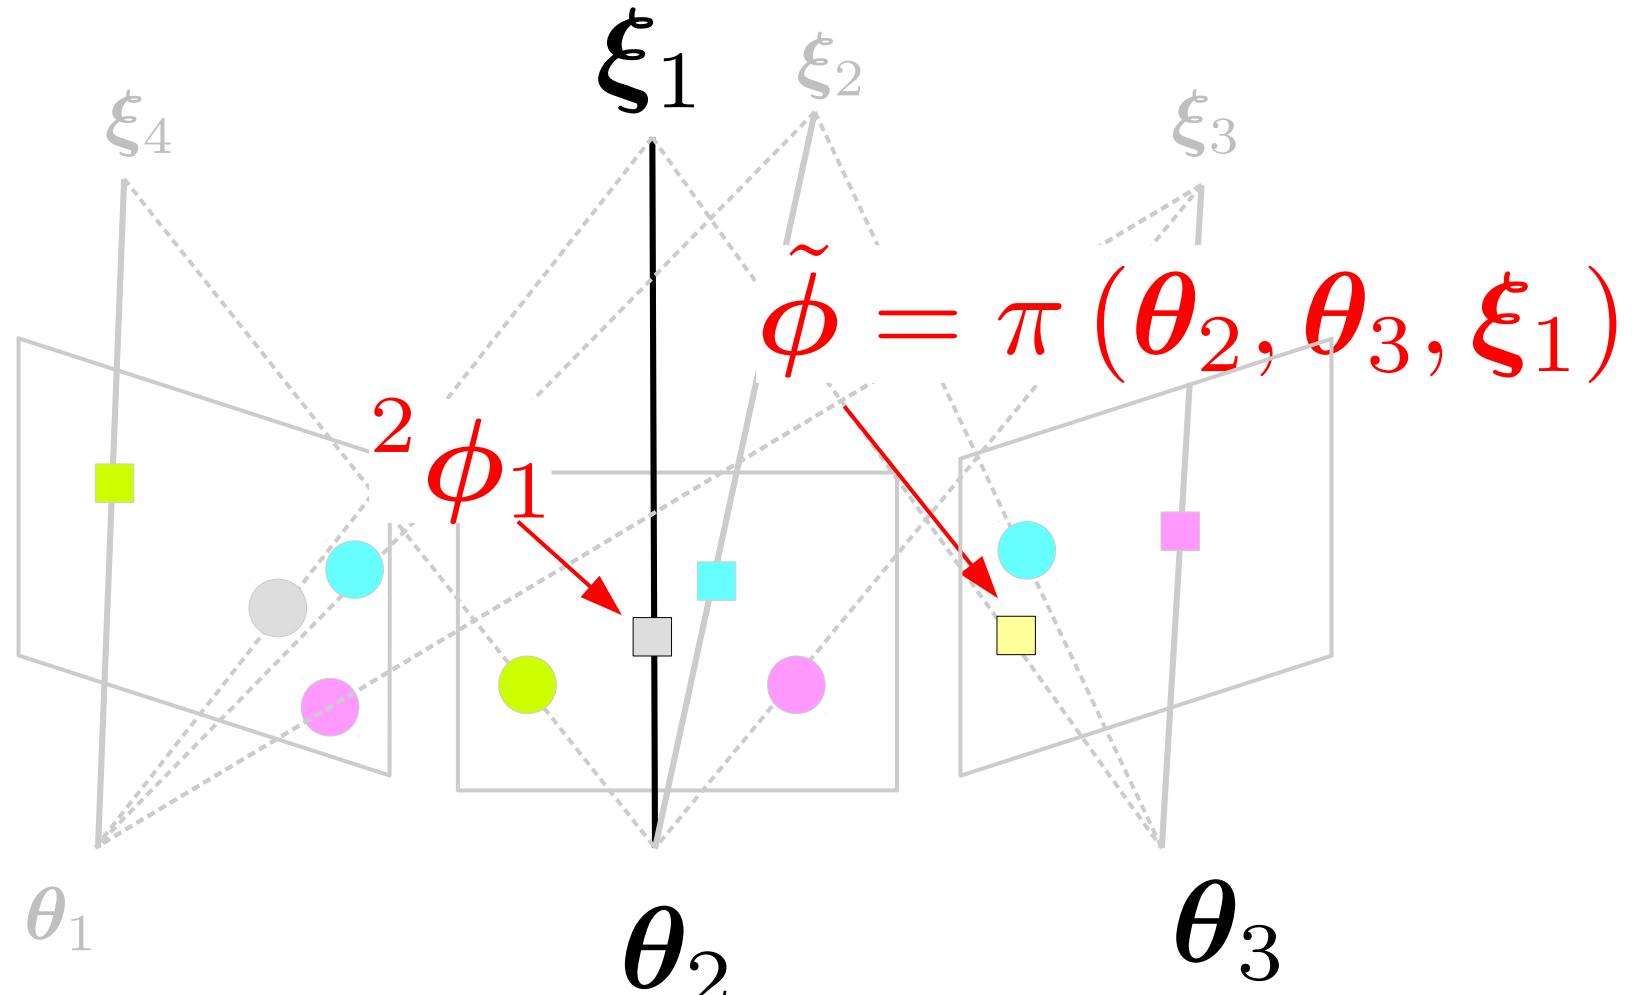
Patch at reference frame



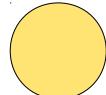
Potential search area

Need to determine  
visibility info

# Photometric Bundle Adjustment



Fixed patch at reference frame



Potential search area

# Objective Per Point

$$\| {}^r \phi_j - \tilde{\phi}(\theta_r, \theta_i, \xi_j) \|_2^2$$

Point ‘j’ with reference frame ‘r’ projected at frame ‘i’

# Objective Per Point

$$\| {}^r \phi_j - \tilde{\phi}(\theta_r, \theta_i, \xi_j) \|_2^2$$

Point ‘j’ with reference frame ‘r’ projected at frame ‘i’

Objective now depends on 2 poses  
Creates additional terms in the Hessian

# Objective Per Point

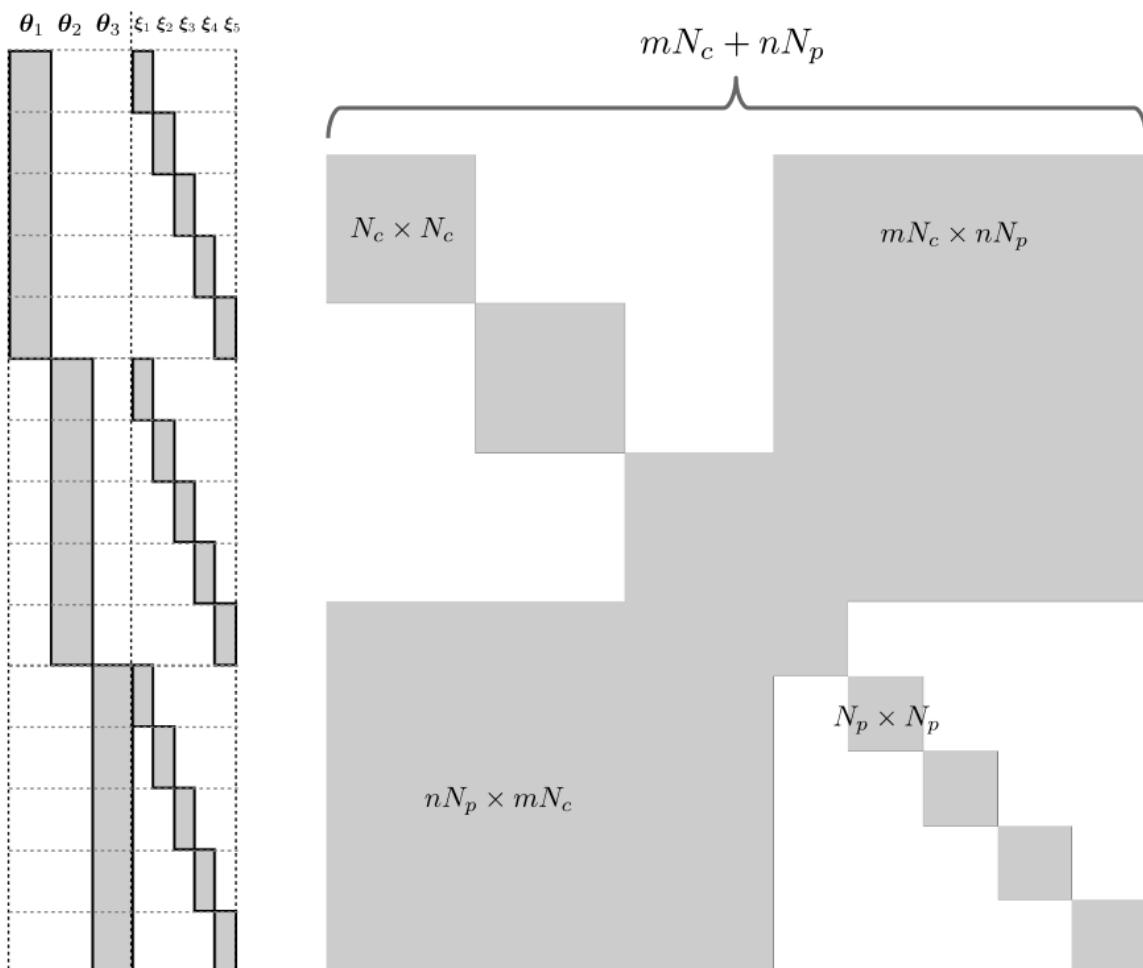
$$\| {}^r \phi_j - \tilde{\phi}(\theta_r, \theta_i, \xi_j) \|_2^2$$

Point ‘j’ with reference frame ‘r’ projected at frame ‘i’

Objective now depends on 2 poses  
Creates additional terms in the Hessian

Make the reference patch **fixed**  
(store it in a buffer)

# Sparsity Pattern



Hessian is Identical to geometric BA

Jacobian is slightly denser

For example, using a 3x3 patch results in 9 instead of 2 residuals

$$\mathbf{J}(\boldsymbol{\theta}) = \nabla \mathbf{I}(\mathbf{u}' + \mathbf{u}) \frac{\partial \mathbf{u}'}{\partial \boldsymbol{\theta}}$$

$$\mathbf{J}(\boldsymbol{\xi}) = \nabla \mathbf{I}(\mathbf{u}' + \mathbf{u}) \frac{\partial \mathbf{u}'}{\partial \boldsymbol{\xi}}$$

The same Jacobian as in geometric BA, but multiplied by warped image gradient

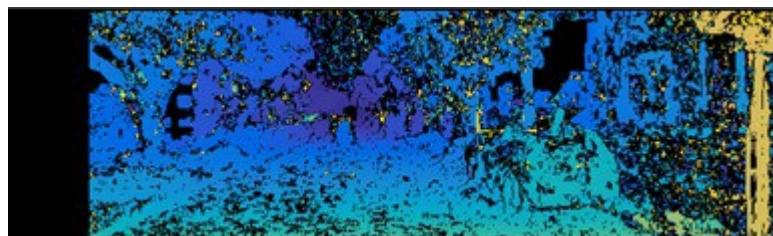
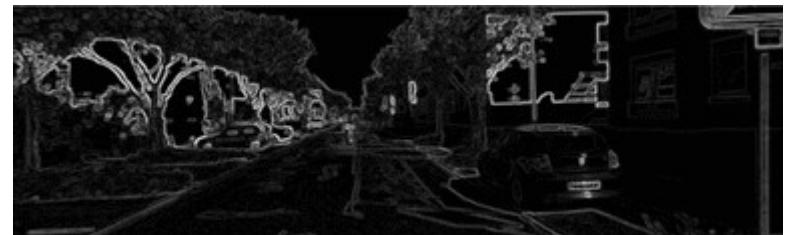
# Point/Pixel Selection

- Has depth and is local max of gradient magnitude

Input image



Gradient magnitude



Depth initialization



Selected pixels

Pixels are selected as **integer** positions

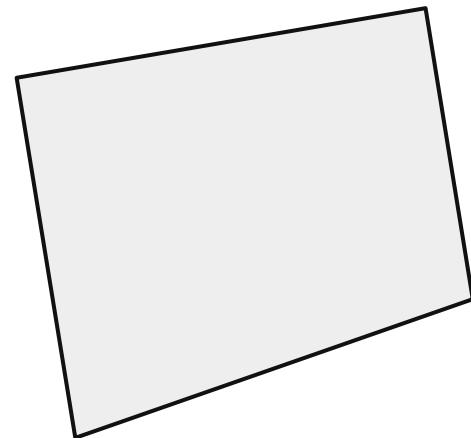
~4096 points per frame on KITTI data

# Visibility

- ▶ Project the current set of 3D points onto new frame

$$\{\xi_j, \phi_j\}_{j=1}^N$$

Initialization of the scene  
points thus far, each with its  
photometric data and a 5x5  
intensity patch



Frame  $(m + 1)$

# Visibility

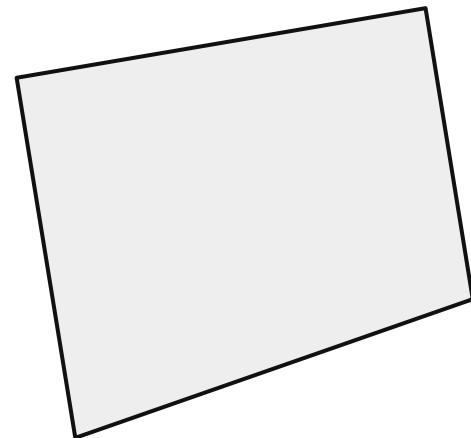
- ▶ Project the current set of 3D points onto new frame

$$\{\xi_j, \phi_j\}_{j=1}^N$$

Initialization of the scene points thus far, each with its photometric data and a 5x5 intensity patch

$$\theta_{m+1}$$

Pose initialization for frame (m+1)



Frame ( $m + 1$ )

# Visibility

- ▶ Project the current set of 3D points onto new frame

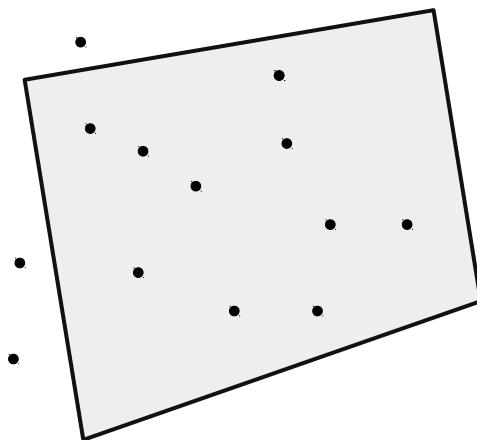
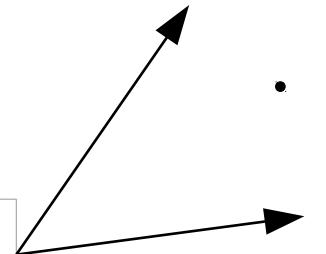
$$\{\xi_j, \phi_j\}_{j=1}^N$$

$$\theta_{m+1}$$

Pose initialization for frame (m+1)

Initialization of the scene points thus far, each with its photometric data and a 5x5 intensity patch

Projections using pose initialization onto new frame



Frame ( $m + 1$ )

# Visibility

- Project the current set of 3D points onto new frame

$$\{\xi_j, \phi_j\}_{j=1}^N$$

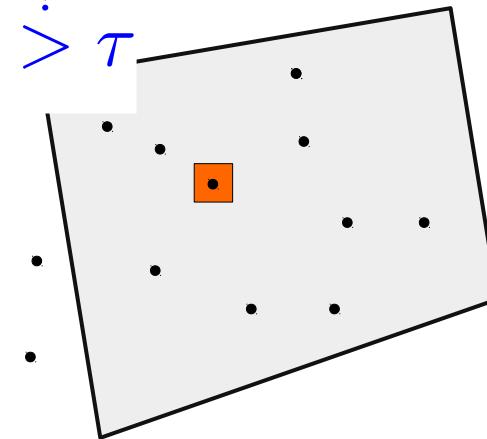
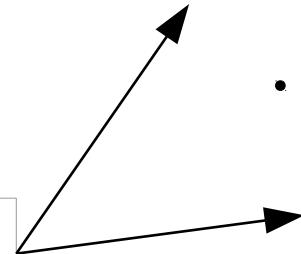
Initialization of the scene points thus far, each with its photometric data and a 5x5 intensity patch

$$\theta_{m+1}$$

Pose initialization for frame (m+1)

$$\text{zncc}(\phi_j, \tilde{\phi}) > \tau$$

Projections using pose initialization onto new frame

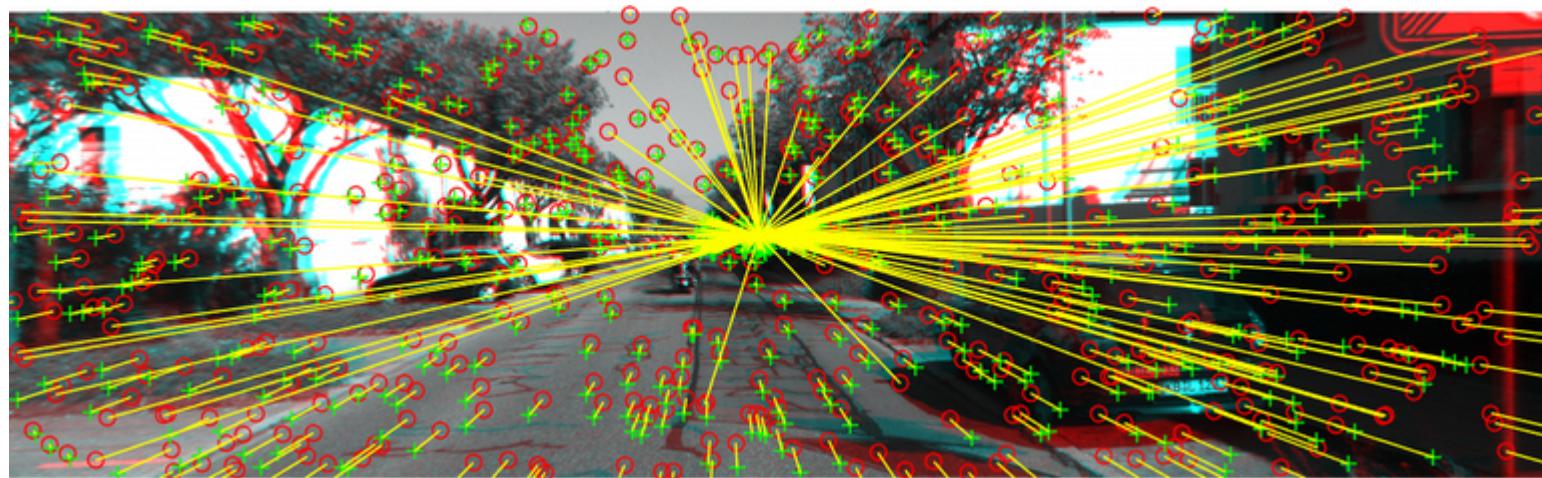


Frame ( $m + 1$ )

# Visibility

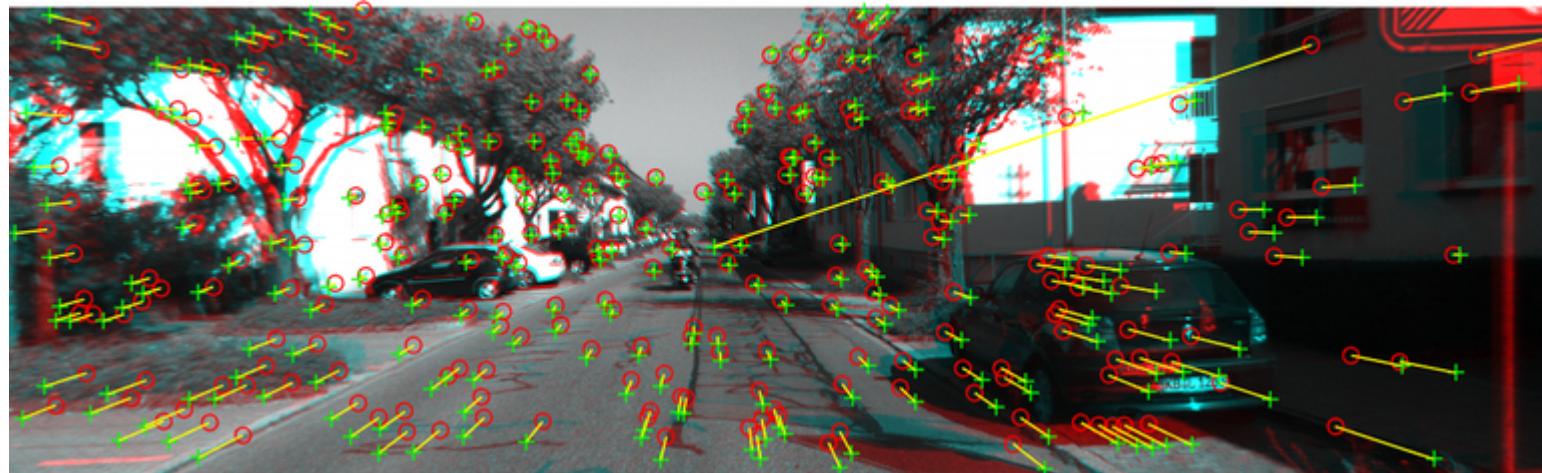
- ▶ Project the set of 3D points onto new frame

Depends on accuracy of pose and  
structure initialization



# Visibility

- ▶ Project the set of 3D points onto new frame
- ▶ Outliers handled with robust minimization
  - Huber loss

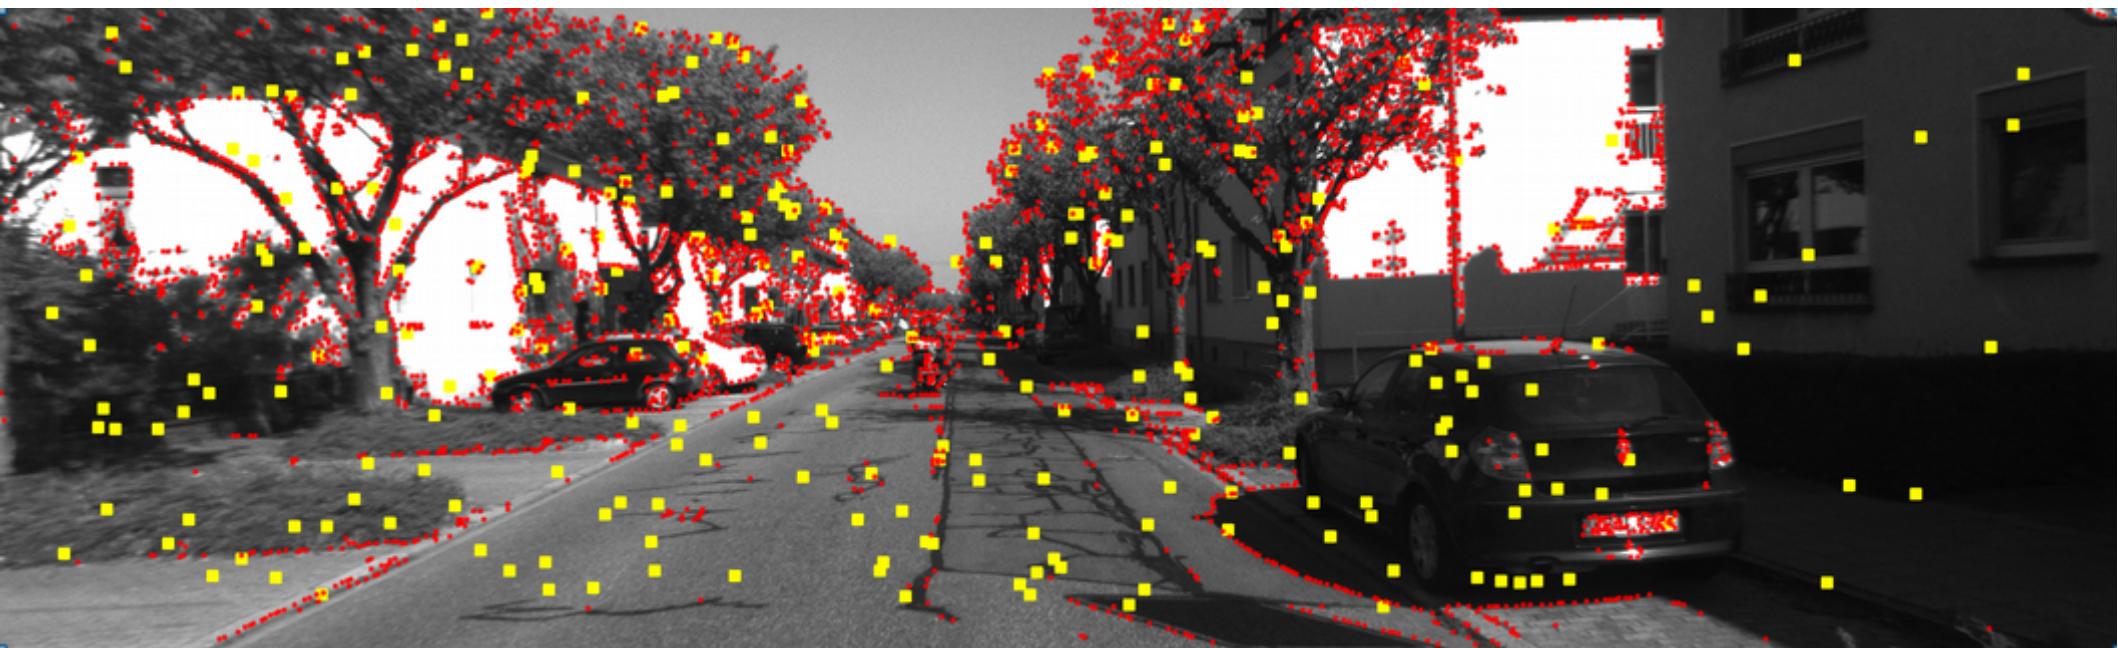


# Implementation details

- ▶ 3x3 patch photometric error
  - Experimented with center-weighted error
    - No significant difference
- ▶ 5x5 patch for zncc (with threshold = 0.6)
- ▶ Sliding window with 5 frames
- ▶ Depth initialization from block matching stereo

# One Last Detail

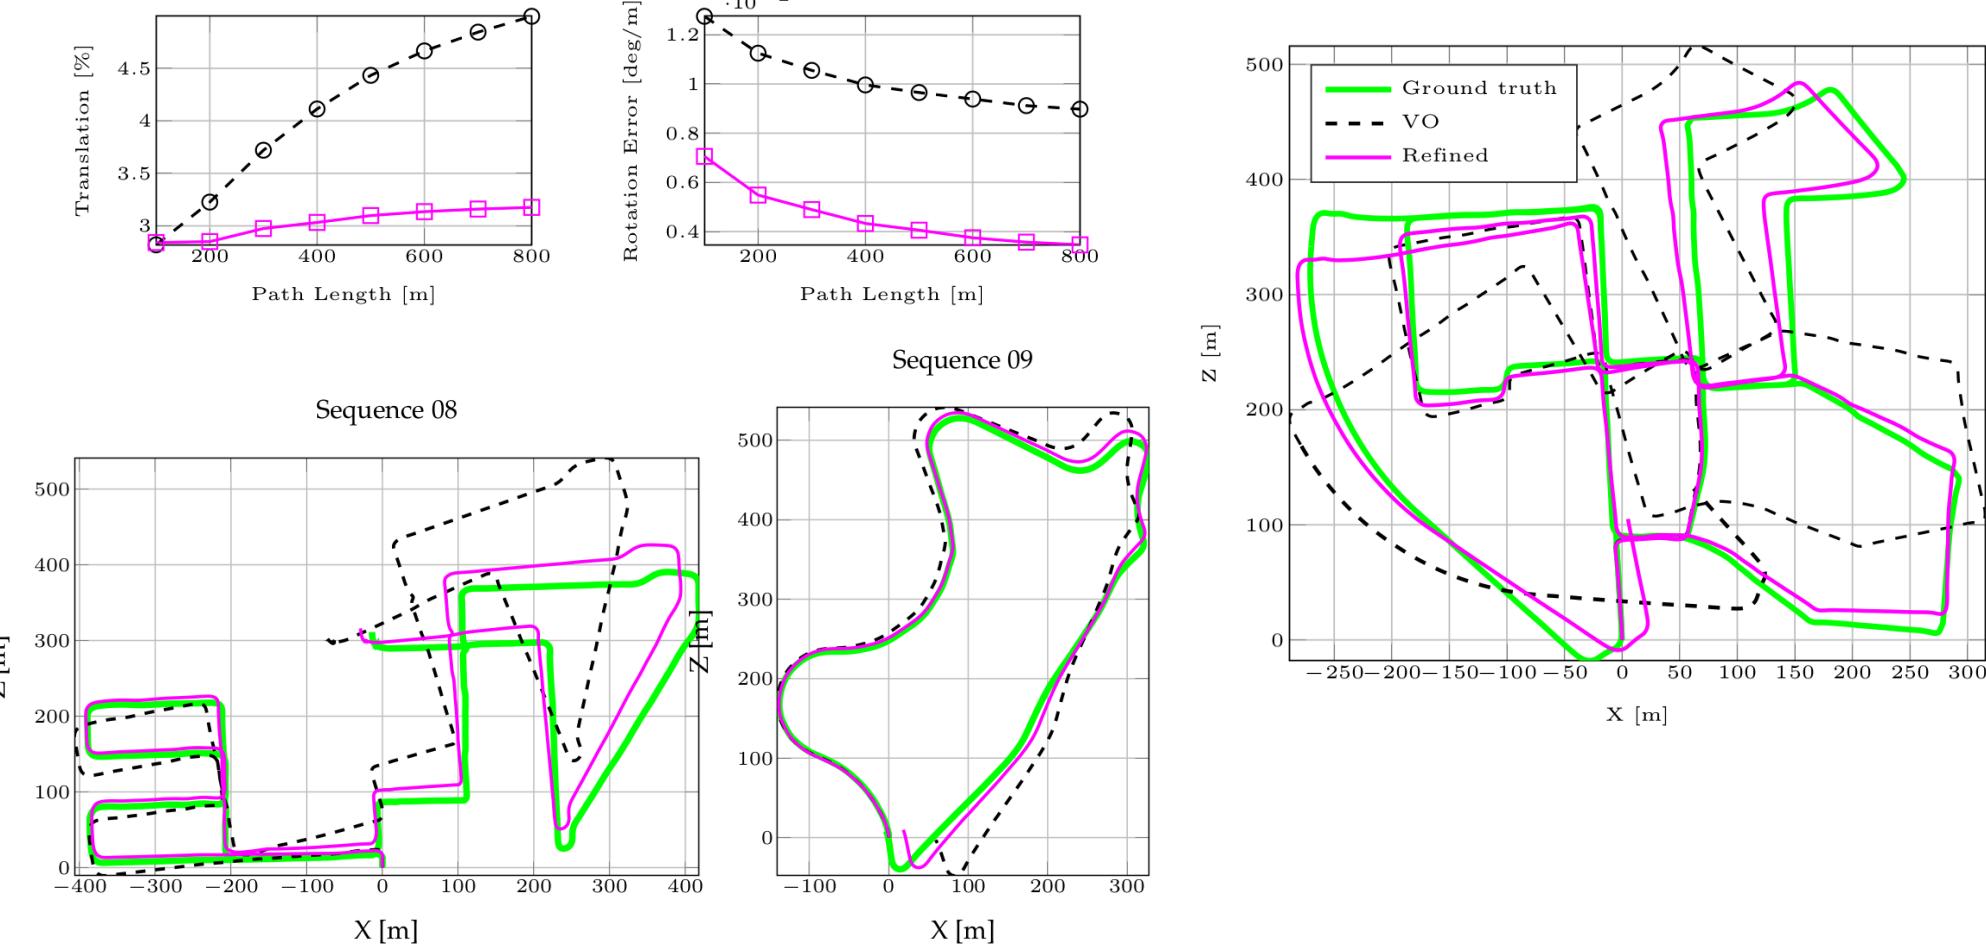
- ▶ Avoid adding duplicate points



- Location of a projected point from the previous frame
- Candidates to initialize new scene points

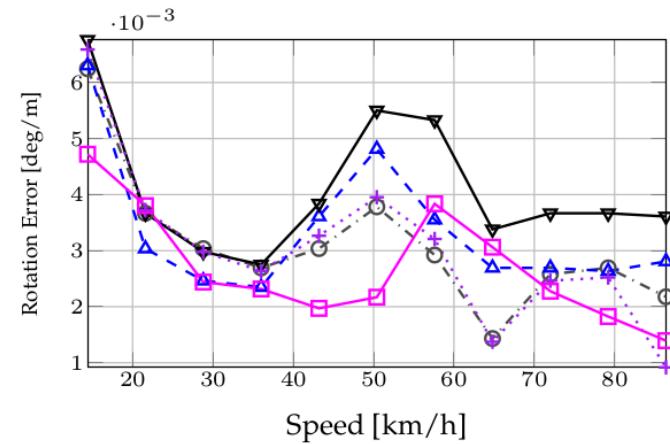
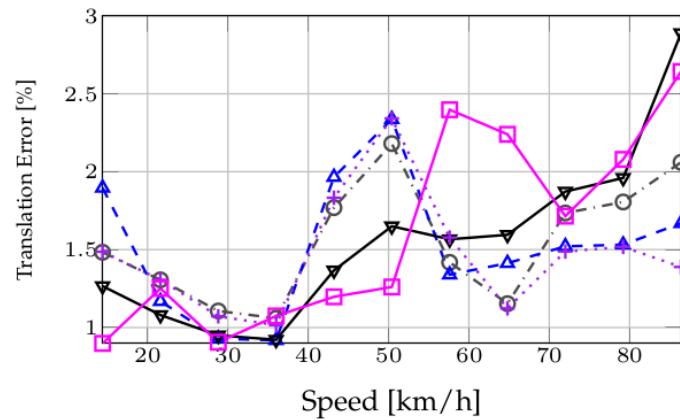
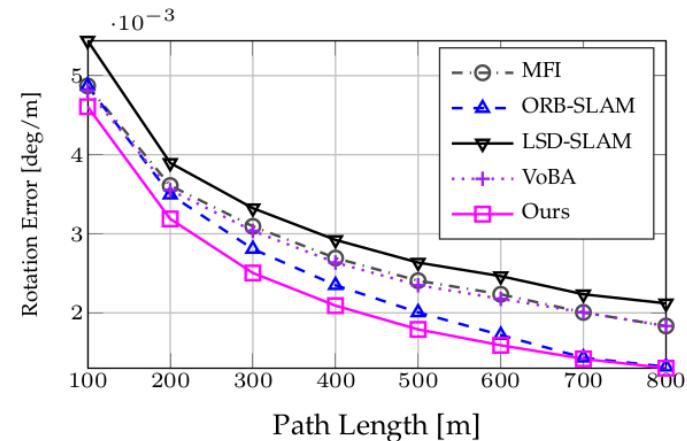
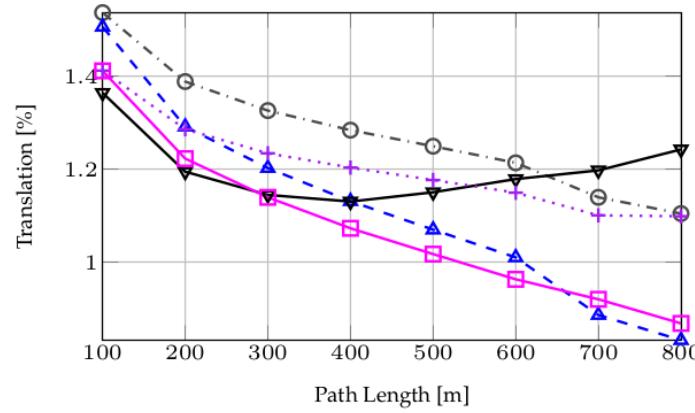
# Improvement on poor initialization

Pose init from direct VO; Depth init from Block Matching stereo



# Improvement on good initialization

Pose init from ORB-SLAM; Depth init from Block Matching stereo

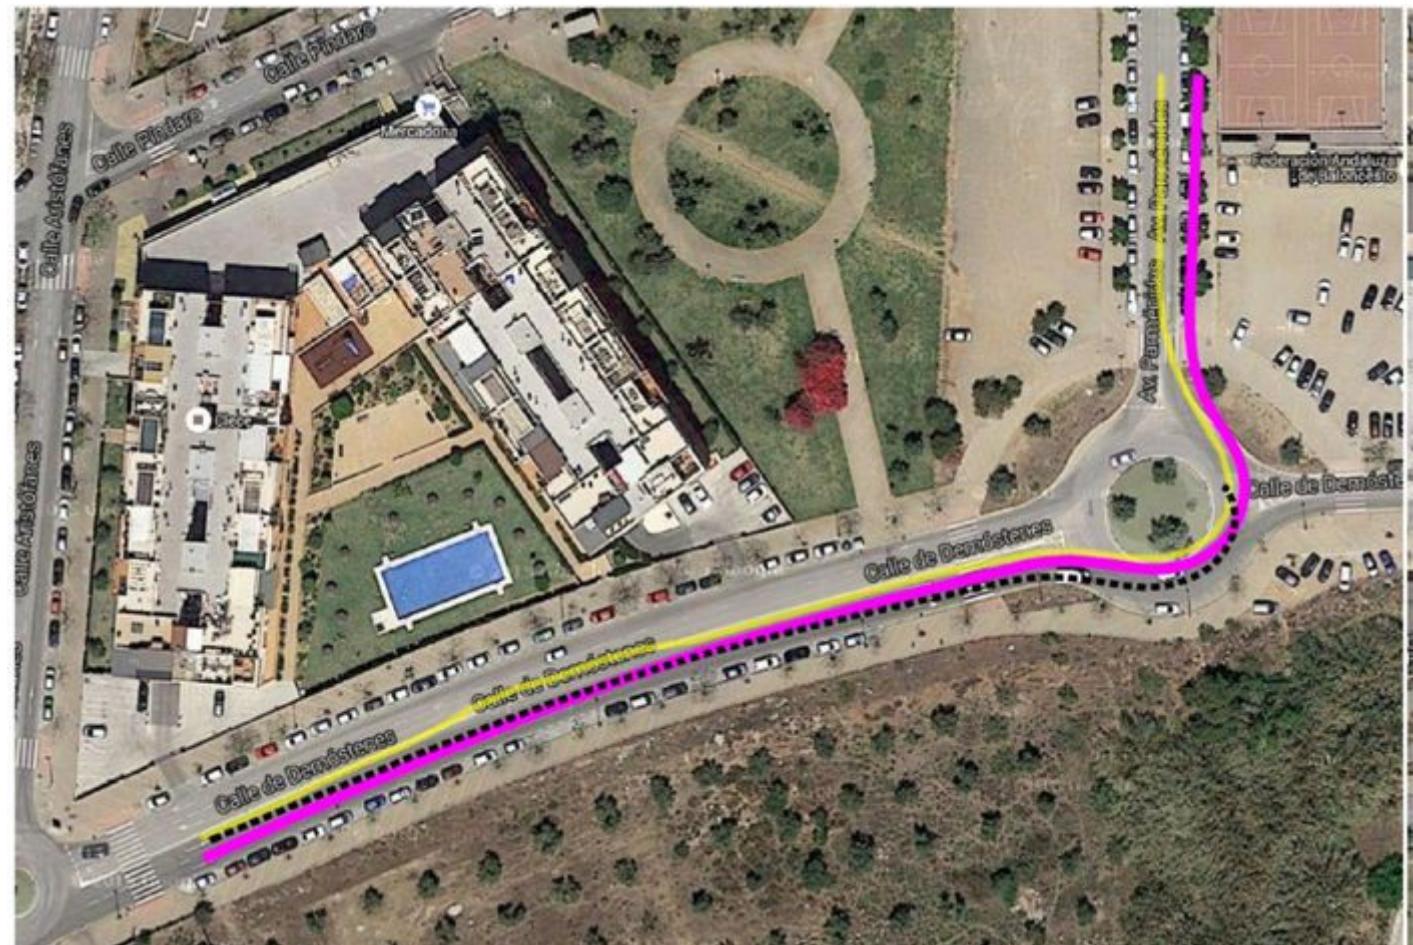


ORB-SLAM does traditional bundle adjustment + loop closure

# Results with no initialization

Higher frame-rate data

----- ORB-SLAM  
\_\_\_\_\_ Ours



# Open Questions

- ▶ Theoretical characterization of photometric BA
  - Convergence basin and conditions, motion...
- ▶ Theoretical/experimental understanding of why it can improve on geometric BA
  - Could improve current geometric BA
- ▶ Good results with simple techniques
  - Other options for: pixel selection, visibility, descriptors, ...

# Summary

- ▶ Extended the state-of-the-art in pose estimation to challenging domains
  - Low light, sudden & drastic appearance change, specular reflections, ...
- ▶ Demonstrated faster than real-time results for template tracking and visual odometry
- ▶ Presented the 1<sup>st</sup> formulation of direct bundle adjustment for VSLAM
  - Improves results obtained with geometric BA

# Conclusions

- ▶ Direct alignment using binary descriptors is a robust solution to pose estimation
  - When the inter-frame displacement is small
  - Use Bit-Planes for correct objective in least-squares
- ▶ Direct (photometric) bundle adjustment is a feasible solution for VSLAM and can improve accuracy

Thank you