

# Image stitching

Digital Visual Effects  
*Yung-Yu Chuang*

*with slides by Richard Szeliski, Steve Seitz, Matthew Brown and Vaclav Hlavac*

## Applications of image stitching

- Video stabilization
- Video summarization
- Video compression
- Video matting
- Panorama creation

DigiVFX

## Image stitching

- Stitching = alignment + blending

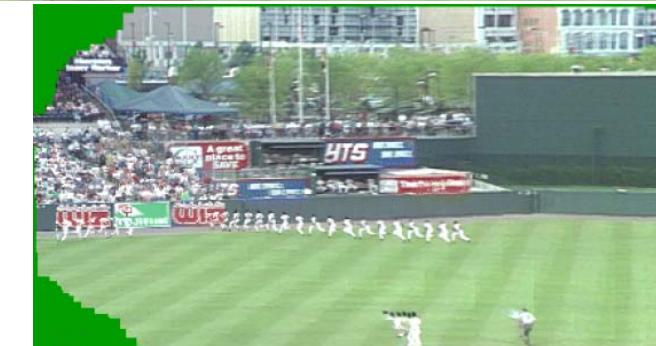
geometrical registration

photometric registration



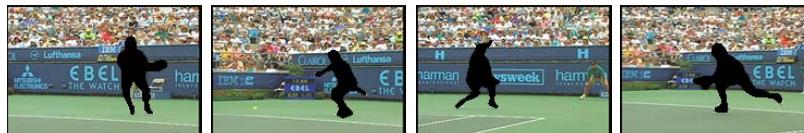
## Video summarization

DigiVFX



## Video compression

DigiVFX



## Object removal

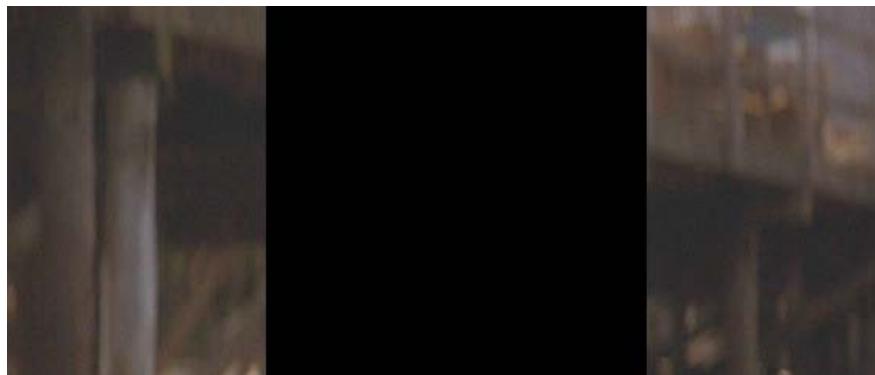
DigiVFX



input video

## Object removal

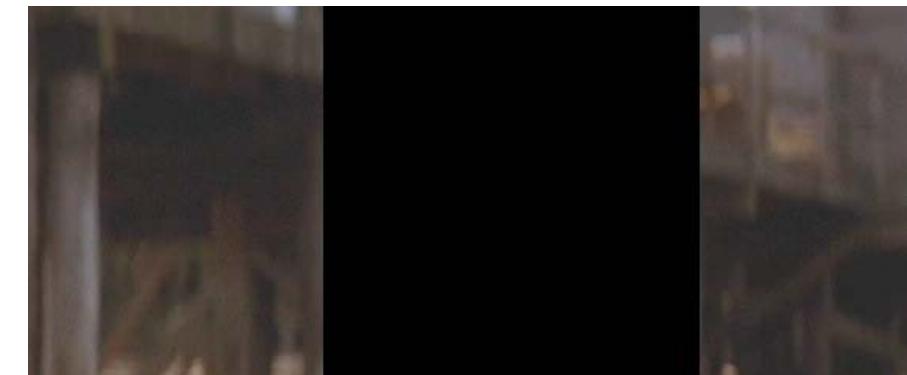
DigiVFX



remove foreground

## Object removal

DigiVFX



estimate background

## Object removal

DigiVFX



background estimation

## Panorama creation

DigiVFX



## Why panorama?

DigiVFX

- Are you getting the whole picture?
  - Compact Camera FOV =  $50 \times 35^\circ$



## Why panorama?

DigiVFX

- Are you getting the whole picture?
  - Compact Camera FOV =  $50 \times 35^\circ$
  - Human FOV =  $200 \times 135^\circ$



## Why panorama?

DigiVFX

- Are you getting the whole picture?
  - Compact Camera FOV =  $50 \times 35^\circ$
  - Human FOV =  $200 \times 135^\circ$
  - Panoramic Mosaic =  $360 \times 180^\circ$



## What can be globally aligned?

DigiVFX

- In image stitching, we seek for a matrix to globally warp one image into another. Are any two images of the same scene can be aligned this way?
  - Images captured with the same center of projection
  - A planar scene or far-away scene

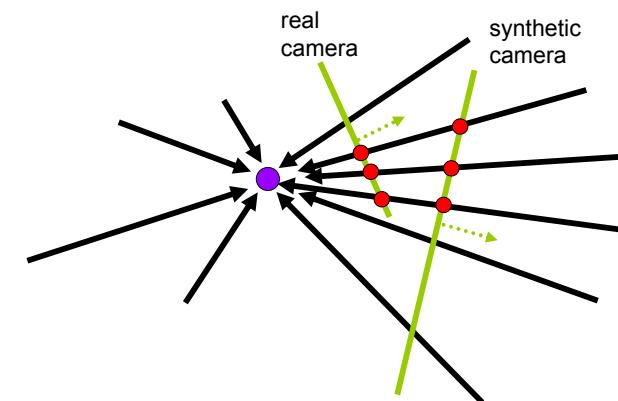
## Panorama examples

DigiVFX

- Like HDR, it is a topic of computational photography, seeking ways to build a better camera mostly in software.
- Most consumer cameras have a panorama mode
- Mars:  
[http://www.panoramas.dk/fullscreen3/f2\\_mars97.html](http://www.panoramas.dk/fullscreen3/f2_mars97.html)
- Earth:  
<http://www.panoramas.dk/new-year-2006/taipei.html>  
<http://www.360cities.net/>

## A pencil of rays contains all views

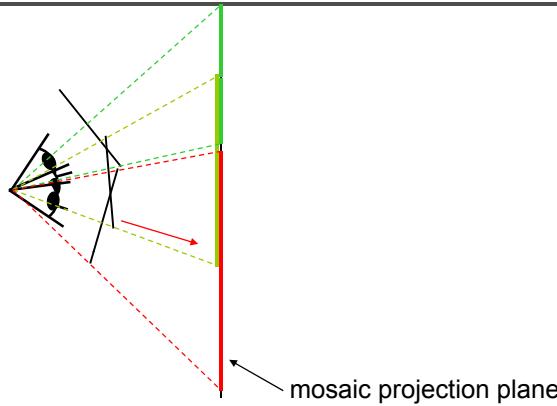
DigiVFX



Can generate any synthetic camera view  
as long as it has the **same center of projection!**

## Mosaic as an image reprojection

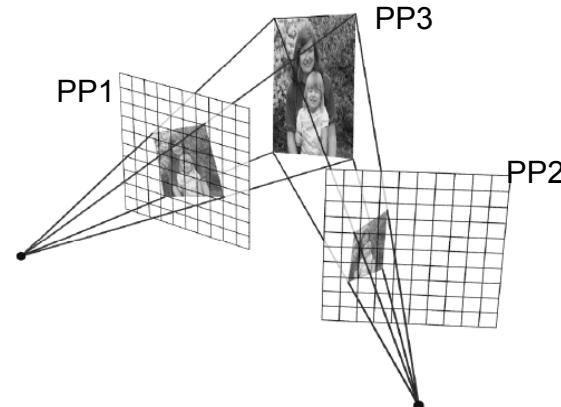
DigiVFX



- The images are reprojected onto a common plane
- The mosaic is formed on this plane
- Mosaic is a *synthetic wide-angle camera*

## Planar scene (or a faraway one)

DigiVFX

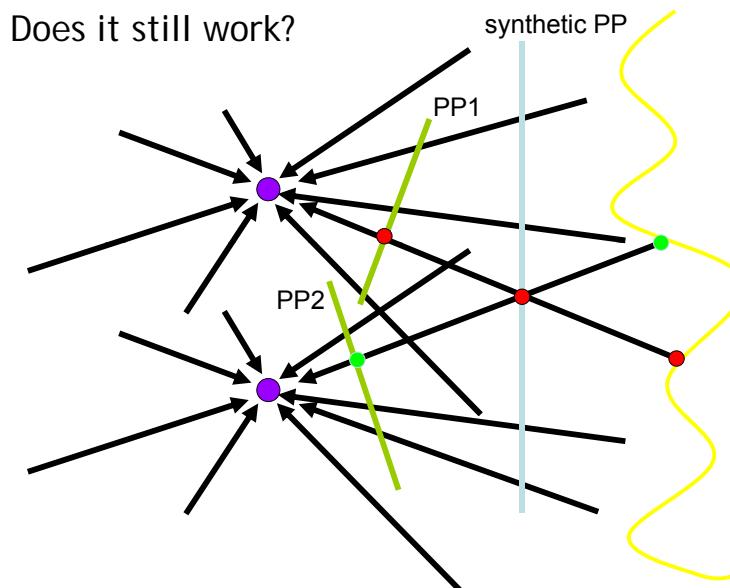


- PP3 is a projection plane of both centers of projection, so we are OK!
- This is how big aerial photographs are made

## Changing camera center

DigiVFX

- Does it still work?



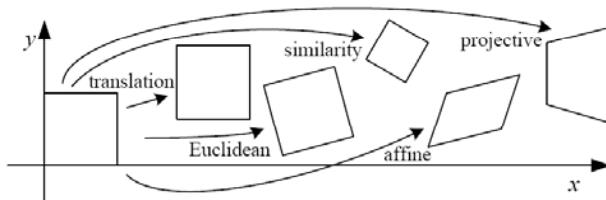
## Motion models

DigiVFX

- Parametric models as the assumptions on the relation between two images.

## 2D Motion models

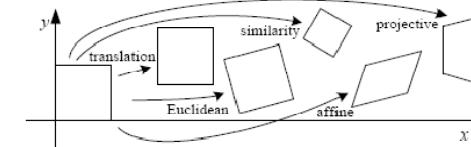
DigiVFX



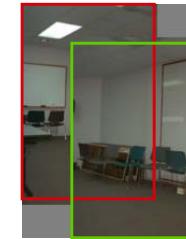
Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$[ I \mid t ]_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$[ R \mid t ]_{2 \times 3}$	3	lengths + ...	
similarity	$[ sR \mid t ]_{2 \times 3}$	4	angles + ...	
affine	$[ A ]_{2 \times 3}$	6	parallelism + ...	
projective	$[ \tilde{H} ]_{3 \times 3}$	8	straight lines	

## Motion models

DigiVFX

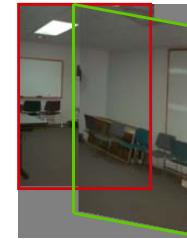


Translation



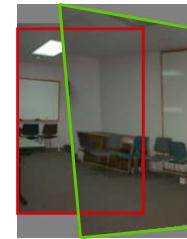
2 unknowns

Affine



6 unknowns

Perspective



8 unknowns

3D rotation

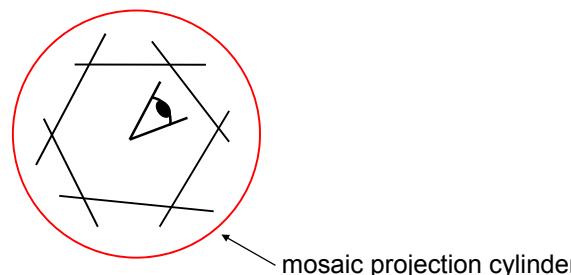


3 unknowns

## A case study: cylindrical panorama

DigiVFX

- What if you want a 360° field of view?



## Cylindrical panoramas

DigiVFX



- Steps

- Reproject each image onto a cylinder
- Blend
- Output the resulting mosaic

## Cylindrical panorama

DigiVFX

1. Take pictures on a tripod (or handheld)
2. Warp to cylindrical coordinate
3. Compute pairwise alignments
4. Fix up the end-to-end alignment
5. Blending
6. Crop the result and import into a viewer

It is required to do radial distortion correction for better stitching results!

## Taking pictures

DigiVFX



Kaidan panoramic tripod head

## Translation model

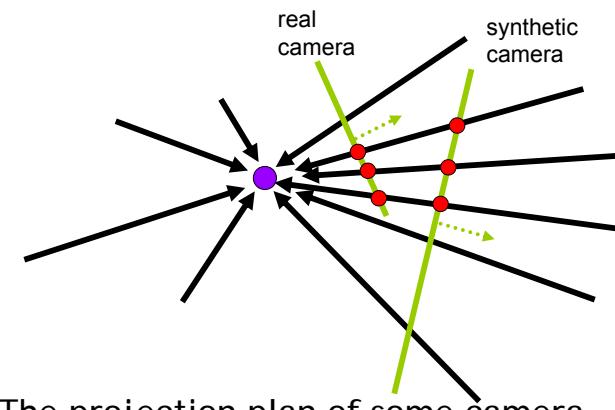
DigiVFX



Try to align this in PaintShop Pro

## Where should the synthetic camera be

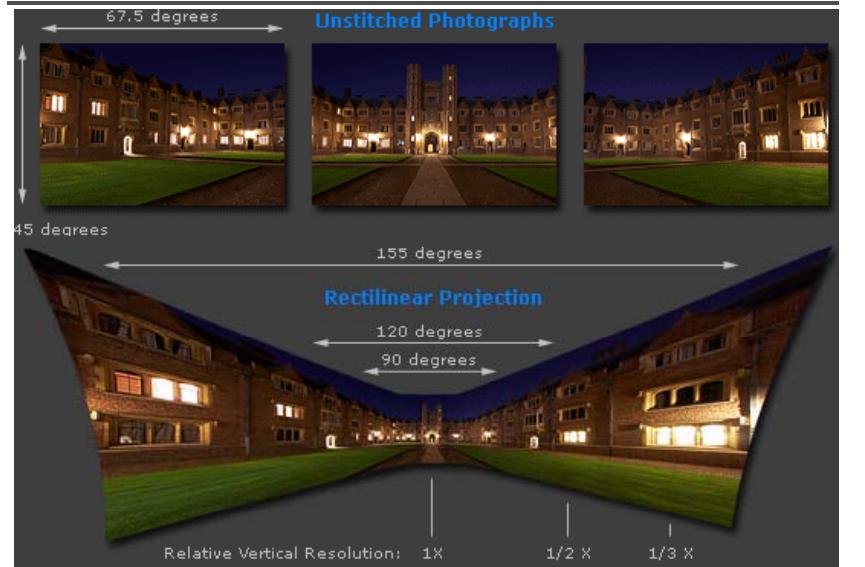
DigiVFX



- The projection plan of some camera
- Onto a cylinder

## Cylindrical projection

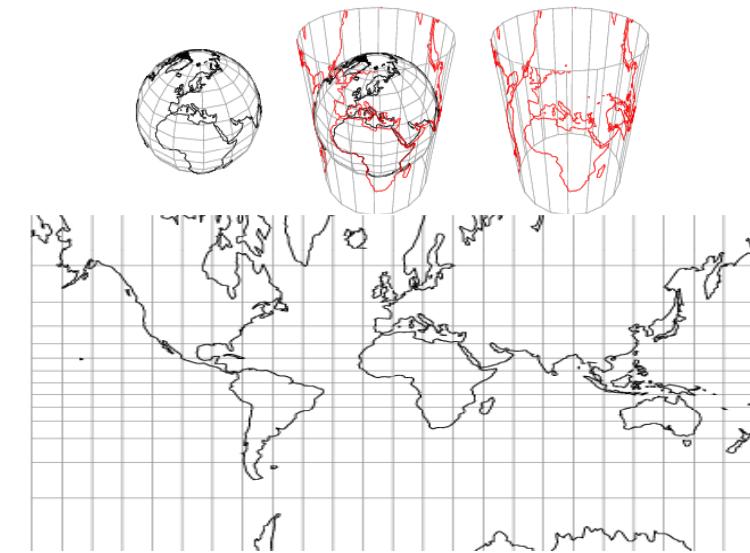
DigiVFX



Adopted from <http://www.cambridgeincolour.com/tutorials/image-projections.htm>

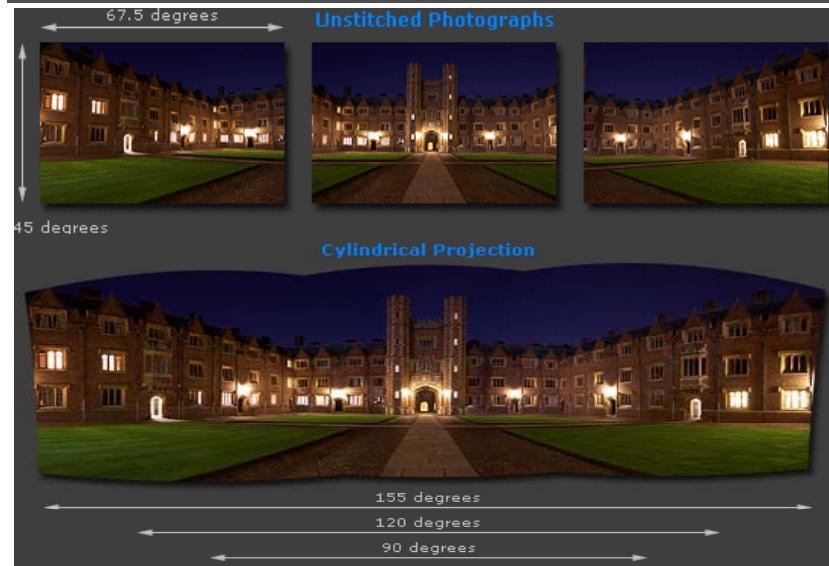
## Cylindrical projection

DigiVFX



## Cylindrical projection

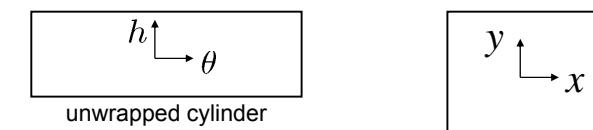
DigiVFX



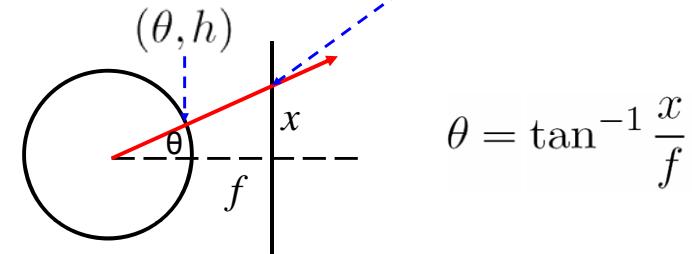
Adopted from <http://www.cambridgeincolour.com/tutorials/image-projections.htm>

## Cylindrical projection

DigiVFX

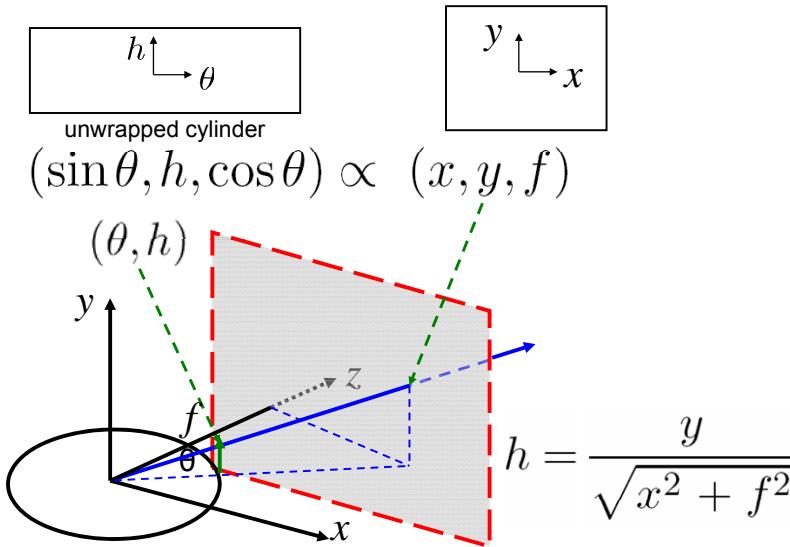


$$(\sin \theta, h, \cos \theta) \propto (x, y, f)$$



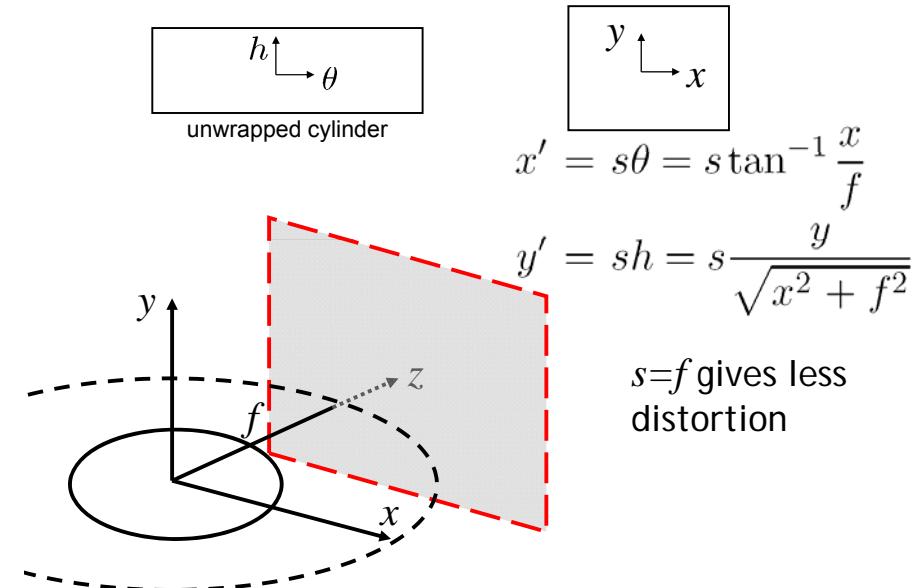
## Cylindrical projection

DigiVFX



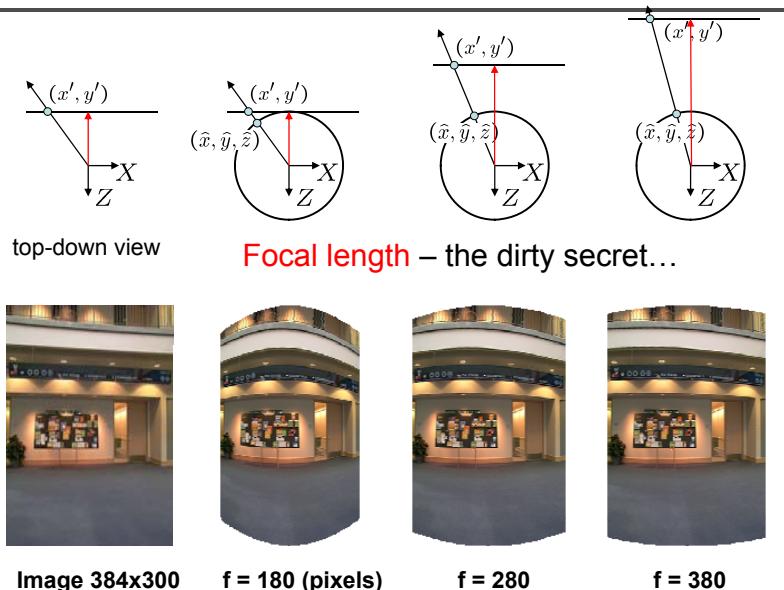
## Cylindrical projection

DigiVFX



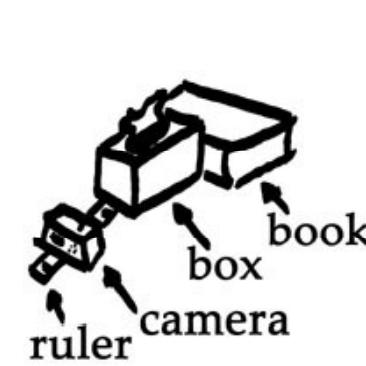
## Cylindrical reprojection

DigiVFX



## A simple method for estimating $f$

DigiVFX



Or, you can use other software, such as AutoStitch, to help.

## Input images

DigiVFX



## Cylindrical warping

DigiVFX



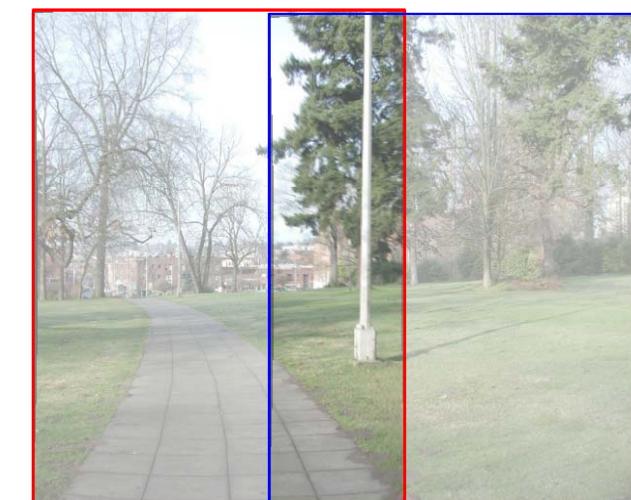
## Blending

DigiVFX

- Why blending: parallax, lens distortion, scene motion, exposure difference

## Blending

DigiVFX



## Blending

DigiVFX



## Blending

DigiVFX



## Assembling the panorama

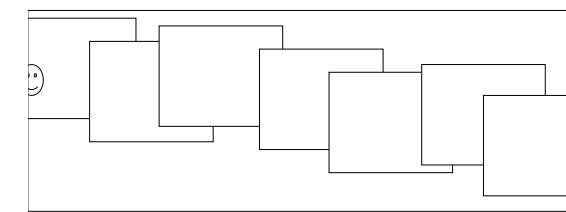
DigiVFX



- Stitch pairs together, blend, then crop

## Problem: Drift

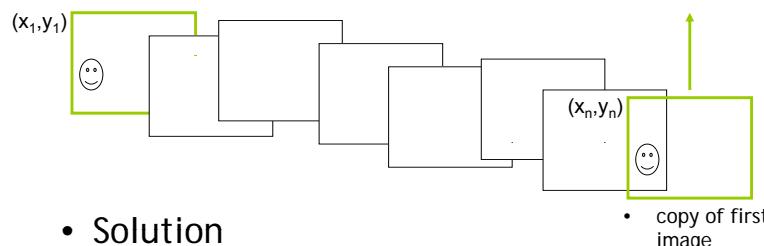
DigiVFX



- Error accumulation
  - small errors accumulate over time

## Problem: Drift

DigiVFX



- Solution

- add another copy of first image at the end
- there are a bunch of ways to solve this problem
  - add displacement of  $(y_1 - y_n)/(n - 1)$  to each image after the first
  - compute a global warp:  $y' = y + ax$
  - run a big optimization problem, incorporating this constraint
    - best solution, but more complicated
    - known as "bundle adjustment"

## End-to-end alignment and crop

DigiVFX



## Viewer: panorama

DigiVFX



example: <http://www.cs.washington.edu/education/courses/cse590ss/01wi/projects/project1/students/dougz/index.html>

## Viewer: texture mapped model

DigiVFX



example: <http://www.panoramas.dk/>

## Cylindrical panorama

DigiVFX

1. Take pictures on a tripod (or handheld)
2. Warp to cylindrical coordinate
3. Compute pairwise alignments
4. Fix up the end-to-end alignment
5. Blending
6. Crop the result and import into a viewer

## Determine pairwise alignment?

DigiVFX

- Feature-based methods: only use feature points to estimate parameters
- We will study the “Recognising panorama” paper published in ICCV 2003
- Run SIFT (or other feature algorithms) for each image, find feature matches.

## Determine pairwise alignment

DigiVFX

- $p' = Mp$ , where  $M$  is a transformation matrix,  $p$  and  $p'$  are feature matches
- It is possible to use more complicated models such as affine or perspective
- For example, assume  $M$  is a  $2 \times 2$  matrix

$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}$$

- Find  $M$  with the least square error

$$\sum_{i=1}^n (Mp - p')^2$$

## Determine pairwise alignment

DigiVFX

$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} \quad \begin{aligned} x_1 m_{11} + y_1 m_{12} &= x'_1 \\ x_1 m_{21} + y_1 m_{22} &= y'_1 \end{aligned}$$

- Overdetermined system

$$\begin{pmatrix} x_1 & y_1 & 0 & 0 \\ 0 & 0 & x_1 & y_1 \\ x_2 & y_2 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ x_n & y_n & 0 & 0 \\ 0 & 0 & x_n & y_n \end{pmatrix} \begin{pmatrix} m_{11} \\ m_{12} \\ m_{21} \\ m_{22} \end{pmatrix} = \begin{pmatrix} x'_1 \\ y'_1 \\ x'_2 \\ \vdots \\ x'_n \\ y'_n \end{pmatrix}$$

## Normal equation

Given an overdetermined system

$$\mathbf{Ax} = \mathbf{b}$$

the normal equation is that which minimizes the sum of the square differences between left and right sides

$$\mathbf{A}^T \mathbf{Ax} = \mathbf{A}^T \mathbf{b}$$

Why?

## Normal equation

$$\mathbf{Ax} - \mathbf{b} = \begin{bmatrix} \sum_{j=1}^m a_{1j}x_j \\ \vdots \\ \sum_{j=1}^m a_{ij}x_j \\ \vdots \\ \sum_{j=1}^m a_{nj}x_j \end{bmatrix} - \begin{bmatrix} b_1 \\ \vdots \\ b_i \\ \vdots \\ b_n \end{bmatrix} = \begin{bmatrix} \left( \sum_{j=1}^m a_{1j}x_j \right) - b_1 \\ \vdots \\ \left( \sum_{j=1}^m a_{ij}x_j \right) - b_i \\ \vdots \\ \left( \sum_{j=1}^m a_{nj}x_j \right) - b_n \end{bmatrix}$$

$$E(\mathbf{x}) = (\mathbf{Ax} - \mathbf{b})^2 = \sum_{i=1}^n \left[ \left( \sum_{j=1}^m a_{ij}x_j \right) - b_i \right]^2$$

## Normal equation

$$E(\mathbf{x}) = (\mathbf{Ax} - \mathbf{b})^2$$

$$\begin{bmatrix} a_{11} & \dots & a_{1m} \\ \vdots & & \vdots \\ \vdots & & \vdots \\ \vdots & & \vdots \\ a_{n1} & \dots & a_{nm} \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix} = \begin{bmatrix} b_1 \\ \vdots \\ b_n \end{bmatrix}$$

$n \times m$ ,  $n$  equations,  $m$  variables

## Normal equation

$$E(\mathbf{x}) = (\mathbf{Ax} - \mathbf{b})^2 = \sum_{i=1}^n \left[ \left( \sum_{j=1}^m a_{ij}x_j \right) - b_i \right]^2$$
$$0 = \frac{\partial E}{\partial x_1} = \sum_{i=1}^n 2 \left[ \left( \sum_{j=1}^m a_{ij}x_j \right) - b_i \right] a_{i1}$$
$$= 2 \sum_{i=1}^n a_{i1} \sum_{j=1}^m a_{ij}x_j - 2 \sum_{i=1}^n a_{i1}b_i$$

$$0 = \frac{\partial E}{\partial \mathbf{x}} = 2(\mathbf{A}^T \mathbf{Ax} - \mathbf{A}^T \mathbf{b}) \rightarrow \mathbf{A}^T \mathbf{Ax} = \mathbf{A}^T \mathbf{b}$$

## Normal equation

$$\begin{aligned} & (\mathbf{Ax} - \mathbf{b})^2 \\ &= (\mathbf{Ax} - \mathbf{b})^T (\mathbf{Ax} - \mathbf{b}) \\ &= ((\mathbf{Ax})^T - \mathbf{b}^T)(\mathbf{Ax} - \mathbf{b}) \\ &= (\mathbf{x}^T \mathbf{A}^T - \mathbf{b}^T)(\mathbf{Ax} - \mathbf{b}) \\ &= \mathbf{x}^T \mathbf{A}^T \mathbf{Ax} - \mathbf{b}^T \mathbf{Ax} - \mathbf{x}^T \mathbf{A}^T \mathbf{b} + \mathbf{b}^T \mathbf{b} \\ &= \mathbf{x}^T \mathbf{A}^T \mathbf{Ax} - (\mathbf{A}^T \mathbf{b})^T \mathbf{x} - (\mathbf{A}^T \mathbf{b})^T \mathbf{x} + \mathbf{b}^T \mathbf{b} \\ &\frac{\partial E}{\partial \mathbf{x}} = 2 \mathbf{A}^T \mathbf{Ax} - 2 \mathbf{A}^T \mathbf{b} \end{aligned}$$

## Determine pairwise alignment

- $\mathbf{p}' = \mathbf{Mp}$ , where  $\mathbf{M}$  is a transformation matrix,  $\mathbf{p}$  and  $\mathbf{p}'$  are feature matches
- For translation model, it is easier.

$$E = \sum_{i=1}^n [(m_1 + x_i - x'_i)^2 + (m_2 + y_i - y'_i)^2]$$

$$0 = \frac{\partial E}{\partial m_1}$$

- What if the match is false? Avoid impact of outliers.

## RANSAC

- RANSAC = Random Sample Consensus
- An algorithm for robust fitting of models in the presence of many data outliers
- Compare to robust statistics
- Given  $N$  data points  $x_i$ , assume that majority of them are generated from a model with parameters  $\Theta$ , try to recover  $\Theta$ .

## RANSAC algorithm

- Run  $k$  times: How many times?
- (1) draw  $n$  samples randomly How big?  
Smaller is better
- (2) fit parameters  $\Theta$  with these  $n$  samples
- (3) for each of other  $N-n$  points, calculate its distance to the fitted model, count the number of inlier points  $c$
- Output  $\Theta$  with the largest  $c$
- How to define?  
Depends on the problem.

## How to determine k

p: probability of real inliers

P: probability of success after k trials

$$P = 1 - (1 - p^n)^k$$

$n$  samples are all inliers  
a failure  
failure after  $k$  trials

$$k = \frac{\log(1 - P)}{\log(1 - p^n)}$$

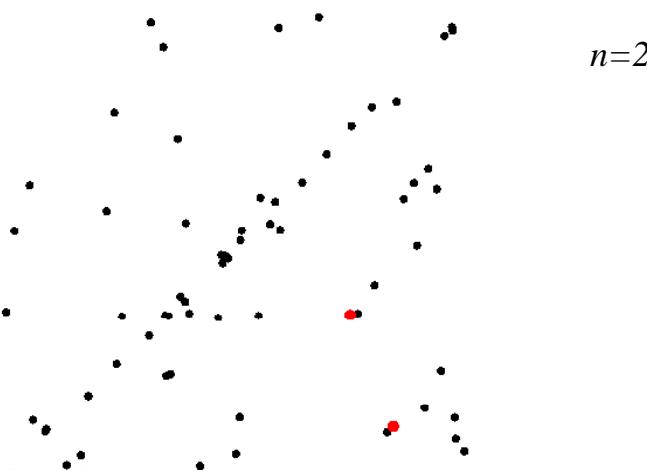
for  $P=0.99$

$n$	$p$	$k$
3	0.5	35
6	0.6	97
6	0.5	293

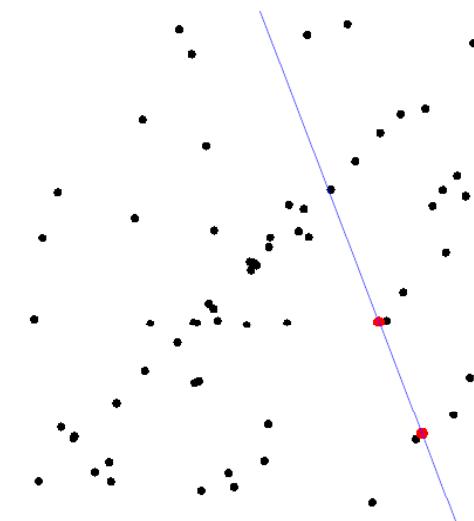
## Example: line fitting



## Example: line fitting

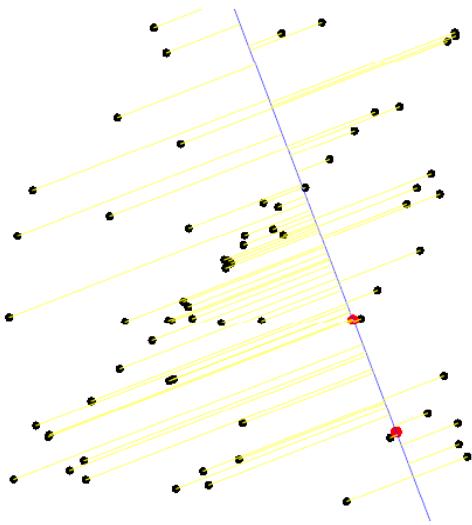


## Model fitting



## Measure distances

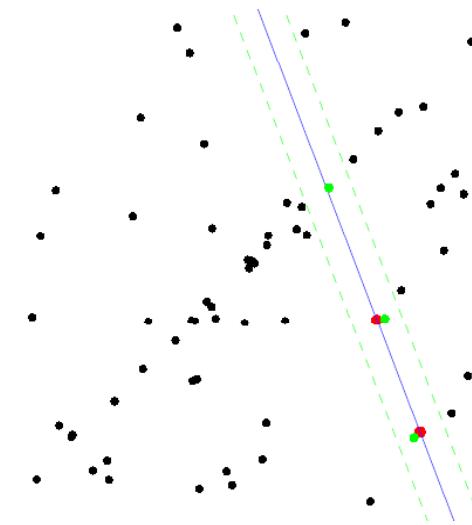
DigiVFX



## Count inliers

DigiVFX

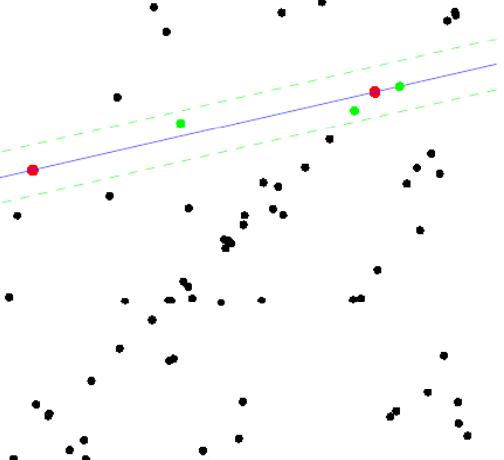
$c=3$



## Another trial

DigiVFX

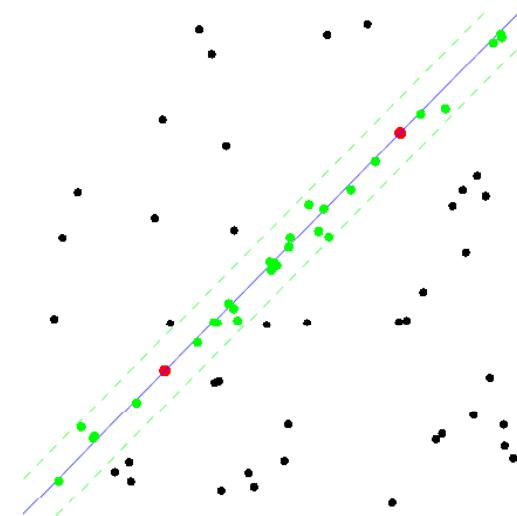
$c=3$



## The best model

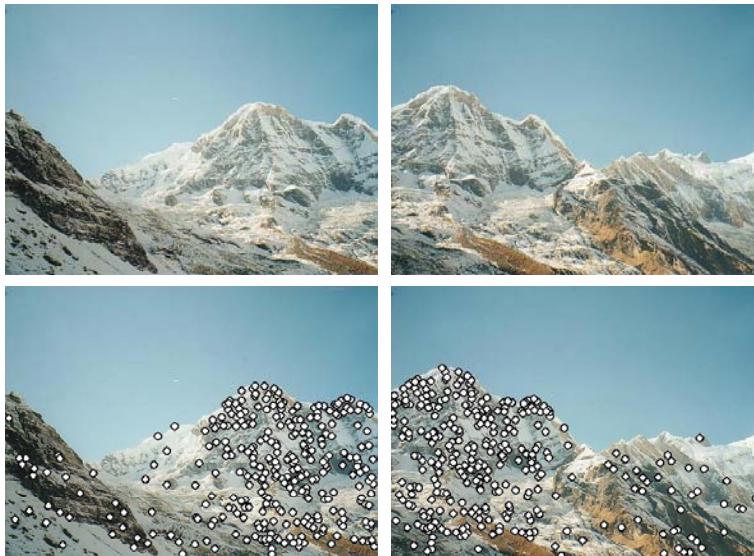
DigiVFX

$c=15$



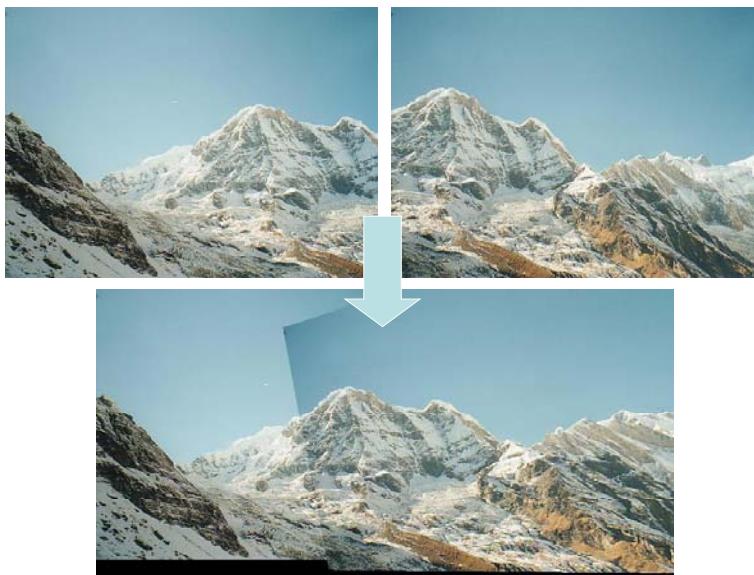
## RANSAC for Homography

DigiVFX



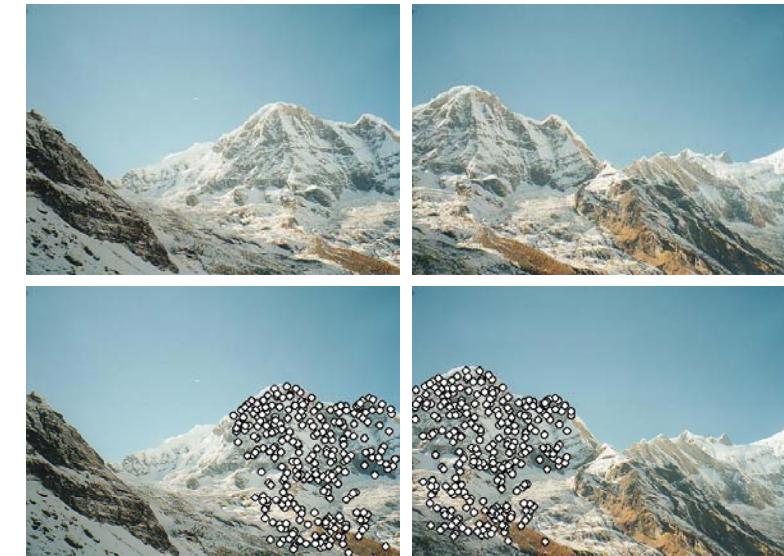
## RANSAC for Homography

DigiVFX



## RANSAC for Homography

DigiVFX



## Applications of panorama in VFX

DigiVFX

- Background plates
- Image-based lighting

## Troy (image-based lighting)

DigiVFX



[http://www.cgnetworks.com/story\\_custom.php?story\\_id=2195&page=4](http://www.cgnetworks.com/story_custom.php?story_id=2195&page=4)

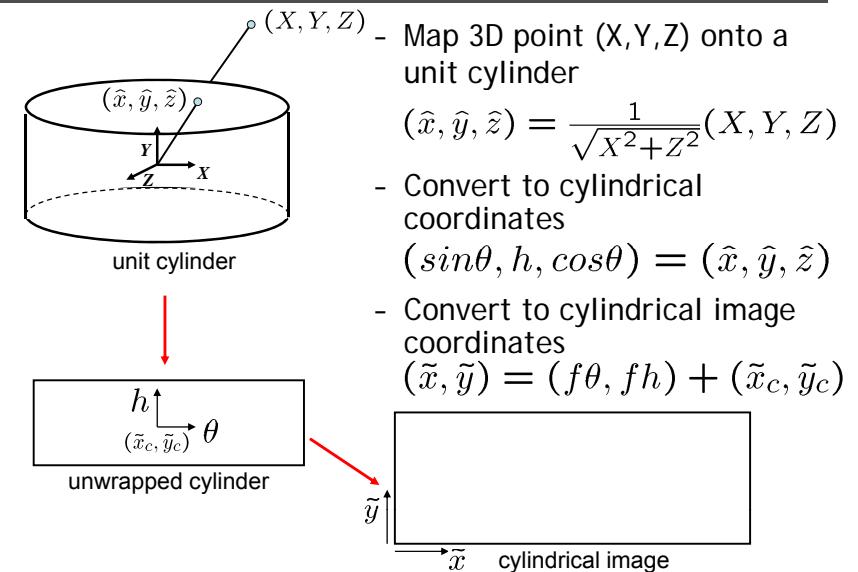
## Spiderman 2 (background plate)

DigiVFX



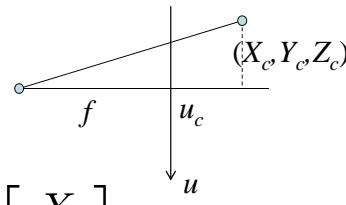
## Cylindrical projection

DigiVFX



## 3D → 2D perspective projection

DigiVFX



$$\begin{bmatrix} X_c \\ Y_c \\ Z_c \end{bmatrix} = [\mathbf{R}]_{3 \times 3} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \mathbf{t}$$

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \sim \begin{bmatrix} U \\ V \\ W \end{bmatrix} = \begin{bmatrix} f & 0 & u_c \\ 0 & f & v_c \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X_c \\ Y_c \\ Z_c \end{bmatrix}$$

## Direct vs feature-based

DigiVFX

- Direct methods use all information and can be very accurate, but they depend on the fragile “brightness constancy” assumption
- Iterative approaches require initialization
- Not robust to illumination change and noise images
- In early days, direct method is better.
  
- Feature based methods are now more robust and potentially faster
- Even better, it can recognize panorama without initialization

## Reference

DigiVFX

- Richard Szeliski, [Image Alignment and Stitching](#), unpublished draft, 2005.
- R. Szeliski and H.-Y. Shum. [Creating full view panoramic image mosaics and texture-mapped models](#), SIGGRAPH 1997, pp251-258.
- M. Brown, D. G. Lowe, [Recognising Panoramas](#), ICCV 2003.

## TODO

DigiVFX

- Bundle adjustment
- LM method
- Direct method vs feature-based method
  
- Frame-rate image alignment for stabilization
  
- Rick's CGA 1995 paper? LM method

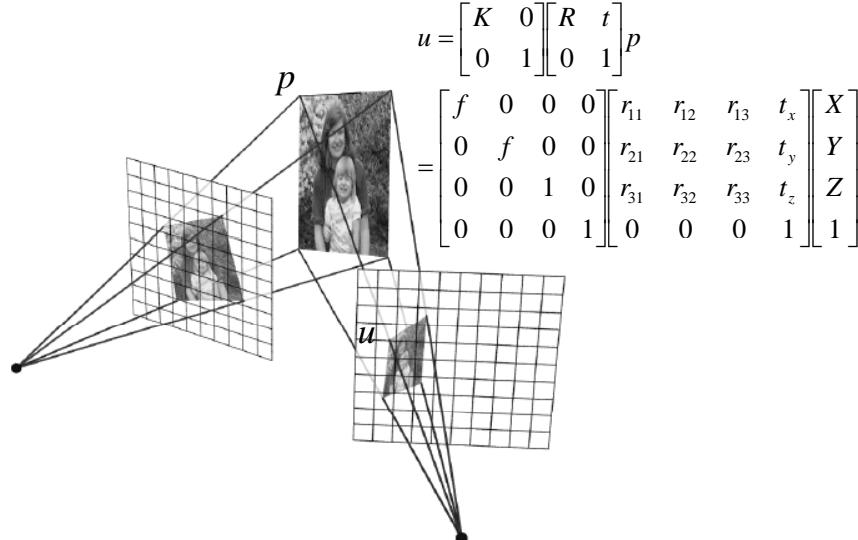
## Project #2 Image stitching

DigiVFX

- camera availability
- Tripod?
- <http://www.tawbaware.com/maxlyons/>
- <http://www.cs.washington.edu/education/courses/cse590ss/CurrentQtr/projects.htm>
- <http://www.cs.ubc.ca/~mbrown/panorama/panorama.html>

## 3D interpretation

DigiVFX



## blending

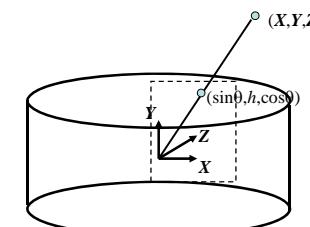
DigiVFX

- Alpha-blending
- Photomontage
- Poisson blending
- Adelson's pyramid blending
- Hdr?

## Cylindrical warping

DigiVFX

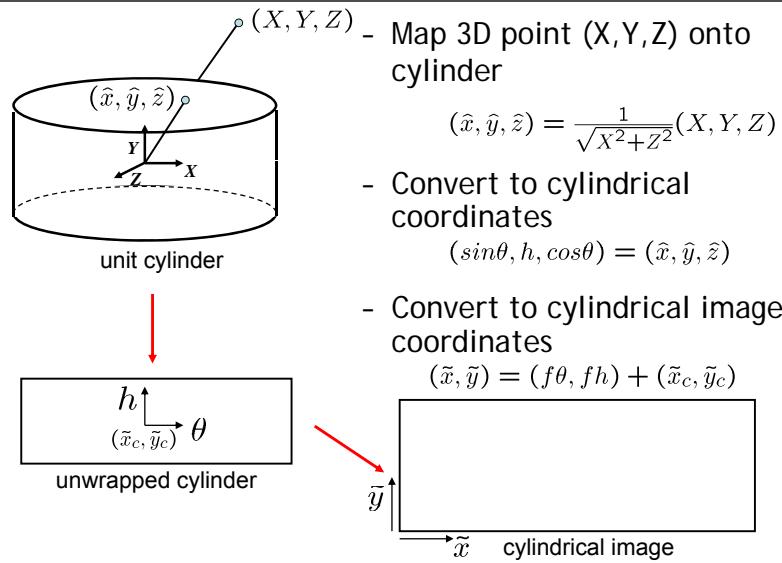
- Given focal length  $f$  and image center  $(x_c, y_c)$



$$\begin{aligned}\theta &= (x_{cyl} - x_c)/f \\ h &= (y_{cyl} - y_c)/f \\ \hat{x} &= \sin \theta \\ \hat{y} &= h \\ \hat{z} &= \cos \theta \\ x &= f\hat{x}/\hat{z} + x_c \\ y &= f\hat{y}/\hat{z} + y_c\end{aligned}$$

## Cylindrical projection

DigiVFX



- Map 3D point  $(X, Y, Z)$  onto cylinder

$$(\hat{x}, \hat{y}, \hat{z}) = \frac{1}{\sqrt{X^2+Z^2}}(X, Y, Z)$$

- Convert to cylindrical coordinates

$$(\sin\theta, h, \cos\theta) = (\hat{x}, \hat{y}, \hat{z})$$

- Convert to cylindrical image coordinates

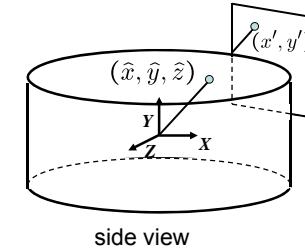
$$(\tilde{x}, \tilde{y}) = (f\theta, fh) + (\tilde{x}_c, \tilde{y}_c)$$



## Cylindrical reprojection

DigiVFX

- How to map from a cylinder to a planar image?



side view

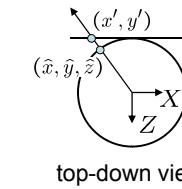
- Apply camera projection matrix

•  $w$  = image width,  $h$  = image height

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} -f & 0 & w/2 & 0 \\ 0 & -f & h/2 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} \hat{x} \\ \hat{y} \\ \hat{z} \\ 1 \end{bmatrix}$$

- Convert to image coordinates

• divide by third coordinate ( $w$ )



top-down view

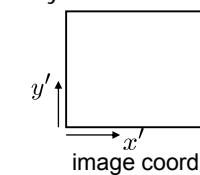
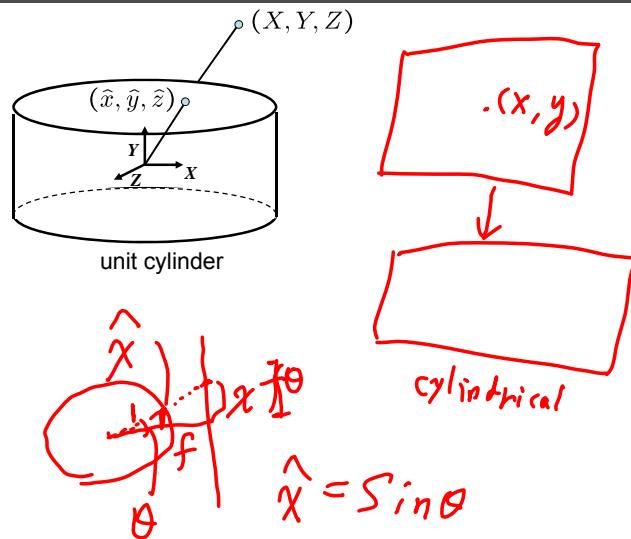


image coords

## Cylindrical projection

DigiVFX



## Levenberg-Marquardt Method

DigiVFX

## Alignment

- a rotation of the camera is a translation of the cylinder!

$$\begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \sum_{x,y} I_x (J(x,y) - I(x,y)) \\ \sum_{x,y} I_y (J(x,y) - I(x,y)) \end{bmatrix}$$

DigiVFX

## LucasKanadeStep

```
void LucasKanadeStep(CBytelImage& img1, CBytelImage& img2, float t[2]) {
    // Transform the image
    Translation(img2, img2t, t);

    // Compute the gradients and summed error by comparing img1 and img2t
    double A[2][2], b[2];
    for (int y = 1; y < height-1; y++) { // ignore borders
        for (int x = 1; x < width-1; x++) {
            // If both have full alphas, then compute and accumulate the error
            double e = img2t.Pixel(x, y, k) - img1.Pixel(x, y, k);
            // Accumulate the matrix entries
            double gx = 0.5*(img2t.Pixel(x+1, y, k) - img2t.Pixel(x-1, y, k));
            double gy = 0.5*(img2t.Pixel(x, y+1, k) - img2t.Pixel(x, y-1, k));

            A[0][0] += gx*gx; A[0][1] += gx*gy;
            A[1][0] += gx*gy; A[1][1] += gy*gy;

            b[0] += e*gx; b[1] += e*gy;
        }
    }
}
```

## LucasKanadeStep (cont.)

```
// Solve for the update At=b and update the vector

double det = 1.0 / (A[0][0]*A[1][1] - A[1][0]*A[0][1]);
t[0] += (A[1][1]*b[0] - A[1][0]*b[1]) * det;
t[1] += (A[0][0]*b[1] - A[1][0]*b[0]) * det;
}
```

DigiVFX

DigiVFX

## PyramidLucasKanade

```
void PyramidLucasKanade(CBytelImage& img1, CBytelImage& img2, float t[2],
                         int nLevels, int nLucasKanadeSteps)
{
    CBytePyramid p1(img1); // Form the two pyramids
    CBytePyramid p2(img2);

    // Process in a coarse-to-fine hierarchy
    for (int l = nLevels-1; l >= 0; l--)
    {
        t[0] /= (1 << l); // scale the t vector
        t[1] /= (1 << l);
        CBytelImage& i1 = p1[l];
        CBytelImage& i2 = p2[l];

        for (int k = 0; k < nLucasKanadeSteps; k++)
            LucasKanadeStep(i1, i2, t);
        t[0] *= (1 << l); // restore the full scaling
        t[1] *= (1 << l);
    }
}
```

## Gaussian pyramid

DigiVFX



## Video matting

DigiVFX



alpha matte

## 2D Motion models

DigiVFX

- translation:  $x' = x + t$   $x = (x, y)$
- rotation:  $x' = R x + t$
- similarity:  $x' = s R x + t$
- affine:  $x' = A x + t$
- perspective:  $\underline{x}' \cong H \underline{x}$   $\underline{x} = (x, y, 1)$   
( $\underline{x}$  is a *homogeneous coordinate*)
- These all form a nested *group* (closed under composition w/ inv.)

## Recognising Panoramas

DigiVFX

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images

## Recognising Panoramas

DigiVFX

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



## Recognising Panoramas

DigiVFX

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



## Recognising Panoramas

DigiVFX

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



- 2D Rotations ( $q, f$ )
  - Ordering  $\not\Rightarrow$  matching images

## Recognising Panoramas

DigiVFX

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



- 2D Rotations ( $q, f$ )
  - Ordering  $\not\Rightarrow$  matching images



## Recognising Panoramas

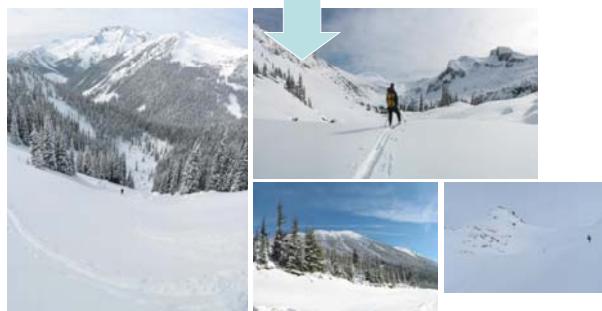
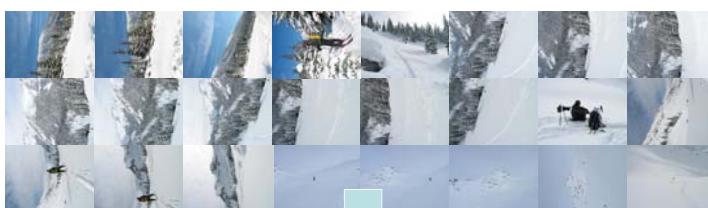
- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



- 2D Rotations ( $q, f$ )
  - Ordering  $\cancel{\Rightarrow}$  matching images



## Recognising Panoramas



## Probabilistic model for verification

- Compare probability that this set of RANSAC inliers/outliers was generated by a correct/false image match
- Choosing values for  $p_1$ ,  $p_0$  and  $p_{\min}$

$$n_i > 5.9 + 0.22n_f$$

## Overview

- SIFT Feature Matching
- Image Matching
- Bundle Adjustment
- Multi-band Blending

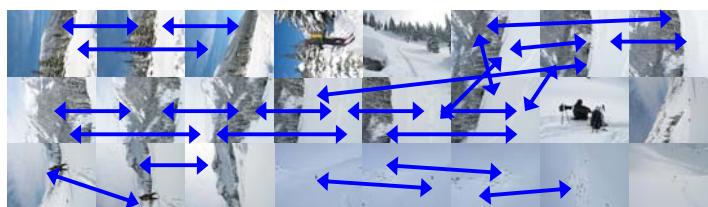
## Nearest Neighbour Matching

DigiVFX

- Find k-NN for each feature
  - $k \approx$  number of overlapping images (we use  $k = 4$ )
- Use k-d tree
  - k-d tree recursively bi-partitions data at mean in the dimension of maximum variance
  - Approximate nearest neighbours found in  $O(n \log n)$

## Finding the panoramas

DigiVFX

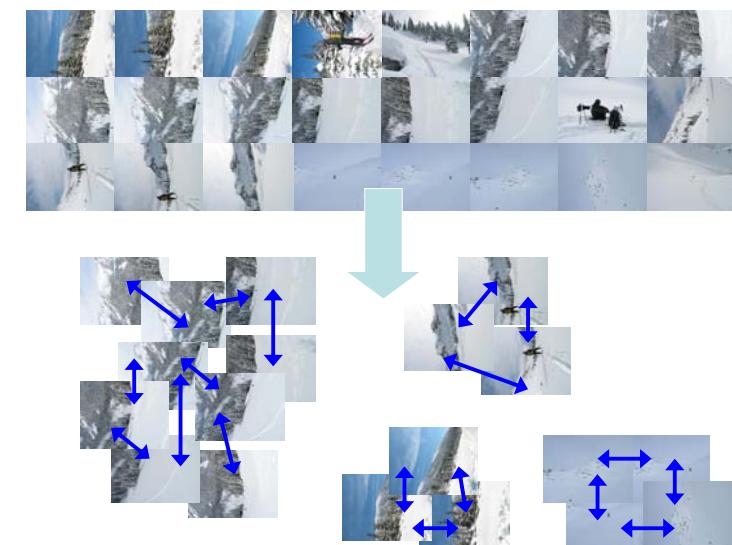


## Overview

- SIFT Feature Matching
- **Image Matching**
  - For each image, use RANSAC to select inlier features from 6 images with most feature matches
- Bundle Adjustment
- Multi-band Blending

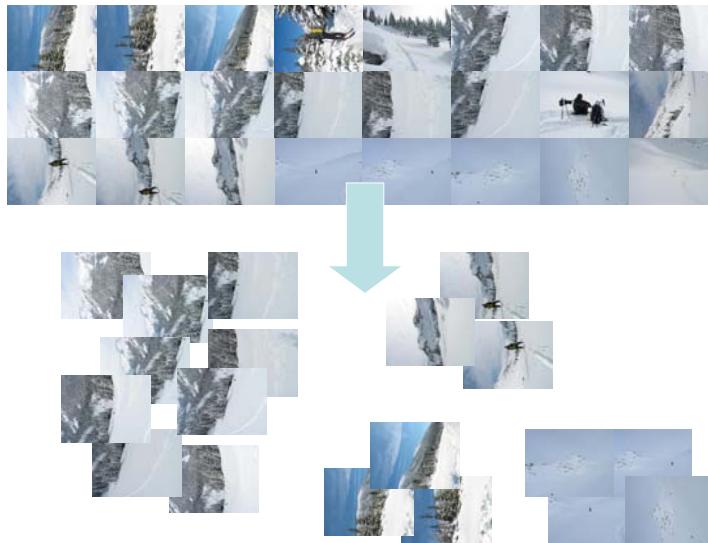
## Finding the panoramas

DigiVFX



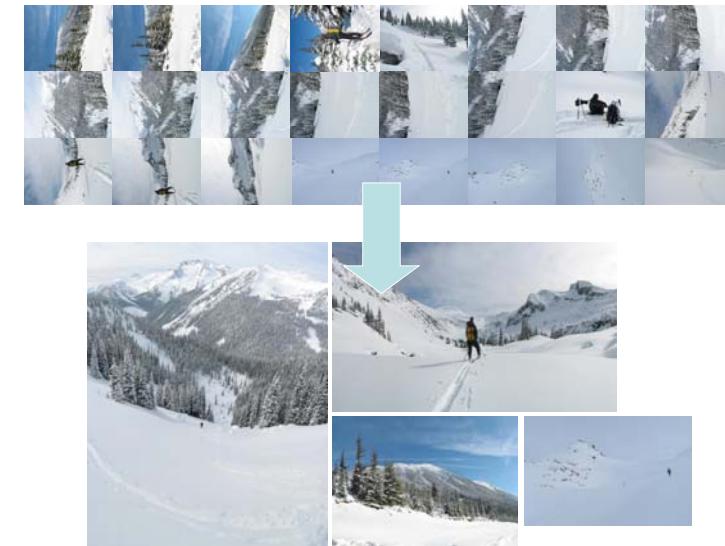
## Finding the panoramas

DigiVFX



## Finding the panoramas

DigiVFX



## Overview

DigiVFX

- SIFT Feature Matching
- Image Matching
- **Bundle Adjustment**
- Multi-band Blending

## Homography for Rotation

DigiVFX

- Parameterise each camera by rotation and focal length

$$\mathbf{R}_i = e^{[\boldsymbol{\theta}_i]_\times}, \quad [\boldsymbol{\theta}_i]_\times = \begin{bmatrix} 0 & -\theta_{i3} & \theta_{i2} \\ \theta_{i3} & 0 & -\theta_{i1} \\ -\theta_{i2} & \theta_{i1} & 0 \end{bmatrix}$$

$$\mathbf{K}_i = \begin{bmatrix} f_i & 0 & 0 \\ 0 & f_i & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- This gives pairwise homographies

$$\tilde{\mathbf{u}}_i = \mathbf{H}_{ij} \tilde{\mathbf{u}}_j, \quad \mathbf{H}_{ij} = \mathbf{K}_i \mathbf{R}_i \mathbf{R}_j^T \mathbf{K}_j^{-1}$$

## Error function

- Sum of squared projection errors

$$e = \sum_{i=1}^n \sum_{j \in \mathcal{I}(i)} \sum_{k \in \mathcal{F}(i,j)} f(\mathbf{r}_{ij}^k)^2$$

- $n$  = #images
  - $\mathcal{I}(i)$  = set of image matches to image  $i$
  - $\mathcal{F}(i, j)$  = set of feature matches between images  $i, j$
  - $\mathbf{r}_{ij}^k$  = residual of  $k^{\text{th}}$  feature match between images  $i, j$
- 
- Robust err<sub>f(x)</sub> =  $\begin{cases} |\mathbf{x}|, & \text{if } |\mathbf{x}| < x_{max} \\ x_{max}, & \text{if } |\mathbf{x}| \geq x_{max} \end{cases}$

## Multi-band Blending

- Burt & Adelson 1983
  - Blend frequency bands over range  $\propto \lambda$



## Overview

- SIFT Feature Matching
- Image Matching
- Bundle Adjustment
- Multi-band Blending

## 2-band Blending

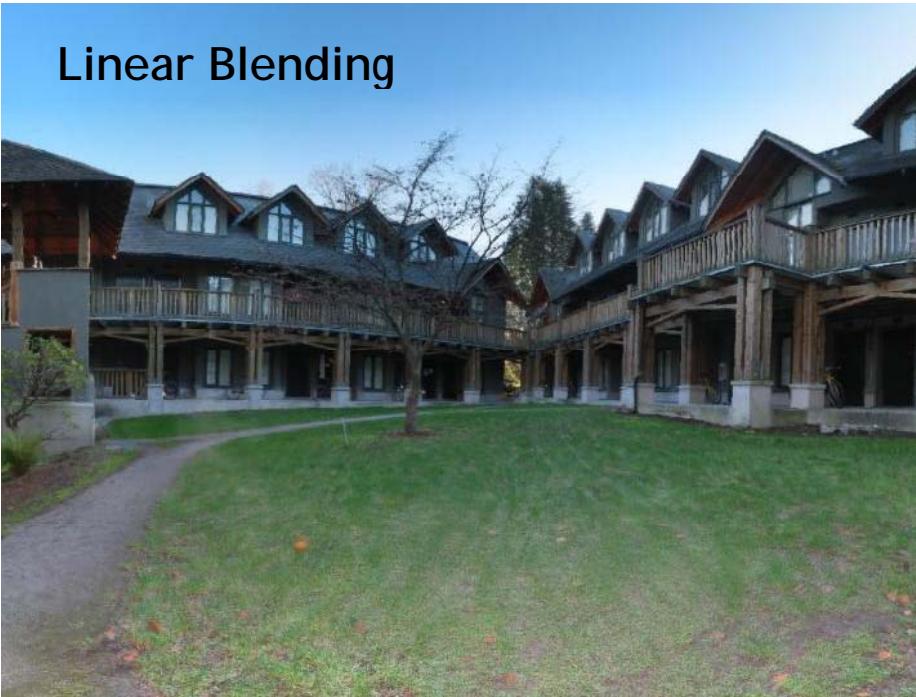


Low frequency ( $\lambda > 2$  pixels)



High frequency ( $\lambda < 2$  pixels)

## Linear Blending

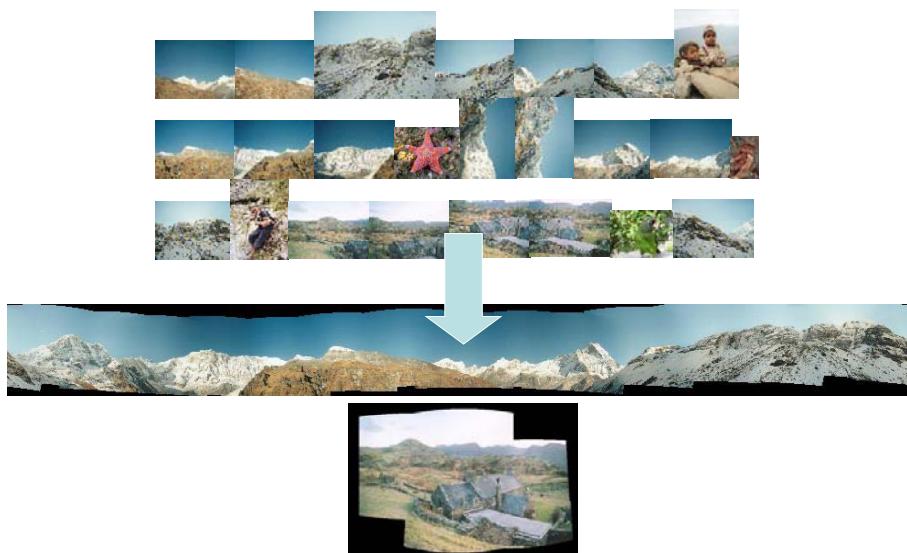


## 2-band Blending



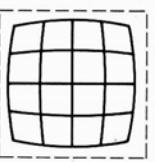
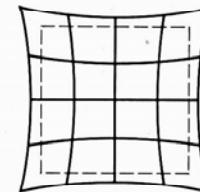
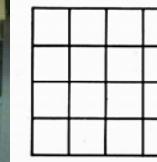
## Results

DigiVFX



## Distortion

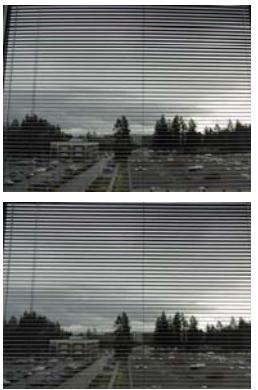
DigiVFX



- Radial distortion of the image
  - Caused by imperfect lenses
  - Deviations are most noticeable for rays that pass through the edge of the lens

## Radial correction

- Correct for “bending” in wide field of view lenses



$$\hat{r}^2 = \hat{x}^2 + \hat{y}^2$$

$$\hat{x}' = \hat{x}/(1 + \kappa_1 \hat{r}^2 + \kappa_2 \hat{r}^4)$$

$$\hat{y}' = \hat{y}/(1 + \kappa_1 \hat{r}^2 + \kappa_2 \hat{r}^4)$$

$$x = f\hat{x}'/\hat{z} + x_c$$

$$y = f\hat{y}'/\hat{z} + y_c$$